

# Curated File Guide: script.js

The public website runtime that renders the portfolio, search, project windows, browser history, rich content, links, and AI chat behavior.

## Before The Code: File Overview

### Abstract

script.js is the public website runtime. It is the JavaScript that turns the static portfolio files into an interactive recruiter-facing site. The HTML page provides stable containers and the JSON catalog provides saved content, but script.js decides what appears, how projects are filtered, how rich text renders, how files become links, how project sections open into full-window views, how browser Back and Refresh interact with those views, how search behaves like a real search tool, and how the AI chat gathers useful context before asking the backend.

The important distinction is that script.js is public and read-only. It should never expose builder-only editing controls. Its job is to present what has already been saved and parsed by the builder. That is why much of the file is about rendering, normalization, routing, search indexing, assistant context selection, and UI state rather than file writes.

### Introduction

When a visitor opens the portfolio, index.html loads styles.css, electronics-search.js, and script.js. script.js immediately grabs the DOM elements it will control: the project grid, search input, filters, profile fields, resume area, contact area, AI assistant, and footer. Then it waits for portfolio data. The data can come from an embedded window.\_\_PORTFOLIO\_CATALOG\_\_ object in previews or from projects.json on the published website. Once the catalog is loaded, script.js normalizes profile data, renders the hero and contact information, creates category filters, renders project cards, builds the search index, and restores any deep-linked project section.

This file is the main reason the same portfolio can feel alive without requiring a database-backed website. It uses static files as the source of truth, then builds an application-like experience in the browser. That makes the public website easy to host on GitHub Pages or Cloudflare, while still giving visitors search, project windows, AI questions, and responsive navigation.

### Background Information

The file works with a catalog generated by the builder parser. That catalog contains categories, projects, profile fields, site content, field styles, rich text blocks, fun facts, and optional site sections. The functions in script.js assume the catalog is already safe enough to display, but they still escape HTML where user-facing strings are inserted into the DOM. Rich text rendering is intentionally structured because portfolio content can include plain text, links, formulas, images, files, code blocks, captions, and nested project subsections.

The website also has a windowing model. Project cards appear inside categories. A project section button opens a full-screen dialog. Inside that dialog, subsections can be clicked like nested folders. Browser history is updated with the project id, section index, and resource path so the browser Back button can move through the same hierarchy. This makes project content feel more like a navigable portfolio than a long flat webpage.

Search is also broader than a normal text filter. The search index contains page sections, project names, category labels, tools, languages, rich text, files, URLs, captions, uploaded text files where fetchable, and electronics keyword expansions. The goal is that a recruiter can type an electronics phrase and land near the relevant artifact, even if the phrase appears inside a project section or file caption rather than on the top card.

Because the website is static, script.js must do work that a database-backed application might normally ask a server to do. It must decide which categories exist, which projects belong to each category, which sections are empty, which files should render as downloads, which links should open externally, and which project paths should be represented in browser history. The price of static hosting is that the browser runtime has to be disciplined. The benefit is that the public site remains easy to deploy, cache, mirror, and serve from GitHub Pages or Cloudflare without a long-running custom web server.

The file also sits between several other files. index.html gives it DOM anchors. styles.css gives visual meaning to the classes it writes. electronics-search.js can provide expanded electronics vocabulary. projects.json provides the portfolio catalog. The AI backend endpoint gives it remote intelligence when available. If one of those contracts changes, script.js is usually the first file that reveals the mismatch because rendering, search, or chat stops behaving correctly.

## Purpose Of The File

The purpose of script.js is to convert portfolio data into a polished public experience. The file decides not only what to render, but also what not to render. Empty sections are hidden. Overview content appears as the first readable content when present. Links that look like websites open as websites rather than downloads. Local files become downloadable resources. Images render inline only when they are meant to be shown. AI links are only displayed when related to the question context.

Without script.js, index.html would be mostly a static shell with placeholders. There would be no dynamic project directory, no generated categories, no rich project windows, no browser-aware section routing, no searchable uploaded file text, no AI context gathering, no responsive project rendering, and no dynamic profile/contact population.

## Primary Responsibilities

The first responsibility is state hydration. loadProjectCatalog reads the portfolio catalog, normalizes the data, stores it in module-level variables, and triggers rendering. This is the startup path for the public website.

The second responsibility is rendering. Functions such as renderProfile, renderSiteContent, renderFunFacts, renderSiteSections, projectCard, parsedSection, and parsedNodeContent turn structured catalog objects into actual HTML.

The third responsibility is interaction. The dialog functions create the full-window project section view, the routing functions keep browser history synchronized, the search functions build and display selectable results, and the AI functions manage the chat panel.

The fourth responsibility is interpretation. The assistant functions classify questions, gather relevant portfolio context, select public source links, decide whether a question is general or portfolio-specific, and prepare the request sent to the AI backend. This prevents the assistant from blindly dumping generic links for every question.

The fifth responsibility is restraint. Public visitors should see only finished, parsed, saved portfolio output. The file therefore hides empty content, avoids exposing builder edit controls, limits which sources are shown with AI answers, and uses normalized link behavior so a website link does not become a fake downloadable file. This restraint is what separates the public website from the local builder application.

The sixth responsibility is mobile and desktop continuity. The same project data must read cleanly on an iPhone and on a wide monitor. The responsive helper functions estimate project density and media presence, while the dialog and assistant functions keep windows and chat panels usable without letting one part of the page resize unrelated parts. script.js is therefore not only a renderer; it is also the public interaction policy for different devices.

## Important Data Objects And Why They Matter

categories, projects, siteSections, profile, siteContent, funFacts, fieldStyles, and the rich variants of those fields are the hydrated catalog state. They are not hardcoded content. They are populated from projects.json or preview data, which is why the same script can run for Maurice's site and also for someone else's freshly installed builder output.

searchableEntries is the in-memory search index. Every item in it has a title, type, context, text, optional project id, optional section index, optional URL, and normalized text. The search functions score and display these entries.

assistantChatHistory and assistantSourceMap are the AI interface memory. assistantChatHistory keeps recent conversation so follow-up questions make sense. assistantSourceMap maps clickable source buttons in an answer back to portfolio search results so a visitor can jump to the relevant project or file.

currentSearchResults and searchResultLimit are the active search session. They remember what the latest query produced and how many items the user wants to see. Without these values, pressing Enter, selecting a result, or changing the result limit would not know which result list is current.

sectionRouteKey and sectionRouteParams are the browser-history contract for project windows. They keep the public URL connected to the currently opened project section and nested subsection. The values are small, but they are what let Refresh reopen the same project view and Back return to the previous view.

## Top-Level Objects And Variables

This section explains the long-lived objects and variables before the function walkthrough. These are the values that give the file memory, configuration, API handles, DOM anchors, path boundaries, cache state, and behavior maps. They are explained by meaning, not by line number.

### Dom Reference Or Ui State

#### **grid**

grid points at the project-grid container. renderProjects writes category sections and project cards into this element, making it the public project's visible mounting point. It anchors UI behavior to a specific browser element or window state. Functions that reference it include mediaGrid, renderSiteSections, parsedChildCards, projectCard, renderProjects, restoreSectionRouteAfterCatalogLoad, loadProjectCatalog.

#### **searchInput**

searchInput points at the project search box. Search handlers read its value, build dropdown results, apply highlights, and decide which projects remain visible. It anchors UI behavior to a specific browser element or window state. Functions that reference it include queueSearchHighlights, ensureSearchPanel, renderSearchResults, updateSearchDropdown, showSearchPanelIfNeeded, renderProjects, restoreSectionRouteAfterCatalogLoad.

#### **filterButtons**

filterButtons stores the current filter button elements. It is updated after dynamic filters are rendered because the original NodeList would not include newly created buttons. It anchors UI behavior to a specific browser element or window state. Functions that reference it include setActiveFilter, bindFilterButtons.

#### **aiClearChatButton**

aiClearChatButton is the explicit Clear chat control. It resets conversation history, source maps, and the visible log. It anchors UI behavior to a specific browser element or window state. Functions that reference it include updateAssistantPanelGrowth, restoreSectionRouteAfterCatalogLoad.

### Portfolio Data State

#### **projectFilters**

projectFilters points at the category filter container. renderCategoryFilters fills it with filter buttons derived from the actual catalog categories. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderCategoryFilters.

#### **projectCount**

projectCount displays the number of visible or total projects. It is updated after catalog load and render operations so the summary band matches the data. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderProjects.

#### **projectTrackCount**

projectTrackCount displays how many project categories/tracks exist. It is derived from the saved categories rather than hardcoded. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderCategoryFilters.

#### **projectTrackLabels**

projectTrackLabels displays the category names. It helps the top summary reflect Analog, Digital, Embedded, or any custom categories the builder creates. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderCategoryFilters.

## **resumeSection**

resumeSection is the public resume panel. renderProfile hides or shows it depending on whether a resume URL exists. Functions that reference it include renderProfile.

## **profilePhoto**

profilePhoto is the large contact-section profile image. It is separate from the header avatar so desktop/mobile layouts can size them differently. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderProfile.

## **projects**

projects is the runtime project array. Rendering, search, assistant context, and project windows all read from it. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include addSiteSectionSearchEntries, rebuildSearchIndex, assistantNamedProjectIds, assistantQuestionHasPortfolioIntent, assistantQuestionTokens, assistantNamedProjectTopicTokens, assistantContextForQuestion, assistantFallbackResults, assistantLocalAnswer, sectionRouteUrl.

## **siteSections**

siteSections stores custom main-page sections such as personal profile, hobbies, sports, social media, or professional material. Functions that reference it include addSiteSectionSearchEntries, assistantPublicProfileLinks, assistantContextForQuestion, renderSiteSections, loadProjectCatalog.

## **profile**

profile stores owner identity: name, email, phone, images, resume, GitHub, LinkedIn, and other public contact details. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderProfile, isWebsiteLinkItem, projectResponsiveClass, responsiveStyleValues, addSiteSectionSearchEntries, assistantUrlsFromTextValue, assistantPublicProfileLinks, assistantQuestionAllowsPublicLookup, assistantContextForQuestion, assistantFallbackResults.

## **profileRich**

profileRich stores rich versions of profile/contact fields so styled text can survive from builder to website. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Functions that reference it include renderProfile, loadProjectCatalog.

## **activeSectionDialogDrag**

activeSectionDialogDrag stores the drag state while a project window is being moved. It anchors UI behavior to a specific browser element or window state. Functions that reference it include beginSectionDialogDrag, moveSectionDialogDrag, endSectionDialogDrag.

## **activeSectionDialogResize**

activeSectionDialogResize stores the resize state while a project window edge is being dragged. It anchors UI behavior to a specific browser element or window state. Functions that reference it include beginSectionDialogResize, moveSectionDialogResize, endSectionDialogResize.

## **sectionDialogDragEnabled**

sectionDialogDragEnabled prevents drag handlers from being installed more than once. It anchors UI behavior to a specific browser element or window state. Functions that reference it include enableSectionDialogDrag.

## **isApplyingSectionRoute**

isApplyingSectionRoute prevents browser history updates from recursively triggering themselves while a saved route is being restored. Functions that reference it include syncSectionRouteToHistory, clearSectionRouteInHistory, applySectionRoute.

### **sectionRouteKey**

sectionRouteKey identifies portfolio section entries in browser history state. Functions that reference it include sectionRouteState, sectionRouteFromState.

### **sectionRouteParams**

sectionRouteParams defines the URL parameter names for project id, section index, and nested subsection path. Visible object fields include path, project, section. Those fields form the contract consumed by related functions. Functions that reference it include sectionBaseUrl, sectionRouteUrl, sectionRouteFromLocation.

### **defaultProfile**

defaultProfile is the empty-owner fallback. It defines all expected profile fields so renderProfile can run even before the user enters contact details. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Visible object fields include brandImage, brandText, contactIntro, displayName, email, githubUrl, herolImage, linkedinUrl, phone, portfolioLabel, profileImage, resumeUrl, websiteUrl. Those fields form the contract consumed by related functions. Functions that reference it include normalizeProfile.

## **Temporary Calculation**

### **year**

year points at the footer year span. Startup code can set it from the current date so the footer remains current without manual HTML edits. Functions that reference it include renderProfile.

### **funFactsCallout**

funFactsCallout is the near-top personality callout. renderFunFacts only reveals it when saved fun facts exist. Functions that reference it include renderFunFacts.

### **heroEyebrow**

heroEyebrow is the small label above the hero title. renderSiteContent can replace it from the builder's front page settings. Functions that reference it include normalizeSiteContent, renderSiteContent, addSiteSectionSearchEntries.

### **heroTitle**

heroTitle is the main public headline. It receives the owner-controlled front-page title instead of forcing a fixed sentence into every portfolio. Functions that reference it include normalizeSiteContent, renderSiteContent, addSiteSectionSearchEntries.

### **heroCopy**

heroCopy is the supporting paragraph in the hero section. It is one of the fields affected by rich text and plain text styling choices. Functions that reference it include normalizeSiteContent, renderSiteContent, addSiteSectionSearchEntries.

### **brandName**

brandName is the bold brand text in the header. renderProfile updates it to match the portfolio owner or site title. Functions that reference it include renderProfile.

### **brandSubtitle**

brandSubtitle is the small subtitle beside the brand. It usually communicates local builder/public portfolio context or a tagline. Functions that reference it include renderProfile.

### **brandIcon**

brandIcon is the image element for the OMB visual mark. Profile settings can change it when the builder is generalized for other users. Functions that reference it include renderProfile.

### **brandText**

brandText is the OMB engraving text in the header logo area. It carries the small trademark-like label. Functions that reference it include normalizeProfile, renderProfile.

### **headerAvatar**

headerAvatar is the clickable top-right profile image container. renderProfile reveals it only when a profile image exists. Functions that reference it include renderProfile.

### **headerAvatarImage**

headerAvatarImage is the actual img element inside the avatar link. It receives the saved profile picture URL. Functions that reference it include renderProfile.

### **heroImage**

heroImage is the main hero/background image element. Site content settings decide whether it is displayed. Functions that reference it include normalizeProfile, renderProfile.

### **contactBand**

contactBand is the final dark contact area. It also contains the AI assistant panel, so layout code must keep chat growth from distorting contact information. Functions that reference it include renderProfile.

### **contactTitle**

contactTitle is the heading for the contact panel. renderProfile updates it from the owner identity. Functions that reference it include renderProfile.

### **contactIntro**

contactIntro is the contact-section introduction text. It can be empty, plain text, or rich content depending on saved profile data. Functions that reference it include normalizeProfile, renderProfile, addSiteSectionSearchEntries.

### **contactDetails**

contactDetails is the list of direct contact methods such as email and phone. renderProfile builds clickable mailto and tel links inside it. Functions that reference it include renderProfile.

### **contactLinks**

contactLinks is the container for external profile links such as GitHub, LinkedIn, resume, or custom websites. Functions that reference it include renderProfile.

### **footerOwner**

footerOwner is the footer ownership text. It lets the footer follow the current profile owner rather than remaining generic. Functions that reference it include renderProfile.

### **searchWrap**

searchWrap is the parent around the search box. The search dropdown attaches near it so results appear where the user is typing. Functions that reference it include ensureSearchPanel.

### **categories**

categories is the runtime copy of project categories loaded from the catalog. It can include more than the original Analog, Digital, and Embedded groups. Functions that reference it include slugLabel, setActiveFilter,

renderCategoryFilters, addSiteSectionSearchEntries, assistantContextForQuestion, assistantFallbackResults, projectCard, renderProjects, openParsedSection, loadProjectCatalog.

### **funFacts**

funFacts stores simple fun fact strings. renderFunFacts turns at most a few lines into a near-top callout. Functions that reference it include renderFunFacts, loadProjectCatalog.

### **funFactsRich**

funFactsRich stores rich-format fun facts when the builder saves styled content instead of plain strings. Functions that reference it include renderFunFacts, loadProjectCatalog.

### **siteContent**

siteContent stores front-page text and site-level settings such as hero copy and titles. Functions that reference it include renderSiteContent, loadProjectCatalog.

### **siteContentRich**

siteContentRich stores rich versions of site-level text fields. Functions that reference it include renderSiteContent, loadProjectCatalog.

### **activeFilter**

activeFilter stores the currently selected project category filter. Functions that reference it include projectVisible, setActiveFilter, renderCategoryFilters, renderProjects.

### **searchPanel**

searchPanel stores the dropdown result panel created for search. Functions that reference it include ensureSearchPanel, renderSearchResults, goToSearchResult, restoreSectionRouteAfterCatalogLoad.

### **searchStatus**

searchStatus stores the status element inside the search dropdown. Functions that reference it include ensureSearchPanel, setSearchMessage.

### **searchLimitSelect**

searchLimitSelect stores the dropdown control that lets the visitor choose 5, 10, 15, or 20 results. Functions that reference it include ensureSearchPanel.

### **currentSearchResults**

currentSearchResults is the current ranked search-result array. It lets click handlers navigate results without recalculating the search. Functions that reference it include ensureSearchPanel, renderSearchResults, restoreSectionRouteAfterCatalogLoad.

### **searchableEntries**

searchableEntries is the full in-browser search index. It contains project titles, sections, files, captions, page sections, uploaded text, and electronics keyword expansions. Functions that reference it include rebuildSearchIndex, searchResultsFor.

### **searchResultLimit**

searchResultLimit stores how many dropdown results to display. Functions that reference it include ensureSearchPanel, renderSearchResults.

### **searchHighlightTimer**

searchHighlightTimer debounces highlight updates so typing does not repaint the page too aggressively. Functions that reference it include queueSearchHighlights.



### **searchHighlightName**

searchHighlightName is the CSS class applied to highlighted search matches. Functions that reference it include clearSearchHighlights, applySearchHighlights.

### **defaultSiteContent**

defaultSiteContent is the fallback front-page copy for a blank install. It keeps the site readable before the builder user customizes hero text and site messaging. Visible object fields include heroCopy, heroEyebrow, heroTitle. Those fields form the contract consumed by related functions. Functions that reference it include normalizeSiteContent.

## **Ai Configuration Or Conversation State**

### **aiAssistantForm**

aiAssistantForm is the chat form. Its submit handler starts answerAssistantQuestion. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include updateAssistantPanelGrowth, setAssistantBusy, restoreSectionRouteAfterCatalogLoad.

### **aiAssistantPanel**

aiAssistantPanel is the whole Ask My Portfolio panel. Growth and layout code use it to keep the assistant right-justified and scrollable. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include assistantDesktopDockedHeight, updateAssistantPanelGrowth.

### **aiAssistantInput**

aiAssistantInput is the question field. It is cleared/focused by chat handlers and intentionally does not keep long instructional text inside the typed value. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include setAssistantBusy, clearAssistantChat, answerAssistantQuestion, restoreSectionRouteAfterCatalogLoad.

### **aiAssistantLog**

aiAssistantLog is the scrollable message history. appendAssistantMessage and replaceAssistantMessage write user and assistant messages into it. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include updateAssistantPanelGrowth, appendAssistantMessage, scrollAssistantAnswerToStart, clearAssistantChat, restoreSectionRouteAfterCatalogLoad.

### **aiAssistantStatus**

aiAssistantStatus is the small status line used for ready/thinking/error states. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include setAssistantStatus.

### **assistantSourceCounter**

assistantSourceCounter creates stable ids for assistant source buttons. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include assistantSourcesMarkup, clearAssistantChat.

### **assistantSourceMap**

assistantSourceMap connects source button ids to source records so clicking a source opens the right place. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include assistantSourcesMarkup, clearAssistantChat, restoreSectionRouteAfterCatalogLoad.

## Appearance Configuration

### fieldStyles

fieldStyles stores plain-field font, size, and color choices. It lets front-page, profile, and contact fields render consistently with builder formatting. Functions that reference it include applyPlainFieldStyle, plainFieldStyleAttribute, loadProjectCatalog.

### legacyTemplateSkins

legacyTemplateSkins maps older template ids to the newer appearance-only template ids. It keeps older saved projects rendering after template behavior changed. Visible object fields include skin-light-blue, skin-clean-white, skin-deep-navy, skin-red-warm, skin-emerald-instrument, skin-amber-power, skin-violet-mixed, skin-graphite-asic, analog-opamp-topology, analog-power-charger, analog-filter-front-end, analog-sensor-interface, analog-mixed-signal-timing, digital-fpga-pipeline, digital-asic-block, digital-verification-suite, digital-interface-controller, digital-signal-processing. Those fields form the contract consumed by related functions. Functions that reference it include canonicalTemplateId.

## Runtime State Or Cache

### assistantChatHistory

assistantChatHistory stores recent user and assistant messages. It gives the chatbot conversation memory. It influences assistant behavior: conversation memory, source display, model prompts, backend routing, or context selection. Functions that reference it include assistantRemember, clearAssistantChat, answerAssistantQuestion.

## Compiler Or Language Configuration

### supportedCodeLanguages

supportedCodeLanguages is the public language list for code rendering and detection. It names C, C++, Verilog/SystemVerilog, Java, JavaScript, Python, HTML, and related aliases. It supports the compile workspace by describing language rules, executable locations, tool status, or compiler output. Functions that reference it include normalizeCodeLanguage, codeLanguageLabel.

## Code Walkthrough

This walkthrough is function-by-function. Each section now follows a textbook rhythm: first the idea of the function, then the syntax, then how to read that syntax, then the implementation and the local objects that make the implementation work. Line numbers are deliberately avoided because the source will keep changing.

### Public runtime utility

The functions in this group all support public runtime utility. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

#### function normalize(value)

To understand normalize, start with the problem it owns. normalize standardizes value data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalize(value) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `normalize` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `normalize`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `normalize`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **function sourceLooksCpp(source = "")**

Begin with the role of `sourceLooksCpp`. `sourceLooksCpp` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sourceLooksCpp(source = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `source`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `source` is source text supplied by the builder, project catalog, GitHub, or an uploaded file. The function inspects, formats, saves, compiles, highlights, or extracts meaning from this text. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `sourceLooksCpp` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `sourceLooksCpp`, the specific rule it owns would disappear from the named workflow. Callers would either skip

that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `sourceLooksCpp` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**text:** text stores a computed value for temporary calculation. It is initialized from `String(source || "")`. Within the function it is passed into test; assigned or updated later in the function.

#### **function sourceLooksC(source = "")**

Begin with the role of `sourceLooksC`. `sourceLooksC` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sourceLooksC(source = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `source`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `source` is source text supplied by the builder, project catalog, GitHub, or an uploaded file. The function inspects, formats, saves, compiles, highlights, or extracts meaning from this text. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `sourceLooksC` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `sourceLooksC`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `sourceLooksC` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**text:** text stores a computed value for temporary calculation. It is initialized from `String(source || "")`. Within the function it is passed into test; assigned or updated later in the function.

#### **function textBlocksFromPlainText(text)**

Begin with the role of `textBlocksFromPlainText`. `textBlocksFromPlainText` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function textBlocksFromPlainText(text) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `text`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `text` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `textBlocksFromPlainText` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `textBlocksFromPlainText`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `textBlocksFromPlainText`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function unwrapFormula(value)**

The best entry point for `unwrapFormula` is its responsibility in the surrounding workflow. `unwrapFormula` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function unwrapFormula(value) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `unwrapFormula` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `unwrapFormula`, the specific rule it owns would disappear from the named workflow. Callers

would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `unwrapFormula` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**text:** text stores a computed value for temporary calculation. It is initialized from `String(value || "").trim()`. Within the function it is returned to the caller; accesses properties/methods match; assigned or updated later in the function.

**wrappers:** wrappers stores a list for temporary calculation. It is initialized from `[]`. Within the function it is assigned or updated later in the function.

**destructured [pattern, group]:** destructured [pattern, group] starts without an immediate value. In this file that usually means it is filled by a loop, branch, callback, or later assignment inside the same function. It unpacks API handles from another object, giving the rest of the file short direct names for those APIs.

**match:** match stores a computed value for temporary calculation. It is initialized from `text.match(pattern)`. Within the function it is returned to the caller; assigned or updated later in the function.

### **function looksLikeBareWebAddress(value = "")**

The best entry point for `looksLikeBareWebAddress` is its responsibility in the surrounding workflow. `looksLikeBareWebAddress` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function looksLikeBareWebAddress(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `looksLikeBareWebAddress` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `looksLikeBareWebAddress`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `looksLikeBareWebAddress` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(value || "").trim()`.

Within the function it is accesses properties/methods `split`; passed into `test`; assigned or updated later in the function.

**host:** `host` stores a computed value for temporary calculation. It is initialized from `clean.split(/[/?#]/)`

`[0].toLowerCase()`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function displayTitle(value, fallback = "Untitled item")**

Begin with the role of `displayTitle`. `displayTitle` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function displayTitle(value, fallback = "Untitled item") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`, `fallback`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `fallback` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `displayTitle` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `displayTitle`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `displayTitle` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `String(value || fallback).trim()`.

Within the function it is assigned or updated later in the function.

**replacements:** `replacements` stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is returned to the caller; assigned or updated later in the function.



### **function slugLabel(value)**

The best entry point for slugLabel is its responsibility in the surrounding workflow. slugLabel belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function slugLabel(value) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: value. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. value is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means slugLabel performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without slugLabel, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of slugLabel is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**category:** category stores a function or callback value for portfolio data state. It is initialized from categories.find((item) => item.id === value). It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is returned to the caller; accesses properties/methods label; assigned or updated later in the function.

### **function parsedItemTerms(item)**

To understand parsedItemTerms, start with the problem it owns. parsedItemTerms interprets dItemTerms text and returns structured meaning. Parsing functions are needed because the rest of the app should not repeatedly scan raw strings when it can work with a stable object, array, or normalized value.

The syntax of the function is:

```
function parsedItemTerms(item) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: item. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.



Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `parsedItemTerms` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `parsedItemTerms`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `parsedItemTerms`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **function pillList(items, className = "")**

To understand `pillList`, start with the problem it owns. `pillList` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function pillList(items, className = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `items`, `className`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `className` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `pillList` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `pillList`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `pillList` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**label:** `label` stores a computed value for temporary calculation. It is initialized from `typeof item === "string" ? item : item.name || item.title || item.label || ""`. Within the function it is returned to the caller; assigned or updated later in the function.

### **function evidenceList(items, renderItem, emptyMessage)**

The best entry point for `evidenceList` is its responsibility in the surrounding workflow. `evidenceList` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function evidenceList(items, renderItem, emptyMessage) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `items`, `renderItem`, `emptyMessage`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `renderItem` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `emptyMessage` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `evidenceList` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `evidenceList`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `evidenceList`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function detailBlock(title, className, content)**

Read `detailBlock` first as a named idea, not as syntax. `detailBlock` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function detailBlock(title, className, content) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `title`, `className`, `content`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `title` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `className` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `content` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `detailBlock` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `detailBlock`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `detailBlock`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function mediaGrid(items)**

Read `mediaGrid` first as a named idea, not as syntax. `mediaGrid` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function mediaGrid(items) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `items`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `escapeHtml`, `normalizeLinkTarget`, `isWebsiteLinkItem`, `linkAttributes`. Those calls show that `mediaGrid` is not isolated; it is one piece of a larger

workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `mediaGrid`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `mediaGrid`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function itemLabel(item = {}, fallback = "Download file")**

The best entry point for `itemLabel` is its responsibility in the surrounding workflow. `itemLabel` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function itemLabel(item = {}, fallback = "Download file") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `item`, `fallback`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `fallback` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `itemUrl`, `fileNameFromUrl`. Those calls show that `itemLabel` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `itemLabel`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `itemLabel` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**url:** `url` stores a computed value for network or routing value. It is initialized from `itemUrl(item)`. Within the function it is passed into `fileNameFromUrl`; assigned or updated later in the function.

### **function parsedChildCards(children, projectId, sectionIndex, basePath = [])**

Begin with the role of `parsedChildCards`. `parsedChildCards` interprets `dChildCards` text and returns structured meaning. Parsing functions are needed because the rest of the app should not repeatedly scan raw strings when it can work with a stable object, array, or normalized value.

The syntax of the function is:

```
function parsedChildCards(children, projectId, sectionIndex, basePath = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `children`, `projectId`, `sectionIndex`, `basePath`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `children` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `basePath` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `nodeHasRenderableContent`, `nodeIsOverview`, `parsedNodeCard`. Those calls show that `parsedChildCards` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `parsedChildCards`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `parsedChildCards` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**visibleChildren:** `visibleChildren` stores a function or callback value for temporary calculation. It is initialized from `children.filter((child) => nodeHasRenderableContent(child) && !child.url && !nodeIsOverview(child))`. Within the function it accesses properties/methods `length`; assigned or updated later in the function.

### **function escapeRegExp(value = "")**

Begin with the role of `escapeRegExp`. `escapeRegExp` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function escapeRegExp(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `escapeRegExp` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `escapeRegExp`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `escapeRegExp`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function scrollToDomTarget(id)**

To understand `scrollToDomTarget`, start with the problem it owns. `scrollToDomTarget` belongs to public runtime utility. This helper supports the public site runtime by normalizing values or preparing data for another renderer. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function scrollToDomTarget(id) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `id`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `id` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser route/history. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: updates browser history so Back, Forward, and refresh can follow the same window state.

Next, trace the helper calls that surround the function. It works with `closeSectionDialogForRoute`, `queueSearchHighlights`. Those calls show that `scrollToDomTarget` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `scrollToDomTarget`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `scrollToDomTarget` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**target:** target stores a DOM element reference. It is initialized from `id ? document.getElementById(id) : null`. Within the function it is accesses properties/methods `scrollIntoView`; assigned or updated later in the function.

## Profile and site content

The functions in this group all support profile and site content. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### function normalizeFunFacts(value)

The best entry point for `normalizeFunFacts` is its responsibility in the surrounding workflow. `normalizeFunFacts` standardizes FunFacts data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeFunFacts(value) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `normalizeFunFacts` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `normalizeFunFacts`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

The easiest way to follow the body of `normalizeFunFacts` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**items:** items stores a computed value for temporary calculation. It is initialized from `Array.isArray(value) ? value : String(value || "").split(/\r?\n/)`. Within the function it is returned to the caller; assigned or updated later in the function.



### **function normalizeSiteContent(content = {})**

To understand `normalizeSiteContent`, start with the problem it owns. `normalizeSiteContent` standardizes `SiteContent` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeSiteContent(content = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `content`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `content` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `normalizeSiteContent` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeSiteContent`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

Inside `normalizeSiteContent`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function normalizeProfile(profileValue = {})**

Read `normalizeProfile` first as a named idea, not as syntax. `normalizeProfile` standardizes `Profile` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeProfile(profileValue = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `profileValue`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `profileValue` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.



With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `normalizeProfile` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `normalizeProfile`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

Inside `normalizeProfile`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function mailComposeLink(email = "")**

The best entry point for `mailComposeLink` is its responsibility in the surrounding workflow. `mailComposeLink` belongs to profile and site content. This function renders or normalizes owner identity, front-page text, contact links, fun facts, or custom main-page sections. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function mailComposeLink(email = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `email`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `email` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `mailComposeLink` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `mailComposeLink`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

The easiest way to follow the body of `mailComposeLink` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(email || "").trim()`.

Within the function it is returned to the caller; passed into `encodeURIComponent`; assigned or updated later in the function.

### **function phoneLink(phone = "")**

Begin with the role of phoneLink. phoneLink belongs to profile and site content. This function renders or normalizes owner identity, front-page text, contact links, fun facts, or custom main-page sections. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function phoneLink(phone = "") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: phone. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. phone is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means phoneLink performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without phoneLink, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

The local state inside phoneLink explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** clean stores a computed value for temporary calculation. It is initialized from `String(phone || "").trim()`. Within the function it is returned to the caller; accesses properties/methods replace; assigned or updated later in the function.

### **function looksLikeWebOrContactText(value = "")**

To understand looksLikeWebOrContactText, start with the problem it owns. looksLikeWebOrContactText belongs to profile and site content. This function renders or normalizes owner identity, front-page text, contact links, fun facts, or custom main-page sections. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function looksLikeWebOrContactText(value = "") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: value. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. value is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and

generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `looksLikeBareWebAddress`. Those calls show that `looksLikeWebOrContactText` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `looksLikeWebOrContactText`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

After the main flow, the working values inside `looksLikeWebOrContactText` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(value || "").trim()`.

Within the function it is returned to the caller; accesses properties/methods `split`; passed into `test`; assigned or updated later in the function.

**function `isWebsiteLinkItem(item = {}, target = "")`**

Begin with the role of `isWebsiteLinkItem`. `isWebsiteLinkItem` is a decision predicate. It answers one yes/no question so the larger workflow can branch clearly instead of embedding the same condition in several places.

The syntax of the function is:

```
function isWebsiteLinkItem(item = {}, target = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `item`, `target`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `target` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `isWebsiteLinkItem` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `isWebsiteLinkItem`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

The local state inside `isWebsiteLinkItem` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(target || "").trim()`. Within the function it is passed into `test`; assigned or updated later in the function.

**typeText:** `typeText` stores a list for temporary calculation. It is initialized from `[item.kind, item.type, item.status, item.meta, item.label, item.title]`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function inferResponsiveProfile(project)**

The best entry point for `inferResponsiveProfile` is its responsibility in the surrounding workflow. `inferResponsiveProfile` makes an educated classification from incomplete evidence. It looks at names, extensions, source text, repository state, or runtime state and returns the app's best decision so the next operation can choose the correct path.

The syntax of the function is:

```
function inferResponsiveProfile(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Its connection to the rest of the file matters as much as its local code. It works with `sectionHasRenderableContent`, `responsiveNodeCount`, `responsiveNodeHasMedia`, `hasPublicTemplate`. Those calls show that `inferResponsiveProfile` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `inferResponsiveProfile`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

The easiest way to follow the body of `inferResponsiveProfile` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**sections:** `sections` stores a computed value for portfolio data state. It is initialized from `project?.portfolioView?.sections || []`. Within the function it is iterated or transformed as a collection; accesses properties/methods `filter`; assigned or updated later in the function.

**visibleSections:** visibleSections stores a function or callback value for portfolio data state. It is initialized from `sections.filter((section) => section?.id !== "brief" && sectionHasRenderableContent(section))`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `reduce`, `some`; assigned or updated later in the function.

**itemCount:** itemCount stores a computed value for temporary calculation. It is initialized from `visibleSections.reduce()`. Within the function it is assigned or updated later in the function.

**maxDepth:** maxDepth stores a computed value for temporary calculation. It is initialized from `visibleSections.reduce()`. Within the function it is assigned or updated later in the function.

**hasMedia:** hasMedia stores a function or callback value for temporary calculation. It is initialized from `visibleSections.some((section) => responsiveNodeHasMedia(section))`. Within the function it is assigned or updated later in the function.

**density:** density stores a computed value for temporary calculation. It is initialized from `itemCount >= 18 || visibleSections.length >= 6 || maxDepth >= 4`. Within the function it is assigned or updated later in the function.

### **function projectResponsiveProfile(project)**

Read `projectResponsiveProfile` first as a named idea, not as syntax. `projectResponsiveProfile` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectResponsiveProfile(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

The function also has to be read through the helpers it calls. It works with `inferResponsiveProfile`. Those calls show that `projectResponsiveProfile` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `projectResponsiveProfile`, owner-specific content would remain generic or malformed. Contact links, profile text, fun facts, and custom sections would not render with the intended public shape.

The local objects and variables inside `projectResponsiveProfile` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**savedProfile:** savedProfile stores a computed value for portfolio data state. It is initialized from `project?.portfolioView?.responsive`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

## Rich content rendering

The functions in this group all support rich content rendering. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### function renderFunFacts()

Read `renderFunFacts` first as a named idea, not as syntax. `renderFunFacts` turns saved data for FunFacts into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderFunFacts() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

The function also has to be read through the helpers it calls. It works with `normalizeFunFacts`, `renderRichContent`, `renderInlineMath`. Those calls show that `renderFunFacts` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `renderFunFacts`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local objects and variables inside `renderFunFacts` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**facts:** facts stores a computed value for temporary calculation. It is initialized from `normalizeFunFacts(funFacts)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `join`, `length`, `map`; assigned or updated later in the function.

**hasRichFacts:** hasRichFacts stores a computed value for temporary calculation. It is initialized from `Boolean(funFactsRich?.blocks?.length)`. Within the function it is assigned or updated later in the function.

### **function escapeHtml(value)**

Read `escapeHtml` first as a named idea, not as syntax. `escapeHtml` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function escapeHtml(value) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `escapeHtml` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `escapeHtml`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `escapeHtml`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function normalizeCodeLanguage(value = "")**

Begin with the role of `normalizeCodeLanguage`. `normalizeCodeLanguage` standardizes `CodeLanguage` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCodeLanguage(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.



The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `normalizeCodeLanguage` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `normalizeCodeLanguage`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local state inside `normalizeCodeLanguage` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(value || "").trim().toLowerCase().replace(/[_-]+/g, " ")`. Within the function it is passed into `includes`; assigned or updated later in the function.

**match:** `match` stores a function or callback value for temporary calculation. It is initialized from `supportedCodeLanguages.find((language) => ...)`. Within the function it is returned to the caller; assigned or updated later in the function.

**function `codeLanguageLabel(value = "")`**

The best entry point for `codeLanguageLabel` is its responsibility in the surrounding workflow. `codeLanguageLabel` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function codeLanguageLabel(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `normalizeCodeLanguage`. Those calls show that `codeLanguageLabel` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `codeLanguageLabel`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `codeLanguageLabel`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **`function normalizeCodePasteMode(value = "")`**

Read `normalizeCodePasteMode` first as a named idea, not as syntax. `normalizeCodePasteMode` standardizes `CodePasteMode` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCodePasteMode(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `normalizeCodePasteMode` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `normalizeCodePasteMode`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `normalizeCodePasteMode`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **`function normalizeCodeText(value = "", pasteMode = "source")`**

To understand `normalizeCodeText`, start with the problem it owns. `normalizeCodeText` standardizes `CodeText` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCodeText(value = "", pasteMode = "source") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`, `pasteMode`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `pasteMode` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalizeCodePasteMode`. Those calls show that `normalizeCodeText` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeCodeText`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside `normalizeCodeText` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**normalized:** `normalized` stores a computed value for temporary calculation. It is initialized from `String(value || "").replace(/r\n?/g, "\n").replace(/u00a0/g, " ")`. Within the function it is returned to the caller; accesses properties/methods `replace`; assigned or updated later in the function.

#### **function detectCodeLanguage(value = "")**

Begin with the role of `detectCodeLanguage`. `detectCodeLanguage` makes an educated classification from incomplete evidence. It looks at names, extensions, source text, repository state, or runtime state and returns the app's best decision so the next operation can choose the correct path.

The syntax of the function is:

```
function detectCodeLanguage(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `sourceLooksCpp`, `sourceLooksC`. Those calls show that `detectCodeLanguage` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `detectCodeLanguage`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local state inside `detectCodeLanguage` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**code:** code stores a computed value for temporary calculation. It is initialized from `String(value || "")`. Within the function it is passed into `sourceLooksC`, `sourceLooksCpp`, `test`; assigned or updated later in the function.

**function tokenizedCodeHtml(code = "", language = "javascript")**

To understand `tokenizedCodeHtml`, start with the problem it owns. `tokenizedCodeHtml` is the lightweight syntax highlighter used by the public site. It escapes the code first, then wraps comments, strings, numbers, keywords, preprocessor directives, and other recognized tokens in spans with CSS classes. It supports the major languages that the builder exposes, including C, C++, SystemVerilog, LTspice-oriented text, Java, JavaScript, Python, and HTML-like code. The function exists so code added to a project reads like code rather than plain paragraph text. Without it, source examples in the portfolio would be visually flat, harder to scan, and less convincing to technical recruiters.

The syntax of the function is:

```
function tokenizedCodeHtml(code = "", language = "javascript") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `code`, `language`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `code` is source text supplied by the builder, project catalog, GitHub, or an uploaded file. The function inspects, formats, saves, compiles, highlights, or extracts meaning from this text. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `language` names the programming or hardware-description language. The function uses it to choose file extensions, highlighting rules, compilers, simulators, labels, or output behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `escapeHtml`, `normalizeCodeLanguage`. Those calls show that `tokenizedCodeHtml` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `tokenizedCodeHtml`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside `tokenizedCodeHtml` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**html:** html stores a computed value for temporary calculation. It is initialized from `escapeHtml(code)`. Within the function it is returned to the caller; accesses properties/methods replace; assigned or updated later in the function.

**tokens:** tokens stores a list for temporary calculation. It is initialized from []. Within the function it is mutated as a collection or DOM container; accesses properties/methods length, push; assigned or updated later in the function.

**protect:** protect stores a function or callback value for temporary calculation. It is initialized from (pattern, className) => {}. Within the function it is assigned or updated later in the function.

**token:** token stores a string value for temporary calculation. It is initialized from `@@CODE\_TOKEN\_\${tokens.length}@@`. Within the function it is returned to the caller; passed into protect; assigned or updated later in the function.

**commentPattern:** commentPattern stores a list for temporary calculation. It is initialized from ["python", "Itspice"].includes(normalizeCodeLanguage(language)). Within the function it is passed into protect; assigned or updated later in the function.

**keywordMap:** keywordMap stores a structured object for temporary calculation. It is initialized from {}. Within the function it is accesses properties/methods javascript; assigned or updated later in the function.

**keywords:** keywords stores a computed value for lookup collection. It is initialized from keywordMap[normalizeCodeLanguage(language)] || keywordMap.javascript. Within the function it is passed into replace; assigned or updated later in the function.

### function renderRichCodeBlock(block = {})

Begin with the role of renderRichCodeBlock. renderRichCodeBlock turns saved data for RichCodeBlock into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that styles.css expects.

The syntax of the function is:

```
function renderRichCodeBlock(block = {}) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: block. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. block is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with normalizeCodeLanguage, detectCodeLanguage, normalizeCodeText, escapeHtml, codeLanguageLabel, tokenizedCodeHtml. Those calls show that renderRichCodeBlock is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without renderRichCodeBlock, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local state inside `renderRichCodeBlock` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**language:** `language` stores a computed value for compiler or language configuration. It is initialized from `normalizeCodeLanguage(block.language || detectCodeLanguage(block.code || ""))`. It supports the compile workspace by describing language rules, executable locations, tool status, or compiler output. Within the function it is passed into `escapeHtml`, `normalizeCodeLanguage`, `tokenizedCodeHtml`; assigned or updated later in the function.

**code:** `code` stores a computed value for temporary calculation. It is initialized from `normalizeCodeText(block.code || "", block.pasteMode || "source")`. Within the function it is passed into `normalizeCodeLanguage`, `normalizeCodeText`, `tokenizedCodeHtml`; assigned or updated later in the function.

### **function renderPlainMultiline(value = "")**

The best entry point for `renderPlainMultiline` is its responsibility in the surrounding workflow. `renderPlainMultiline` turns saved data for `PlainMultiline` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderPlainMultiline(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `escapeHtml`. Those calls show that `renderPlainMultiline` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `renderPlainMultiline`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `renderPlainMultiline`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderSiteContent()**

Begin with the role of `renderSiteContent`. `renderSiteContent` turns saved data for `SiteContent` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderSiteContent() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `normalizeSiteContent`, `renderRichFieldContent`, `applyPlainFieldStyle`. Those calls show that `renderSiteContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `renderSiteContent`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local state inside `renderSiteContent` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**content:** content stores a computed value for temporary calculation. It is initialized from `normalizeSiteContent(siteContent || {})`. Within the function it is accesses properties/methods `heroCopy`, `heroEyebrow`, `heroTitle`; passed into `renderRichFieldContent`; assigned or updated later in the function.

### **function renderProfile()**

To understand `renderProfile`, start with the problem it owns. `renderProfile` applies owner identity to the public page. It updates the displayed name, tagline, profile photo, brand/title text, email link, phone link, GitHub/resume links, contact cards, resume viewer, metadata hints, and social/contact links. It is the bridge between saved profile data and the visible recruiter-facing contact section. The function is important because the same website shell can serve Maurice's portfolio or a fresh user's portfolio. Profile data cannot be hardcoded if the builder is intended to be reusable. Without `renderProfile`, the page would either remain generic or require manual HTML editing for every owner.

The syntax of the function is:

```
function renderProfile() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.



Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

Next, trace the helper calls that surround the function. It works with `normalizeProfile`, `mailComposeLink`, `phoneLink`, `normalizeLinkTarget`, `renderRichFieldContent`, `applyPlainFieldStyle`, `plainFieldStyleAttribute`. Those calls show that `renderProfile` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `renderProfile`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside `renderProfile` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**current:** `current` stores a computed value for temporary calculation. It is initialized from `normalizeProfile(profile || {})`. Within the function it is accesses properties/methods `brandImage`, `brandText`, `contactIntro`, `displayName`, `email`, `githubUrl`, `heroImage`, `linkedinUrl`; passed into `mailComposeLink`, `normalizeLinkTarget`, `phoneLink`, `renderRichFieldContent`; assigned or updated later in the function.

**displayName:** `displayName` stores a computed value for temporary calculation. It is initialized from `current.displayName || "Portfolio"`. Within the function it is accesses properties/methods `slice`; passed into `renderRichFieldContent`, `setAttribute`; assigned or updated later in the function.

**label:** `label` stores a computed value for temporary calculation. It is initialized from `current.portfolioLabel || "Portfolio"`. Within the function it is passed into `applyPlainFieldStyle`, `renderRichFieldContent`, `setAttribute`; assigned or updated later in the function.

**emailLink:** `emailLink` stores a computed value for temporary calculation. It is initialized from `mailComposeLink(current.email)`. Within the function it is assigned or updated later in the function.

**callLink:** `callLink` stores a computed value for temporary calculation. It is initialized from `phoneLink(current.phone)`. Within the function it is assigned or updated later in the function.

**links:** `links` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; passed into `renderRichFieldContent`; assigned or updated later in the function.

**heroGithubLink:** `heroGithubLink` stores a list for Git publishing and authorization state. It is initialized from `[...document.querySelectorAll(".hero-actions a")]`. It affects publishing, where mistakes can change the public website or expose the wrong files. Within the function it is accesses properties/methods `hidden`, `href`; assigned or updated later in the function.

**resumeObject:** `resumeObject` stores a DOM element reference. It is initialized from `resumeSection.querySelector("object")`. Within the function it is accesses properties/methods `data`; assigned or updated later in the function.

**fallbackLink:** `fallbackLink` stores a DOM element reference. It is initialized from `resumeSection.querySelector("object a")`. Within the function it is accesses properties/methods `href`; assigned or updated later in the function.

**styleId:** styleId stores a computed value for appearance configuration. It is initialized from link.label === "GitHub". Within the function it is passed into plainFieldStyleAttribute; assigned or updated later in the function.

### function renderInlineMath(text)

To understand renderInlineMath, start with the problem it owns. renderInlineMath turns saved data for InlineMath into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that styles.css expects.

The syntax of the function is:

```
function renderInlineMath(text) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: text. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. text is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Next, trace the helper calls that surround the function. It works with escapeHtml, normalizeLinkTarget. Those calls show that renderInlineMath is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without renderInlineMath, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside renderInlineMath reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**escaped:** escaped stores a computed value for temporary calculation. It is initialized from escapeHtml(text). Within the function it is accesses properties/methods replace; assigned or updated later in the function.

**withMath:** withMath stores a computed value for temporary calculation. It is initialized from escaped.replace(/\\$([^\\$]+)\\$/g, '<span class="rich-inline-formula">\$1</span>'). Within the function it is returned to the caller; accesses properties/methods replace; assigned or updated later in the function.

**trailing:** trailing stores a computed value for temporary calculation. It is initialized from match.match(/I),... Within the function it is accesses properties/methods length; passed into slice; assigned or updated later in the function.

**clean:** clean stores a computed value for temporary calculation. It is initialized from trailing ? match.slice(0, -trailing.length) : match. Within the function it is passed into normalizeLinkTarget; assigned or updated later in the function.

**href:** href stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(clean, { assumeWeb: true })`. Within the function it is assigned or updated later in the function.

### **function renderMultilineInlineText(text = "")**

The best entry point for `renderMultilineInlineText` is its responsibility in the surrounding workflow.

`renderMultilineInlineText` turns saved data for `MultilineInlineText` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderMultilineInlineText(text = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `text`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `text` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `renderInlineMath`. Those calls show that `renderMultilineInlineText` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `renderMultilineInlineText`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `renderMultilineInlineText`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function sanitizeRichInlineHtml(value = "")**

To understand `sanitizeRichInlineHtml`, start with the problem it owns. `sanitizeRichInlineHtml` is the safety filter for rich inline HTML. Builder editors may store spans, links, bold text, colors, and other inline markup, but the public website should not blindly inject arbitrary HTML. This function removes unsafe tags and attributes while preserving the formatting the portfolio actually needs. The function prevents a content field from becoming a script injection surface. It works with `renderRichContent` and `renderRichFieldContent`, which need formatted HTML but should only receive controlled markup.

The syntax of the function is:

```
function sanitizeRichInlineHtml(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: creates DOM elements instead of relying only on static HTML; writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

Next, trace the helper calls that surround the function. It works with `normalizeLinkTarget`, `cleanFontFamily`, `normalizeFontPx`, `normalizeTextColor`, `linkifyRichTextNodes`. Those calls show that `sanitizeRichInlineHtml` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `sanitizeRichInlineHtml`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside `sanitizeRichInlineHtml` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**template:** `template` stores a DOM element reference. It is initialized from `document.createElement("template")`. Within the function it is returned to the caller; accesses properties/methods `content`, `innerHTML`; passed into `createElement`, `linkifyRichTextNodes`; assigned or updated later in the function.

**allowedTags:** `allowedTags` stores a constructed object for lookup collection. It is initialized from `new Set(["A", "B", "BR", "EM", "I", "SPAN", "STRONG", "U"])`. Within the function it is accesses properties/methods `has`; assigned or updated later in the function.

**href:** `href` stores a computed value for temporary calculation. It is initialized from `node.tagName === "A" ? String(node.getAttribute("href") || "").trim() : ""`. Within the function it is passed into `String`, `normalizeLinkTarget`, `setAttribute`; assigned or updated later in the function.

**normalizedHref:** `normalizedHref` stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(href, { assumeWeb: true })`. Within the function it is passed into `test`; assigned or updated later in the function.

**safeHref:** `safeHref` stores a computed value for temporary calculation. It is initialized from `/(https?:|mailto:|tel:| #|V)/i.test(normalizedHref) ? normalizedHref : ""`. Within the function it is passed into `setAttribute`; assigned or updated later in the function.

**fontFamily:** `fontFamily` stores a computed value for temporary calculation. It is initialized from `cleanFontFamily(node.style.fontFamily || "")`. Within the function it is passed into `cleanFontFamily`; assigned or updated later in the function.

**fontPx:** `fontPx` stores a computed value for temporary calculation. It is initialized from `normalizeFontPx(parseFloat(node.style.fontSize || ""))`. Within the function it is assigned or updated later in the function.

**color:** `color` stores a computed value for appearance configuration. It is initialized from `normalizeTextColor(node.style.color || "")`. Within the function it is passed into `normalizeTextColor`; assigned or updated later in the function.

**rawFontWeight:** rawFontWeight stores a computed value for temporary calculation. It is initialized from `node.style.fontWeight || ""`. Within the function it is passed into test; assigned or updated later in the function.

**fontWeight:** fontWeight stores a computed value for temporary calculation. It is initialized from `/(bold|[6-9]00)/i.test(rawFontWeight)`. Within the function it is assigned or updated later in the function.

**fontStyle:** fontStyle stores a computed value for appearance configuration. It is initialized from `node.style.fontStyle === "italic" ? "italic" : ""`. Within the function it is assigned or updated later in the function.

**textDecoration:** textDecoration stores a computed value for temporary calculation. It is initialized from `node.style.textDecoration.includes("underline") ? "underline" : ""`. Within the function it is accesses properties/methods includes; assigned or updated later in the function.

### **function richTextStyle(block = {})**

Begin with the role of richTextStyle. richTextStyle belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function richTextStyle(block = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `block`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `block` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `cleanFontFamily`, `normalizeFontPx`, `normalizeTextColor`, `escapeHtml`. Those calls show that richTextStyle is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without richTextStyle, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local state inside richTextStyle explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**styles:** styles stores a list for appearance configuration. It is initialized from `[]`. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods `join`, `length`, `push`; passed into `escapeHtml`; assigned or updated later in the function.

**fontFamily:** fontFamily stores a computed value for temporary calculation. It is initialized from `cleanFontFamily(block.fontFamily || "Arial") || "Arial"`. Within the function it is passed into `cleanFontFamily`, `push`; assigned or updated later in the function.

**fontPx:** fontPx stores a computed value for temporary calculation. It is initialized from `normalizeFontPx(block.fontPx)`. Within the function it is passed into `normalizeFontPx`, push; assigned or updated later in the function.

**color:** color stores a computed value for appearance configuration. It is initialized from `normalizeTextColor(block.color || "")`. Within the function it is passed into `normalizeTextColor`, push; assigned or updated later in the function.

### **function richImageCropStyle(block = {})**

Read `richImageCropStyle` first as a named idea, not as syntax. `richImageCropStyle` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function richImageCropStyle(block = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `block`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `block` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `normalizeCropAspect`, `normalizeCropZoom`, `normalizeCropPosition`, `escapeHtml`. Those calls show that `richImageCropStyle` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `richImageCropStyle`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local objects and variables inside `richImageCropStyle` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**aspect:** aspect stores a computed value for temporary calculation. It is initialized from `normalizeCropAspect(block.cropAspect)`. Within the function it is passed into `push`; assigned or updated later in the function.

**styles:** styles stores a list for appearance configuration. It is initialized from `[]`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `join`, `push`; passed into `escapeHtml`; assigned or updated later in the function.

### **function richImageDownloadLink(block = {})**

The best entry point for `richImageDownloadLink` is its responsibility in the surrounding workflow. `richImageDownloadLink` belongs to rich content rendering. This function turns saved builder content into public

HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function richImageDownloadLink(block = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `block`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `block` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `escapeHtml`, `cleanRichImageTitle`, `normalizeLinkTarget`, `isWebsiteLinkItem`, `linkAttributes`, `downloadAttribute`. Those calls show that `richImageDownloadLink` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `richImageDownloadLink`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The easiest way to follow the body of `richImageDownloadLink` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**label:** `label` stores a computed value for temporary calculation. It is initialized from `escapeHtml(cleanRichImageTitle(block) || block.caption || "Image file")`. Within the function it is assigned or updated later in the function.

**url:** `url` stores a computed value for network or routing value. It is initialized from `block.url || "#"`. Within the function it is passed into `normalizeLinkTarget`; assigned or updated later in the function.

**target:** `target` stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(url, { assumeWeb: isWebsiteLinkItem(block, url) })`. Within the function it is passed into `downloadAttribute`, `escapeHtml`, `linkAttributes`; assigned or updated later in the function.

**function `cleanRichImageTitle(block = {})`**

Read `cleanRichImageTitle` first as a named idea, not as syntax. `cleanRichImageTitle` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function cleanRichImageTitle(block = {}) { ... }
```



There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `block`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `block` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `cleanRichImageTitle` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `cleanRichImageTitle`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local objects and variables inside `cleanRichImageTitle` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `String(block.title || "").trim()`. Within the function it is passed into `String, test`; assigned or updated later in the function.

#### **function renderRichContent(rich, fallbackText = "")**

To understand `renderRichContent`, start with the problem it owns. `renderRichContent` is the renderer for builder-created rich content. It accepts blocks representing paragraphs, images, formulas, files, links, code blocks, captions, and styled text, then converts those blocks into safe public HTML. It decides whether an image should appear inline, whether a file should be downloadable, how captions attach, and how text styles survive into the website. This function matters because the builder is not just saving plain text. It supports formatted explanations, pasted images, formulas, code, and files. Without `renderRichContent`, much of what the builder collects would be flattened or lost when the portfolio is published.

The syntax of the function is:

```
function renderRichContent(rich, fallbackText = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rich`, `fallbackText`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `rich` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `fallbackText` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Next, trace the helper calls that surround the function. It works with `textBlocksFromPlainText`, `richImageDownloadLink`, `cleanRichImageTitle`, `normalizeLinkTarget`, `normalizeCropAspect`, `richImageCropStyle`, `escapeHtml`, `unwrapFormula`. Those calls show that `renderRichContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `renderRichContent`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside `renderRichContent` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**blocks:** `blocks` stores a computed value for temporary calculation. It is initialized from `rich?.blocks?.length ? rich.blocks : textBlocksFromPlainText(fallbackText)`. Within the function it is iterated or transformed as a collection; accesses properties/methods map; assigned or updated later in the function.

**align:** `align` stores a list for temporary calculation. It is initialized from `["left", "center", "right"].includes(block.align) ? block.align : "left"`. Within the function it is passed into `includes`; assigned or updated later in the function.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `cleanRichImageTitle(block)`. Within the function it is passed into `escapeHtml`; assigned or updated later in the function.

**imageSrc:** `imageSrc` stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(block.url, { assumeWeb: true })`. Within the function it is passed into `escapeHtml`; assigned or updated later in the function.

**formula:** `formula` stores a computed value for temporary calculation. It is initialized from `unwrapFormula(block.formula)`. Within the function it is passed into `escapeHtml`, `renderInlineMath`, `unwrapFormula`; assigned or updated later in the function.

**size:** `size` stores a list for temporary calculation. It is initialized from `["small", "normal", "large"].includes(block.fontSize) ? block.fontSize : "normal"`. Within the function it is assigned or updated later in the function.

**content:** `content` stores a computed value for temporary calculation. It is initialized from `block.html`. Within the function it is assigned or updated later in the function.

### **function renderRichFieldContent(rich, fallbackText = "")**

To understand `renderRichFieldContent`, start with the problem it owns. `renderRichFieldContent` turns saved data for `RichFieldContent` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderRichFieldContent(rich, fallbackText = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rich`, `fallbackText`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `rich` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `fallbackText` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Next, trace the helper calls that surround the function. It works with `textBlocksFromPlainText`, `sanitizeRichInlineHtml`, `renderInlineMath`, `richTextStyle`. Those calls show that `renderRichFieldContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `renderRichFieldContent`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

After the main flow, the working values inside `renderRichFieldContent` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**blocks:** `blocks` stores a computed value for temporary calculation. It is initialized from `rich?.blocks?.length ? rich.blocks : textBlocksFromPlainText(fallbackText)`. Within the function it is returned to the caller; iterated or transformed as a collection; accesses properties/methods map; assigned or updated later in the function.

**align:** `align` stores a list for temporary calculation. It is initialized from `["left", "center", "right"].includes(block.align) ? block.align : "left"`. Within the function it is passed into `includes`; assigned or updated later in the function.

**content:** `content` stores a computed value for temporary calculation. It is initialized from `block.html`. Within the function it is assigned or updated later in the function.

### function `richTextTerms(rich)`

Read `richTextTerms` first as a named idea, not as syntax. `richTextTerms` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function richTextTerms(rich) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rich`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `rich` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `richTextTerms` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `richTextTerms`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `richTextTerms`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderDynamicLinks(items = [])**

Read `renderDynamicLinks` first as a named idea, not as syntax. `renderDynamicLinks` turns saved data for `DynamicLinks` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderDynamicLinks(items = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `items`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `resourceLink`. Those calls show that `renderDynamicLinks` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `renderDynamicLinks`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local objects and variables inside `renderDynamicLinks` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**links:** links stores a function or callback value for temporary calculation. It is initialized from `items.filter((item) => item?.url && (item.label || item.title))`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; assigned or updated later in the function.

### **function collectRichDownloads(rich, fallbackType = "File", output = [])**

The best entry point for `collectRichDownloads` is its responsibility in the surrounding workflow.

`collectRichDownloads` gathers scattered nested data into a shape that another function can loop over. This is important because portfolio data is hierarchical while rendering, search, and AI context often need a flat list.

The syntax of the function is:

```
function collectRichDownloads(rich, fallbackType = "File", output = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rich`, `fallbackType`, `output`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `rich` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `fallbackType` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `output` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `cleanRichImageTitle`, `fileNameFromUrl`. Those calls show that `collectRichDownloads` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `collectRichDownloads`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `collectRichDownloads`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderDownloadBlock(title, items = [])**

The best entry point for `renderDownloadBlock` is its responsibility in the surrounding workflow.

`renderDownloadBlock` turns saved data for `DownloadBlock` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderDownloadBlock(title, items = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `title`, `items`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `title` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `uniqueDownloads`, `detailBlock`, `escapeHtml`, `resourceLink`, `renderMultilineInlineText`. Those calls show that `renderDownloadBlock` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `renderDownloadBlock`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The easiest way to follow the body of `renderDownloadBlock` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**downloads:** `downloads` stores a computed value for temporary calculation. It is initialized from `uniqueDownloads(items)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; assigned or updated later in the function.

**function `renderParsedBriefBlock(section, fallbackSummary = "")`**

Read `renderParsedBriefBlock` first as a named idea, not as syntax. `renderParsedBriefBlock` turns saved data for `ParsedBriefBlock` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderParsedBriefBlock(section, fallbackSummary = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `section`, `fallbackSummary`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `section` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `fallbackSummary` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `renderRichContent`. Those calls show that `renderParsedBriefBlock` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `renderParsedBriefBlock`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local objects and variables inside `renderParsedBriefBlock` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**briefItem:** `briefItem` stores a computed value for temporary calculation. It is initialized from `section?.items?.[0] || {}`. Within the function it is accesses properties/methods `description`, `rich`; passed into `renderRichContent`; assigned or updated later in the function.

**briefText:** `briefText` stores a computed value for temporary calculation. It is initialized from `briefItem.description || fallbackSummary || ""`. Within the function it is passed into `renderRichContent`; assigned or updated later in the function.

### **function richHasRenderableContent(rich)**

To understand `richHasRenderableContent`, start with the problem it owns. `richHasRenderableContent` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function richHasRenderableContent(rich) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rich`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `rich` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `richHasRenderableContent` performs its rule directly from parameters, module state, standard APIs, or local calculations.



The practical importance of the function is easiest to see by imagining it removed. Without `richHasRenderableContent`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `richHasRenderableContent`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function richContainsMedia(rich)**

To understand `richContainsMedia`, start with the problem it owns. `richContainsMedia` belongs to rich content rendering. This function turns saved builder content into public HTML: styled text, formulas, code blocks, images, files, links, and safe markup. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function richContainsMedia(rich) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rich`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `rich` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `richContainsMedia` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `richContainsMedia`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

Inside `richContainsMedia`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderProjects()**

Read `renderProjects` first as a named idea, not as syntax. `renderProjects` rebuilds the visible project directory from current state. It checks the active category filter and search query, groups visible projects by category, creates the category sections, and updates the summary count. It is called after catalog load, filter changes, and search changes. The function is the public directory refresh mechanism. Without it, filters and searches might update state internally but the page would not visually change.

The syntax of the function is:

```
function renderProjects() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local

constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

The function also has to be read through the helpers it calls. It works with `normalize`, `projectVisible`, `queueSearchHighlights`, `categorySection`. Those calls show that `renderProjects` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `renderProjects`, builder-created content would not survive cleanly into the website. Formatting, images, formulas, links, code blocks, or sanitized HTML could be lost or rendered unsafely.

The local objects and variables inside `renderProjects` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**query:** query stores a computed value for temporary calculation. It is initialized from `normalize(searchInput.value).trim()`. Within the function it is passed into `projectVisible`; assigned or updated later in the function.

**visible:** visible stores a function or callback value for temporary calculation. It is initialized from `projects.filter((project) => projectVisible(project, query))`. Within the function it is iterated or transformed as a collection; accesses properties/methods `filter`, `length`; assigned or updated later in the function.

**visibleProjects:** visibleProjects stores a function or callback value for portfolio data state. It is initialized from `visible.filter((project) => project.category === category.id)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is returned to the caller; accesses properties/methods `length`; passed into `categorySection`; assigned or updated later in the function.

**shouldShowEmptyCategory:** `shouldShowEmptyCategory` stores a computed value for portfolio data state. It is initialized from `!query && (activeFilter === "all" || activeFilter === category.id)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

## Project windows and navigation

The functions in this group all support project windows and navigation. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### **function clampSectionDialogPosition(left, top, dialog)**

The best entry point for `clampSectionDialogPosition` is its responsibility in the surrounding workflow. `clampSectionDialogPosition` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function clampSectionDialogPosition(left, top, dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `left`, `top`, `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `left` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `top` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `clampSectionDialogPosition` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `clampSectionDialogPosition`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `clampSectionDialogPosition` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**rect:** `rect` stores a computed value for temporary calculation. It is initialized from `dialog.getBoundingClientRect()`. Within the function it is accesses properties/methods `height`, `width`; passed into `max`; assigned or updated later in the function.

**margin:** `margin` stores a numeric value for temporary calculation. It is initialized from `12`. Within the function it is passed into `max`, `min`; assigned or updated later in the function.

**maxLeft:** `maxLeft` stores a computed value for temporary calculation. It is initialized from `Math.max(margin, window.innerWidth - rect.width - margin)`. Within the function it is assigned or updated later in the function.

**maxTop:** `maxTop` stores a computed value for temporary calculation. It is initialized from `Math.max(margin, window.innerHeight - rect.height - margin)`. Within the function it is assigned or updated later in the function.

### **function anchorSectionDialog(dialog)**

Begin with the role of `anchorSectionDialog`. `anchorSectionDialog` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function anchorSectionDialog(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `clampSectionDialogPosition`. Those calls show that `anchorSectionDialog` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `anchorSectionDialog`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `anchorSectionDialog` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**rect:** `rect` stores a computed value for temporary calculation. It is initialized from `dialog.getBoundingClientRect()`. Within the function it is accesses properties/methods `left`, `top`; passed into `clampSectionDialogPosition`; assigned or updated later in the function.

**position:** `position` stores a computed value for temporary calculation. It is initialized from `clampSectionDialogPosition(rect.left, rect.top, dialog)`. Within the function it is accesses properties/methods `left`, `top`; assigned or updated later in the function.

### **function ensureSectionResizeHandles(dialog)**

Begin with the role of `ensureSectionResizeHandles`. `ensureSectionResizeHandles` creates or repairs something only when it is needed. This pattern avoids duplicate DOM nodes, duplicate repositories, duplicate handlers, or missing folders while still letting later code assume the resource exists.

The syntax of the function is:

```
function ensureSectionResizeHandles(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing

history, or mutating controlled runtime state. A close reading of the body shows these implementation details: creates DOM elements instead of relying only on static HTML.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `ensureSectionResizeHandles` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `ensureSectionResizeHandles`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `ensureSectionResizeHandles` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**handle:** handle stores a DOM element reference. It is initialized from `document.createElement("div")`. Within the function it is accesses properties/methods `className`, `dataset`, `setAttribute`; passed into `append`; assigned or updated later in the function.

### **function anchorSectionDialogForResize(dialog)**

To understand `anchorSectionDialogForResize`, start with the problem it owns. `anchorSectionDialogForResize` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function anchorSectionDialogForResize(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

Next, trace the helper calls that surround the function. It works with `anchorSectionDialog`. Those calls show that `anchorSectionDialogForResize` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `anchorSectionDialogForResize`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `anchorSectionDialogForResize` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**rect:** rect stores a computed value for temporary calculation. It is initialized from `dialog.getBoundingClientRect()`. Within the function it is accesses properties/methods height, width; assigned or updated later in the function.

### **function sectionDialogResizeBounds(state, event)**

Read `sectionDialogResizeBounds` first as a named idea, not as syntax. `sectionDialogResizeBounds` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionDialogResizeBounds(state, event) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `state`, `event`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `state` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `event` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `sectionDialogResizeBounds` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `sectionDialogResizeBounds`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `sectionDialogResizeBounds` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**margin:** margin stores a numeric value for temporary calculation. It is initialized from 12. Within the function it is passed into `min`; assigned or updated later in the function.

**minWidth:** `minWidth` stores a computed value for temporary calculation. It is initialized from `Math.min(320, Math.max(180, window.innerWidth - margin * 2))`. Within the function it is passed into `min`; assigned or updated later in the function.

**minHeight:** `minHeight` stores a computed value for temporary calculation. It is initialized from `Math.min(170, Math.max(120, window.innerHeight - margin * 2))`. Within the function it is passed into `min`; assigned or updated later in the function.

**direction:** `direction` stores a computed value for filesystem location. It is initialized from `state.direction`. It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler

artifacts, or published assets. Within the function it is accesses properties/methods includes; assigned or updated later in the function.

**destructured { left, top, width, height }:** destructured { left, top, width, height } stores a computed value for imported or unpacked API handles. It is initialized from state.rect. It unpacks API handles from another object, giving the rest of the file short direct names for those APIs.

**dx:** dx stores a computed value for temporary calculation. It is initialized from event.clientX - state.startX. Within the function it is passed into min; assigned or updated later in the function.

**dy:** dy stores a computed value for temporary calculation. It is initialized from event.clientY - state.startY. Within the function it is passed into min; assigned or updated later in the function.

**right:** right stores a computed value for temporary calculation. It is initialized from state.rect.left + state.rect.width. Within the function it is assigned or updated later in the function.

**bottom:** bottom stores a computed value for temporary calculation. It is initialized from state.rect.top + state.rect.height. Within the function it is assigned or updated later in the function.

### **function sectionDialogFromDragEvent(event)**

To understand sectionDialogFromDragEvent, start with the problem it owns. sectionDialogFromDragEvent belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionDialogFromDragEvent(event) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: event. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. event is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means sectionDialogFromDragEvent performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without sectionDialogFromDragEvent, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside sectionDialogFromDragEvent reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.



**handle:** handle stores a computed value for temporary calculation. It is initialized from `event.target.closest("[data-section-dialog-drag='true']")`. Within the function it is accesses properties/methods `closest`; assigned or updated later in the function.

**dialog:** dialog stores a computed value for DOM reference or UI state. It is initialized from `handle.closest("dialog")`. It anchors UI behavior to a specific browser element or window state. Within the function it is returned to the caller; passed into `closest`; assigned or updated later in the function.

### **function beginSectionDialogDrag(event)**

Begin with the role of `beginSectionDialogDrag`. `beginSectionDialogDrag` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function beginSectionDialogDrag(event) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `event`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `event` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `sectionDialogFromDragEvent`, `anchorSectionDialog`. Those calls show that `beginSectionDialogDrag` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `beginSectionDialogDrag`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `beginSectionDialogDrag` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**dialog:** dialog stores a computed value for DOM reference or UI state. It is initialized from `sectionDialogFromDragEvent(event)`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods `classList`, `getBoundingClientRect`; passed into `add`, `anchorSectionDialog`; assigned or updated later in the function.

**rect:** rect stores a computed value for temporary calculation. It is initialized from `dialog.getBoundingClientRect()`. Within the function it is accesses properties/methods `left`, `top`; assigned or updated later in the function.

### **function moveSectionDialogDrag(event)**

Read `moveSectionDialogDrag` first as a named idea, not as syntax. `moveSectionDialogDrag` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function moveSectionDialogDrag(event) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `event`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `event` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `clampSectionDialogPosition`. Those calls show that `moveSectionDialogDrag` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `moveSectionDialogDrag`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `moveSectionDialogDrag` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**next:** `next` stores a computed value for temporary calculation. It is initialized from `clampSectionDialogPosition()`. Within the function it is accesses properties/methods `left`, `top`; assigned or updated later in the function.

### **function endSectionDialogDrag()**

Read `endSectionDialogDrag` first as a named idea, not as syntax. `endSectionDialogDrag` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function endSectionDialogDrag() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `endSectionDialogDrag` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `endSectionDialogDrag`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `endSectionDialogDrag`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function beginSectionDialogResize(event)**

Begin with the role of `beginSectionDialogResize`. `beginSectionDialogResize` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function beginSectionDialogResize(event) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `event`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `event` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `anchorSectionDialogForResize`. Those calls show that `beginSectionDialogResize` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `beginSectionDialogResize`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `beginSectionDialogResize` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**handle:** handle stores a computed value for temporary calculation. It is initialized from `event.target.closest("[data-section-dialog-resize-handle]")`. Within the function it is accesses properties/methods dataset; passed into `closest`; assigned or updated later in the function.

**dialog:** dialog stores a computed value for DOM reference or UI state. It is initialized from `handle?.closest("dialog")`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods `classList`, `getBoundingClientRect`; passed into `add`, `anchorSectionDialogForResize`, `closest`; assigned or updated later in the function.

### **function moveSectionDialogResize(event)**

Read `moveSectionDialogResize` first as a named idea, not as syntax. `moveSectionDialogResize` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function moveSectionDialogResize(event) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `event`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `event` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `sectionDialogResizeBounds`. Those calls show that `moveSectionDialogResize` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `moveSectionDialogResize`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `moveSectionDialogResize` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**next:** next stores a computed value for temporary calculation. It is initialized from `sectionDialogResizeBounds(activeSectionDialogResize, event)`. Within the function it is accesses properties/methods `height`, `left`, `top`, `width`; assigned or updated later in the function.

### **function endSectionDialogResize()**

Read `endSectionDialogResize` first as a named idea, not as syntax. `endSectionDialogResize` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function endSectionDialogResize() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `endSectionDialogResize` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `endSectionDialogResize`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `endSectionDialogResize`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function updateSectionDialogMinimize(dialog)**

Read `updateSectionDialogMinimize` first as a named idea, not as syntax. `updateSectionDialogMinimize` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function updateSectionDialogMinimize(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `updateSectionDialogMinimize` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `updateSectionDialogMinimize`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `updateSectionDialogMinimize` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**minimized:** `minimized` stores a computed value for temporary calculation. It is initialized from `dialog.classList.contains("is-minimized-dialog")`. Within the function it is passed into `contains`, `setAttribute`; assigned or updated later in the function.

**button:** `button` stores a DOM element reference. It is initialized from `dialog.querySelector(".section-view-minimize")`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods `setAttribute`, `textContent`, `title`; assigned or updated later in the function.

### **function toggleSectionDialogMinimized(dialog)**

To understand `toggleSectionDialogMinimized`, start with the problem it owns. `toggleSectionDialogMinimized` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function toggleSectionDialogMinimized(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

Next, trace the helper calls that surround the function. It works with `anchorSectionDialog`, `updateSectionDialogMinimize`. Those calls show that `toggleSectionDialogMinimized` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `toggleSectionDialogMinimized`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `toggleSectionDialogMinimized`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function enableSectionDialogDrag()**

Read `enableSectionDialogDrag` first as a named idea, not as syntax. `enableSectionDialogDrag` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function enableSectionDialogDrag() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: attaches event handlers that turn user actions into behavior.

The function also has to be read through the helpers it calls. It works with `clampSectionDialogPosition`. Those calls show that `enableSectionDialogDrag` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `enableSectionDialogDrag`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `enableSectionDialogDrag` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**rect:** `rect` stores a computed value for temporary calculation. It is initialized from `dialog.getBoundingClientRect()`. Within the function it is accesses properties/methods `left`, `top`; passed into `clampSectionDialogPosition`; assigned or updated later in the function.

**next:** `next` stores a computed value for temporary calculation. It is initialized from `clampSectionDialogPosition(rect.left, rect.top, dialog)`. Within the function it is accesses properties/methods `left`, `top`; assigned or updated later in the function.

### **function linkifyRichTextNodes(root)**

Read `linkifyRichTextNodes` first as a named idea, not as syntax. `linkifyRichTextNodes` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:



```
function linkifyRichTextNodes(root) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `root`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `root` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: creates DOM elements instead of relying only on static HTML.

The function also has to be read through the helpers it calls. It works with `normalizeLinkTarget`. Those calls show that `linkifyRichTextNodes` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `linkifyRichTextNodes`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `linkifyRichTextNodes` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**walker:** `walker` stores a DOM element reference. It is initialized from `document.createTreeWalker(root, NodeFilter.SHOW_TEXT, {`. Within the function it is accesses properties/methods `nextNode`; assigned or updated later in the function.

**nodes:** `nodes` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is iterated or transformed as a collection; mutated as a collection or DOM container; accesses properties/methods `forEach`, `push`; assigned or updated later in the function.

**node:** `node` stores a computed value for temporary calculation. It is initialized from `walker.nextNode()`. Within the function it is returned to the caller; accesses properties/methods `parentElement`, `textContent`; passed into `createTreeWalker`, `push`, `test`; assigned or updated later in the function.

**text:** `text` stores a computed value for temporary calculation. It is initialized from `textNode.textContent || ""`. Within the function it is accesses properties/methods `length`, `replace`, `slice`; passed into `append`; assigned or updated later in the function.

**fragment:** `fragment` stores a DOM element reference. It is initialized from `document.createDocumentFragment()`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `append`; passed into `replaceWith`; assigned or updated later in the function.

**cursor:** `cursor` stores a numeric value for temporary calculation. It is initialized from `0`. Within the function it is passed into `append`; assigned or updated later in the function.

**trailing:** `trailing` stores a computed value for temporary calculation. It is initialized from `match.match(/I),..`. Within the function it is accesses properties/methods `length`; passed into `append`, `slice`; assigned or updated later in the function.

**clean:** clean stores a computed value for temporary calculation. It is initialized from `trailing ? match.slice(0, -trailing.length) : match`. Within the function it is passed into `normalizeLinkTarget`; assigned or updated later in the function.

**link:** link stores a DOM element reference. It is initialized from `document.createElement("a")`. Within the function it accesses properties/methods `href`, `rel`, `target`, `textContent`; passed into `append`; assigned or updated later in the function.

### **function siteSectionHasContent(section)**

To understand `siteSectionHasContent`, start with the problem it owns. `siteSectionHasContent` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function siteSectionHasContent(section) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `section`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `section` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `siteSectionHasContent` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `siteSectionHasContent`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `siteSectionHasContent`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function siteSectionRenderable(section)**

Read `siteSectionRenderable` first as a named idea, not as syntax. `siteSectionRenderable` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function siteSectionRenderable(section) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `section`. Read those names as the contract

between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. section is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `siteSectionHasContent`. Those calls show that `siteSectionRenderable` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `siteSectionRenderable`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `siteSectionRenderable`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderSiteSections()**

To understand `renderSiteSections`, start with the problem it owns. `renderSiteSections` turns saved data for `SiteSections` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderSiteSections() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

Next, trace the helper calls that surround the function. It works with `escapeHtml`, `renderRichContent`, `renderDynamicLinks`. Those calls show that `renderSiteSections` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `renderSiteSections`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `renderSiteSections` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**mount:** `mount` stores a DOM element reference. It is initialized from `document.querySelector("#dynamic-sections")`. Within the function it is accesses properties/methods `innerHTML`; assigned or updated later in the function.

**visibleSections:** `visibleSections` stores a function or callback value for portfolio data state. It is initialized from `(siteSections || []).filter(siteSectionRenderable)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `map`; assigned or updated later in the function.

**style:** `style` stores a computed value for appearance configuration. It is initialized from `section.backgroundImage ? `style="--section-bg: url('${escapeHtml(section.backgroundImage)}')'` : ""`. Within the function it is assigned or updated later in the function.

**links:** `links` stores a list for temporary calculation. It is initialized from `[...(section.links || []), ...(section.assets || [])]`. Within the function it is passed into `renderDynamicLinks`; assigned or updated later in the function.

**subsections:** `subsections` stores a function or callback value for portfolio data state. It is initialized from `(section.subsections || []).filter((item) => item.title || item.description || item.richDescription?.blocks?.length || (item.links || []).length)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; assigned or updated later in the function.

### function customSectionBlocks(project)

Read `customSectionBlocks` first as a named idea, not as syntax. `customSectionBlocks` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function customSectionBlocks(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `detailBlock`, `evidenceList`, `resourceLink`. Those calls show that `customSectionBlocks` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `customSectionBlocks`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `customSectionBlocks`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderCollectedDownloadSections(project)**

The best entry point for `renderCollectedDownloadSections` is its responsibility in the surrounding workflow. `renderCollectedDownloadSections` turns saved data for `CollectedDownloadSections` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderCollectedDownloadSections(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `uniqueDownloads`, `collectDownloadsFromValue`, `renderDownloadBlock`, `displayTitle`. Those calls show that `renderCollectedDownloadSections` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `renderCollectedDownloadSections`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `renderCollectedDownloadSections` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**blocks:** `blocks` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods `join`, `push`; assigned or updated later in the function.

**addBlock:** `addBlock` stores a function or callback value for temporary calculation. It is initialized from `(title, source, fallbackType = "File") => {}`. Within the function it is assigned or updated later in the function.

**downloads:** `downloads` stores a computed value for temporary calculation. It is initialized from `uniqueDownloads(collectDownloadsFromValue(source, fallbackType, []))`. Within the function it is accesses properties/methods `length`; passed into `addBlock`, `push`; assigned or updated later in the function.

### **function pathToString(path = [])**

The best entry point for pathToString is its responsibility in the surrounding workflow. pathToString belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function pathToString(path = []) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: path. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. path points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: uses Node path utilities so Windows path separators and relative paths are handled consistently.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means pathToString performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without pathToString, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside pathToString, there are no local const, let, or var values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function pathFromString(value = "")**

Begin with the role of pathFromString. pathFromString belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function pathFromString(value = "") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: value. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. value is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and

generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `pathFromString` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `pathFromString`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `pathFromString`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **`function sectionRouteState(projectId, sectionIndex, resourcePath = "")`**

To understand `sectionRouteState`, start with the problem it owns. `sectionRouteState` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionRouteState(projectId, sectionIndex, resourcePath = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `projectId`, `sectionIndex`, `resourcePath`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `resourcePath` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `sectionRouteState` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `sectionRouteState`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.



Inside `sectionRouteState`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function canUseSectionHistory()**

Read `canUseSectionHistory` first as a named idea, not as syntax. `canUseSectionHistory` is a decision predicate. It answers one yes/no question so the larger workflow can branch clearly instead of embedding the same condition in several places.

The syntax of the function is:

```
function canUseSectionHistory() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser route/history. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: updates browser history so Back, Forward, and refresh can follow the same window state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `canUseSectionHistory` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `canUseSectionHistory`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `canUseSectionHistory`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function sectionBaseUrl()**

The best entry point for `sectionBaseUrl` is its responsibility in the surrounding workflow. `sectionBaseUrl` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionBaseUrl() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `sectionBaseUrl` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `sectionBaseUrl`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `sectionBaseUrl` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**url:** url stores a constructed object for network or routing value. It is initialized from new `URL(window.location.href)`. Within the function it is returned to the caller; accesses properties/methods `searchParams`; assigned or updated later in the function.

**function `sectionRouteUrl(projectId, sectionIndex, resourcePath = "")`**

The best entry point for `sectionRouteUrl` is its responsibility in the surrounding workflow. `sectionRouteUrl` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionRouteUrl(projectId, sectionIndex, resourcePath = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `projectId`, `sectionIndex`, `resourcePath`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `resourcePath` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `sectionBaseUrl`. Those calls show that `sectionRouteUrl` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `sectionRouteUrl`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `sectionRouteUrl` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**url:** url stores a computed value for network or routing value. It is initialized from `sectionBaseUrl()`. Within the function it is returned to the caller; accesses properties/methods `hash`, `searchParams`; assigned or updated later in the function.

### **function sectionRouteFromState(state)**

To understand `sectionRouteFromState`, start with the problem it owns. `sectionRouteFromState` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionRouteFromState(state) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `state`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `state` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `sectionRouteFromState` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `sectionRouteFromState`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `sectionRouteFromState`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function sectionRouteFromLocation()**

To understand `sectionRouteFromLocation`, start with the problem it owns. `sectionRouteFromLocation` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser

history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionRouteFromLocation() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser route/history. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects. A close reading of the body shows these implementation details: reads or writes URL query parameters so browser history can represent the current view.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `sectionRouteFromLocation` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `sectionRouteFromLocation`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `sectionRouteFromLocation` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**params:** `params` stores a constructed object for temporary calculation. It is initialized from `new URLSearchParams(window.location.search)`. Within the function it is accessed properties/methods `get`; assigned or updated later in the function.

**projectId:** `projectId` stores a computed value for portfolio data state. It is initialized from `params.get(sectionRouteParams.project)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**sectionIndex:** `sectionIndex` stores a computed value for portfolio data state. It is initialized from `params.get(sectionRouteParams.section)`. Within the function it is assigned or updated later in the function.

### **function sameSectionRoute(left, right)**

Read `sameSectionRoute` first as a named idea, not as syntax. `sameSectionRoute` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sameSectionRoute(left, right) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `left`, `right`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `left` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `right` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `sameSectionRoute` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `sameSectionRoute`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `sameSectionRoute`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

**`function syncSectionRouteToHistory(projectId, sectionIndex, resourcePath = "", mode = "push")`**

Read `syncSectionRouteToHistory` first as a named idea, not as syntax. `syncSectionRouteToHistory` reconciles two states that can drift apart. In this project, sync functions connect local drafts, route state, Git branches, publish mirrors, or frontend state to the current truth.

The syntax of the function is:

```
function syncSectionRouteToHistory(projectId, sectionIndex, resourcePath = "", mode = "push")
{ ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `projectId`, `sectionIndex`, `resourcePath`, `mode`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `resourcePath` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `mode` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser route/history. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `canUseSectionHistory`, `sectionRouteState`, `sectionRouteUrl`, `sectionRouteFromState`, `sectionRouteFromLocation`, `sameSectionRoute`. Those calls show that `syncSectionRouteToHistory` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `syncSectionRouteToHistory`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `syncSectionRouteToHistory` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**route:** route stores a resolved filesystem or route value. It is initialized from `sectionRouteState(projectId, sectionIndex, resourcePath)`. Within the function it is passed into `sameSectionRoute`; assigned or updated later in the function.

**url:** url stores a resolved filesystem or route value. It is initialized from `sectionRouteUrl(projectId, sectionIndex, resourcePath)`. Within the function it is assigned or updated later in the function.

**currentRoute:** `currentRoute` stores a computed value for temporary calculation. It is initialized from `sectionRouteFromState(history.state) || sectionRouteFromLocation()`. Within the function it is passed into `sameSectionRoute`; assigned or updated later in the function.

**method:** method stores a computed value for temporary calculation. It is initialized from `mode === "replace" || sameSectionRoute(currentRoute, route) ? "replaceState" : "pushState"`. Within the function it is assigned or updated later in the function.

### function clearSectionRouteInHistory()

Begin with the role of `clearSectionRouteInHistory`. `clearSectionRouteInHistory` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function clearSectionRouteInHistory() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

The body is where the contract above becomes behavior. Inside the body, it updates browser route/history. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: updates browser history so Back, Forward, and refresh can follow the same window state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `canUseSectionHistory`, `sectionBaseUrl`. Those calls show that `clearSectionRouteInHistory` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `clearSectionRouteInHistory`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `clearSectionRouteInHistory`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function initializeSectionHistoryState()**

Read `initializeSectionHistoryState` first as a named idea, not as syntax. `initializeSectionHistoryState` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function initializeSectionHistoryState() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser route/history. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: updates browser history so Back, Forward, and refresh can follow the same window state.

The function also has to be read through the helpers it calls. It works with `canUseSectionHistory`, `sectionRouteFromLocation`, `sectionRouteFromState`, `sectionRouteState`, `sectionRouteUrl`. Those calls show that `initializeSectionHistoryState` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `initializeSectionHistoryState`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `initializeSectionHistoryState` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**route:** `route` stores a computed value for temporary calculation. It is initialized from `sectionRouteFromLocation()` || `sectionRouteFromState(history.state)`. Within the function it accesses properties/methods `projectId`, `resourcePath`, `sectionIndex`; passed into `replaceState`, `sectionRouteUrl`; assigned or updated later in the function.



### **function nodeAtPath(section, path = [])**

Begin with the role of nodeAtPath. nodeAtPath belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function nodeAtPath(section, path = []) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: section, path. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. section is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. path points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means nodeAtPath performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without nodeAtPath, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside nodeAtPath explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**node:** node stores a computed value for temporary calculation. It is initialized from section. Within the function it is returned to the caller; accesses properties/methods children; assigned or updated later in the function.

**children:** children stores a computed value for temporary calculation. It is initialized from section?.items || []. Within the function it is assigned or updated later in the function.

**index:** index starts without an immediate value. In this file that usually means it is filled by a loop, branch, callback, or later assignment inside the same function.

### **function nodeChildren(node)**

The best entry point for nodeChildren is its responsibility in the surrounding workflow. nodeChildren belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function nodeChildren(node) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `nodeChildren` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `nodeChildren`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `nodeChildren`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function nodeHasRenderableContent(node)**

To understand `nodeHasRenderableContent`, start with the problem it owns. `nodeHasRenderableContent` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function nodeHasRenderableContent(node) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `richHasRenderableContent`, `nodeChildren`. Those calls show that `nodeHasRenderableContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `nodeHasRenderableContent`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `nodeHasRenderableContent` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**children:** `children` stores a computed value for temporary calculation. It is initialized from `nodeChildren(node)`. Within the function it is accesses properties/methods some; assigned or updated later in the function.

### **function sectionHasRenderableContent(section)**

Begin with the role of `sectionHasRenderableContent`. `sectionHasRenderableContent` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function sectionHasRenderableContent(section) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `section`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `section` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `richHasRenderableContent`. Those calls show that `sectionHasRenderableContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `sectionHasRenderableContent`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `sectionHasRenderableContent`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function responsiveNodeHasMedia(node)**

Begin with the role of `responsiveNodeHasMedia`. `responsiveNodeHasMedia` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function responsiveNodeHasMedia(node) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `richContainsMedia`, `responsiveChildren`. Those calls show that `responsiveNodeHasMedia` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `responsiveNodeHasMedia`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `responsiveNodeHasMedia` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**url:** url stores a computed value for network or routing value. It is initialized from `String(node.url || "")`. Within the function it is passed into `String`, `test`; assigned or updated later in the function.

### **function responsiveNodeCount(node)**

To understand `responsiveNodeCount`, start with the problem it owns. `responsiveNodeCount` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function responsiveNodeCount(node) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `nodeHasRenderableContent`, `responsiveChildren`. Those calls show that `responsiveNodeCount` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `responsiveNodeCount`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `responsiveNodeCount`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function responsiveNodeDepth(node)**

To understand `responsiveNodeDepth`, start with the problem it owns. `responsiveNodeDepth` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function responsiveNodeDepth(node) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `nodeHasRenderableContent`, `responsiveChildren`. Those calls show that `responsiveNodeDepth` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `responsiveNodeDepth`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `responsiveNodeDepth` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**children:** `children` stores a computed value for temporary calculation. It is initialized from `responsiveChildren(node).filter(nodeHasRenderableContent)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; passed into `max`; assigned or updated later in the function.

### **function nodeSummary(title, rich, text, emptyMessage = "No summary has been added yet.")**

Begin with the role of nodeSummary. nodeSummary belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function nodeSummary(title, rich, text, emptyMessage = "No summary has been added yet.") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: title, rich, text, emptyMessage. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. title is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. rich is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. text is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. emptyMessage is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with escapeHtml, renderRichContent. Those calls show that nodeSummary is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without nodeSummary, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside nodeSummary, there are no local const, let, or var values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function nodeIsOverview(item = {})**

To understand nodeIsOverview, start with the problem it owns. nodeIsOverview belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function nodeIsOverview(item = {}) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: item. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalize`, `displayTitle`. Those calls show that `nodeOverview` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `nodeOverview`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `nodeOverview` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `normalize(displayTitle(item.title || item.label || ""))`. Within the function it is accessed properties/methods `endsWith`; passed into `normalize`; assigned or updated later in the function.

### **function nodeOverviewDetails(node, children = [])**

The best entry point for `nodeOverviewDetails` is its responsibility in the surrounding workflow.

`nodeOverviewDetails` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function nodeOverviewDetails(node, children = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`, `children`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered. `children` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `renderRichContent`, `renderInlineSectionFiles`, `nodeChildren`. Those calls show that `nodeOverviewDetails` is not isolated; it is one



piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `nodeOverviewDetails`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `nodeOverviewDetails` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**overviewItems:** `overviewItems` stores a computed value for temporary calculation. It is initialized from `children.filter(nodesOverview)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `forEach`; assigned or updated later in the function.

**blocks:** `blocks` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `join`, `length`, `push`; assigned or updated later in the function.

**itemFiles:** `itemFiles` stores a computed value for temporary calculation. It is initialized from `renderInlineSectionFiles(nodeChildren(item))`. Within the function it is passed into `push`; assigned or updated later in the function.

### **function `parsedNodeCard(node, projectId, sectionIndex, path)`**

The best entry point for `parsedNodeCard` is its responsibility in the surrounding workflow. `parsedNodeCard` interprets `dNodeCard` text and returns structured meaning. Parsing functions are needed because the rest of the app should not repeatedly scan raw strings when it can work with a stable object, array, or normalized value.

The syntax of the function is:

```
function parsedNodeCard(node, projectId, sectionIndex, path) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`, `projectId`, `sectionIndex`, `path`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `path` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `displayTitle`, `pathToString`, `escapeHtml`. Those calls show that `parsedNodeCard` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `parsedNodeCard`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `parsedNodeCard` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `displayTitle(node.title)`. Within the function it is passed into `displayTitle`, `escapeHtml`; assigned or updated later in the function.

**function `parsedNodeContent(node, projectId, sectionIndex, path = [])`**

Begin with the role of `parsedNodeContent`. `parsedNodeContent` interprets `dNodeContent` text and returns structured meaning. Parsing functions are needed because the rest of the app should not repeatedly scan raw strings when it can work with a stable object, array, or normalized value.

The syntax of the function is:

```
function parsedNodeContent(node, projectId, sectionIndex, path = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`, `projectId`, `sectionIndex`, `path`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `path` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `renderInlineSectionFiles`, `nodeChildren`, `nodesOverview`, `nodeOverviewDetails`, `parsedChildCards`. Those calls show that `parsedNodeContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `parsedNodeContent`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `parsedNodeContent` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**isContainer:** `isContainer` stores a computed value for temporary calculation. It is initialized from `!node.kind || node.kind === "subsection"`. Within the function it is assigned or updated later in the function.

**children:** children stores a computed value for temporary calculation. It is initialized from `nodeChildren(node)`. Within the function it is iterated or transformed as a collection; accesses properties/methods filter; passed into `nodeOverviewDetails`, `parsedChildCards`; assigned or updated later in the function.

**contentChildren:** `contentChildren` stores a function or callback value for temporary calculation. It is initialized from `children.filter((child) => !modelsOverview(child))`. Within the function it is passed into `renderInlineSectionFiles`; assigned or updated later in the function.

### **function parsedSectionContent(section, projectId, sectionIndex, path = [])**

The best entry point for `parsedSectionContent` is its responsibility in the surrounding workflow. `parsedSectionContent` interprets `dSectionContent` text and returns structured meaning. Parsing functions are needed because the rest of the app should not repeatedly scan raw strings when it can work with a stable object, array, or normalized value.

The syntax of the function is:

```
function parsedSectionContent(section, projectId, sectionIndex, path = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `section`, `projectId`, `sectionIndex`, `path`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `section` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `path` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `nodeAtPath`, `parsedNodeContent`. Those calls show that `parsedSectionContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `parsedSectionContent`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `parsedSectionContent` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**node:** node stores a resolved filesystem or route value. It is initialized from `path.length ? nodeAtPath(section, path)` : section. Within the function it is passed into `parsedNodeContent`; assigned or updated later in the function.

### **function parsedSection(section, index, project)**

Begin with the role of `parsedSection`. `parsedSection` interprets `dSection` text and returns structured meaning. Parsing functions are needed because the rest of the app should not repeatedly scan raw strings when it can work with a stable object, array, or normalized value.

The syntax of the function is:

```
function parsedSection(section, index, project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `section`, `index`, `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `section` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `index` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `escapeHtml`, `displayTitle`. Those calls show that `parsedSection` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `parsedSection`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `parsedSection` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**visibleItems:** `visibleItems` stores a function or callback value for temporary calculation. It is initialized from `(section.items || []).filter(nodeHasRenderableContent)`. Within the function it is accesses properties/methods some; assigned or updated later in the function.

**hasImages:** `hasImages` stores a function or callback value for temporary calculation. It is initialized from `visibleItems.some((item) => item.kind === "image")`. Within the function it is assigned or updated later in the function.

### **function categorySection(category, visibleProjects)**

The best entry point for `categorySection` is its responsibility in the surrounding workflow. `categorySection` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function categorySection(category, visibleProjects) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `category`, `visibleProjects`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `category` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `visibleProjects` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `escapeHtml`. Those calls show that `categorySection` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `categorySection`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `categorySection`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function ensureSectionDialog()**

The best entry point for `ensureSectionDialog` is its responsibility in the surrounding workflow.

`ensureSectionDialog` creates and wires the reusable full-window dialog used for project sections. It builds the top controls, back/forward behavior, close behavior, content container, and event delegation for child cards. It creates the window once and then later calls `reuse` it. The function exists so every project section window behaves consistently. If each click created its own markup by hand, close buttons, back buttons, keyboard behavior, and nested navigation would drift apart.

The syntax of the function is:

```
function ensureSectionDialog() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: serializes structured data before sending or saving it; creates DOM elements instead of relying only on static

HTML; writes generated markup into the page, which is why escaping and sanitizing helpers around it matter; toggles CSS classes to express UI state; attaches event handlers that turn user actions into behavior.

Its connection to the rest of the file matters as much as its local code. It works with `closeSectionDialog`, `clearSectionRouteInHistory`, `updateSectionDialogMinimize`, `pathFromString`, `sectionForwardStack`, `pathToString`, `openParsedSection`. Those calls show that `ensureSectionDialog` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `ensureSectionDialog`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `ensureSectionDialog` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**dialog:** `dialog` stores a DOM element reference. It is initialized from `document.querySelector("#section-view-dialog")`. It anchors UI behavior to a specific browser element or window state. Within the function it is returned to the caller; accesses properties/methods `addEventListener`, `classList`, `className`, `dataset`, `id`, `innerHTML`, `querySelector`; passed into `append`, `closeSectionDialog`, `createElement`, `openParsedSection`, `pathFromString`, `querySelector`, `remove`; assigned or updated later in the function.

**returnTarget:** `returnTarget` stores a DOM element reference. It is initialized from `document.querySelector()`. Within the function it is assigned or updated later in the function.

**path:** `path` stores a resolved filesystem or route value. It is initialized from `pathFromString(dialog.dataset.resourcePath || "")`. It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler artifacts, or published assets. Within the function it accesses properties/methods `length`, `pop`; passed into `closest`, `openParsedSection`, `push`; assigned or updated later in the function.

**stack:** `stack` stores a computed value for runtime state or cache. It is initialized from `sectionForwardStack(dialog)`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `pop`, `push`; passed into `stringify`; assigned or updated later in the function.

**stack:** `stack` stores a computed value for runtime state or cache. It is initialized from `sectionForwardStack(dialog)`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `pop`, `push`; passed into `stringify`; assigned or updated later in the function.

**nextPath:** `nextPath` stores a computed value for filesystem location. It is initialized from `stack.pop()`. It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler artifacts, or published assets. Within the function it is passed into `openParsedSection`; assigned or updated later in the function.

**nodeButton:** `nodeButton` stores a computed value for DOM reference or UI state. It is initialized from `event.target.closest("[data-resource-path]")`. It anchors UI behavior to a specific browser element or window state. Within the function it accesses properties/methods `dataset`; passed into `openParsedSection`; assigned or updated later in the function.

### **function sectionForwardStack(dialog)**

Begin with the role of `sectionForwardStack`. `sectionForwardStack` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.



The syntax of the function is:

```
function sectionForwardStack(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates. A close reading of the body shows these implementation details: parses JSON rather than treating incoming text as already trusted data.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `sectionForwardStack` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `sectionForwardStack`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `sectionForwardStack`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **function updateSectionNavigation(dialog, path)**

Begin with the role of `updateSectionNavigation`. `updateSectionNavigation` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function updateSectionNavigation(dialog, path) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`, `path`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered. `path` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.



The surrounding workflow becomes clearer when you notice its collaborators. It works with `sectionForwardStack`. Those calls show that `updateSectionNavigation` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `updateSectionNavigation`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local state inside `updateSectionNavigation` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**back:** back stores a DOM element reference. It is initialized from `dialog.querySelector(".section-view-back")`. Within the function it is accesses properties/methods `hidden`, `setAttribute`, `title`; passed into `querySelector`; assigned or updated later in the function.

### **function closeSectionDialog(dialog)**

The best entry point for `closeSectionDialog` is its responsibility in the surrounding workflow. `closeSectionDialog` reverses an open window or route state. It does more than hide markup: it must also clean navigation state, restore the previous view when appropriate, and avoid trapping the visitor inside a dialog.

The syntax of the function is:

```
function closeSectionDialog(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Its connection to the rest of the file matters as much as its local code. It works with `clearSectionRouteInHistory`. Those calls show that `closeSectionDialog` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `closeSectionDialog`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

Inside `closeSectionDialog`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function closeOrStepBackSectionDialog(dialog)**

To understand `closeOrStepBackSectionDialog`, start with the problem it owns. `closeOrStepBackSectionDialog` reverses an open window or route state. It does more than hide markup: it must also clean navigation state, restore the previous view when appropriate, and avoid trapping the visitor inside a dialog.

The syntax of the function is:

```
function closeOrStepBackSectionDialog(dialog) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `dialog`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `dialog` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: serializes structured data before sending or saving it.

Next, trace the helper calls that surround the function. It works with `pathFromString`, `sectionForwardStack`, `pathToString`, `openParsedSection`, `closeSectionDialog`. Those calls show that `closeOrStepBackSectionDialog` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `closeOrStepBackSectionDialog`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `closeOrStepBackSectionDialog` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**path:** `path` stores a resolved filesystem or route value. It is initialized from `pathFromString(dialog.dataset.resourcePath || "")`. It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler artifacts, or published assets. Within the function it accesses properties/methods `length`, `pop`; passed into `openParsedSection`, `push`; assigned or updated later in the function.

**stack:** `stack` stores a computed value for runtime state or cache. It is initialized from `sectionForwardStack(dialog)`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `push`; passed into `stringify`; assigned or updated later in the function.

**function `openParsedSection(projectId, sectionIndex, resourcePath = "", options = {})`**

Read `openParsedSection` first as a named idea, not as syntax. `openParsedSection` opens a project section or nested subsection as a full-window public view. It finds the project, finds the section, walks the requested subsection path, applies the selected appearance template, renders the overview and child sections, updates navigation buttons, and syncs browser history. This function is the heart of the website's project-window behavior. Without it, clicking Design, Simulation, Documentation, or any user-created section would either scroll the page awkwardly or render content in the wrong place.

The syntax of the function is:

```
function openParsedSection(projectId, sectionIndex, resourcePath = "", options = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `projectId`, `sectionIndex`, `resourcePath`, `options`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `projectId` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `sectionIndex` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `resourcePath` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `options` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter; toggles CSS classes to express UI state.

The function also has to be read through the helpers it calls. It works with `sectionHasRenderableContent`, `pathFromString`, `nodeAtPath`, `nodeHasRenderableContent`, `ensureSectionDialog`, `applyProjectTemplateToElement`, `pathToString`, `syncSectionRouteToHistory`. Those calls show that `openParsedSection` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `openParsedSection`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `openParsedSection` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**project:** `project` stores a function or callback value for portfolio data state. It is initialized from `projects.find((item) => item.id === projectId)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods `category`; passed into `applyProjectTemplateToElement`; assigned or updated later in the function.

**sections:** `sections` stores a function or callback value for portfolio data state. It is initialized from `(project?.portfolioView?.sections || []).filter((section) => section.id !== "brief" && sectionHasRenderableContent(section))`. Within the function it is assigned or updated later in the function.

**section:** `section` stores a computed value for portfolio data state. It is initialized from `sections[Number(sectionIndex)]`. Within the function it is accesses properties/methods `id`, `title`; passed into `displayTitle`, `filter`, `nodeAtPath`, `parsedSectionContent`, `querySelector`, `sectionHasRenderableContent`; assigned or updated later in the function.

**path:** `path` stores a resolved filesystem or route value. It is initialized from `pathFromString(resourcePath)`. It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler artifacts, or published assets. Within the function it is accesses properties/methods `length`; passed into

nodeAtPath, pathToString, syncSectionRouteToHistory, updateSectionNavigation; assigned or updated later in the function.

**node:** node stores a resolved filesystem or route value. It is initialized from path.length ? nodeAtPath(section, path) : section. Within the function it is accesses properties/methods title; passed into displayTitle; assigned or updated later in the function.

**dialog:** dialog stores a computed value for DOM reference or UI state. It is initialized from ensureSectionDialog(). It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods dataset, open, querySelector, scrollTop, showModal; passed into applyProjectTemplateToElement, updateSectionNavigation; assigned or updated later in the function.

**category:** category stores a function or callback value for portfolio data state. It is initialized from categories.find((item) => item.id === project.category) || {}. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods accent; passed into applyProjectTemplateToElement; assigned or updated later in the function.

### function closeSectionDialogForRoute()

To understand closeSectionDialogForRoute, start with the problem it owns. closeSectionDialogForRoute reverses an open window or route state. It does more than hide markup: it must also clean navigation state, restore the previous view when appropriate, and avoid trapping the visitor inside a dialog.

The syntax of the function is:

```
function closeSectionDialogForRoute() { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means closeSectionDialogForRoute performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without closeSectionDialogForRoute, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside closeSectionDialogForRoute reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**dialog:** dialog stores a DOM element reference. It is initialized from document.querySelector("#section-view-dialog"). It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods close; passed into querySelector; assigned or updated later in the function.

### **function applySectionRoute(route)**

Read `applySectionRoute` first as a named idea, not as syntax. `applySectionRoute` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function applySectionRoute(route) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `route`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `route` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `closeSectionDialogForRoute`, `openParsedSection`. Those calls show that `applySectionRoute` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `applySectionRoute`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The local objects and variables inside `applySectionRoute` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**opened:** `opened` stores a resolved filesystem or route value. It is initialized from `openParsedSection(route.projectId, route.sectionIndex, route.resourcePath, {})`. Within the function it is assigned or updated later in the function.

### **function restoreSectionRouteAfterCatalogLoad()**

To understand `restoreSectionRouteAfterCatalogLoad`, start with the problem it owns. `restoreSectionRouteAfterCatalogLoad` belongs to project windows and navigation. This function supports full-window project sections, nested subsections, browser history, back/forward behavior, or content tree traversal. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function restoreSectionRouteAfterCatalogLoad() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser route/history. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: attaches event handlers that turn user actions into behavior; updates browser history so Back, Forward, and refresh can follow the same window state.

Next, trace the helper calls that surround the function. It works with `initializeSectionHistoryState`, `applySectionRoute`, `sectionRouteFromLocation`, `sectionRouteFromState`, `ensureSearchPanel`, `goToSearchResult`, `clearSearchHighlights`, `closeSectionDialogForRoute`. Those calls show that `restoreSectionRouteAfterCatalogLoad` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `restoreSectionRouteAfterCatalogLoad`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

After the main flow, the working values inside `restoreSectionRouteAfterCatalogLoad` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**topLink:** `topLink` stores a computed value for temporary calculation. It is initialized from `event.target.closest('a[href="#top"]')`. Within the function it is assigned or updated later in the function.

**question:** `question` stores a computed value for temporary calculation. It is initialized from `aiAssistantInput?.value || ""`. Within the function it is passed into `answerAssistantQuestion`; assigned or updated later in the function.

**button:** `button` stores a computed value for DOM reference or UI state. It is initialized from `event.target.closest("[data-ai-source-id]")`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods `dataset`; passed into `get`; assigned or updated later in the function.

**source:** `source` stores a computed value for lookup collection. It is initialized from `assistantSourceMap.get(button.dataset.aiSourceId)`. Within the function it is passed into `closest`, `goToSearchResult`; assigned or updated later in the function.

**sectionButton:** `sectionButton` stores a computed value for portfolio data state. It is initialized from `event.target.closest("[data-section-project]")`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods `dataset`; passed into `openParsedSection`; assigned or updated later in the function.

### **function** `renderInlineSectionFiles(items = [])`

The best entry point for `renderInlineSectionFiles` is its responsibility in the surrounding workflow. `renderInlineSectionFiles` turns saved data for `InlineSectionFiles` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderInlineSectionFiles(items = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `items`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `uniqueDownloads`, `normalizeDownloadAsset`, `escapeHtml`, `normalizeLinkTarget`, `isWebsiteLinkItem`, `linkAttributes`, `downloadAttribute`, `fileNameFromUrl`. Those calls show that `renderInlineSectionFiles` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `renderInlineSectionFiles`, nested project navigation would break down. A visitor might click a section and see nothing, lose Back/Forward behavior, remain trapped in a dialog, or see content rendered in the wrong hierarchy.

The easiest way to follow the body of `renderInlineSectionFiles` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**downloads:** `downloads` stores a function or callback value for lookup collection. It is initialized from `uniqueDownloads(items.map((item) => normalizeDownloadAsset(item, "File")).filter(Boolean))`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; assigned or updated later in the function.

## Resources and links

The functions in this group all support resources and links. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### **function normalizeLinkTarget(target = "", options = {})**

Read `normalizeLinkTarget` first as a named idea, not as syntax. `normalizeLinkTarget` standardizes `LinkTarget` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeLinkTarget(target = "", options = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `target`, `options`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.



Those syntax pieces become practical once you connect them to the data being passed in. `target` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `options` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `looksLikeBareWebAddress`. Those calls show that `normalizeLinkTarget` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `normalizeLinkTarget`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `normalizeLinkTarget` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(target || "").trim()`. Within the function it is returned to the caller; assigned or updated later in the function.

**function linkAttributes(url, item = {})**

To understand `linkAttributes`, start with the problem it owns. `linkAttributes` belongs to resources and links. This function decides whether content is a website link, local download, uploaded file, or resource card. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function linkAttributes(url, item = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `url`, `item`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `url` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalizeLinkTarget`, `isWebsiteLinkItem`. Those calls show that `linkAttributes` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `linkAttributes`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `linkAttributes` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**target:** `target` stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(url, { assumeWeb: isWebsiteLinkItem(item, url) })`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function isLocalDownloadTarget(target = "")**

Read `isLocalDownloadTarget` first as a named idea, not as syntax. `isLocalDownloadTarget` is a decision predicate. It answers one yes/no question so the larger workflow can branch clearly instead of embedding the same condition in several places.

The syntax of the function is:

```
function isLocalDownloadTarget(target = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `target`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `target` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalizeLinkTarget`. Those calls show that `isLocalDownloadTarget` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `isLocalDownloadTarget`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `isLocalDownloadTarget` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**value:** value stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(target, { assumeWeb: true })`. Within the function it is passed into `Boolean, test`; assigned or updated later in the function.

#### **function downloadAttribute(target = "", item = {})**

Begin with the role of `downloadAttribute`. `downloadAttribute` belongs to resources and links. This function decides whether content is a website link, local download, uploaded file, or resource card. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function downloadAttribute(target = "", item = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `target`, `item`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `target` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `normalizeLinkTarget`, `isWebsiteLinkItem`, `isLocalDownloadTarget`. Those calls show that `downloadAttribute` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `downloadAttribute`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `downloadAttribute` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**value:** value stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(target, { assumeWeb: isWebsiteLinkItem(item, target) })`. Within the function it is passed into `isLocalDownloadTarget`; assigned or updated later in the function.

#### **function fileNameFromUrl(url = "")**

Begin with the role of `fileNameFromUrl`. `fileNameFromUrl` belongs to resources and links. This function decides whether content is a website link, local download, uploaded file, or resource card. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function fileNameFromUrl(url = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `url`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `url` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `fileNameFromUrl` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `fileNameFromUrl`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `fileNameFromUrl` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(url || "").split(/[?#]/)[0].split("/").pop() || "Download file"`. Within the function it is returned to the caller; passed into `decodeURIComponent`; assigned or updated later in the function.

**function resourceLink(item = {}, label = item.label || item.title || item.name || fileNameFromUrl(item.url || item.artifact))**

Read `resourceLink` first as a named idea, not as syntax. `resourceLink` belongs to resources and links. This function decides whether content is a website link, local download, uploaded file, or resource card. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function resourceLink(item = {}, label = item.label || item.title || item.name ||
  fileNameFromUrl(item.url || item.artifact)) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `item`, `label`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `label` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `normalizeLinkTarget`, `isWebsiteLinkItem`, `escapeHtml`, `linkAttributes`, `downloadAttribute`, `fileNameFromUrl`. Those calls show that `resourceLink` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `resourceLink`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `resourceLink` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**rawTarget:** `rawTarget` stores a computed value for temporary calculation. It is initialized from `item.url || item.artifact || item.href || item.file || item.path || item.src || ""`. Within the function it is passed into `normalizeLinkTarget`; assigned or updated later in the function.

**target:** `target` stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(rawTarget, { assumeWeb: isWebsiteLinkItem(item, rawTarget) })`. Within the function it is passed into `downloadAttribute`, `escapeHtml`, `linkAttributes`; assigned or updated later in the function.

### **function itemUrl(item = {})**

Read `itemUrl` first as a named idea, not as syntax. `itemUrl` belongs to resources and links. This function decides whether content is a website link, local download, uploaded file, or resource card. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function itemUrl(item = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `item`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `itemUrl` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `itemUrl`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `itemUrl`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

**`function normalizeDownloadAsset(item = {}, fallbackType = "File")`**

To understand `normalizeDownloadAsset`, start with the problem it owns. `normalizeDownloadAsset` standardizes `DownloadAsset` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeDownloadAsset(item = {}, fallbackType = "File") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `item`, `fallbackType`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `item` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `fallbackType` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Next, trace the helper calls that surround the function. It works with `itemUrl`, `itemLabel`. Those calls show that `normalizeDownloadAsset` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeDownloadAsset`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `normalizeDownloadAsset` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**url:** url stores a computed value for network or routing value. It is initialized from `itemUrl(item)`. Within the function it is assigned or updated later in the function.

### **function collectDownloadsFromValue(value, fallbackType = "File", output = [])**

Read `collectDownloadsFromValue` first as a named idea, not as syntax. `collectDownloadsFromValue` gathers scattered nested data into a shape that another function can loop over. This is important because portfolio data is hierarchical while rendering, search, and AI context often need a flat list.

The syntax of the function is:

```
function collectDownloadsFromValue(value, fallbackType = "File", output = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`, `fallbackType`, `output`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `fallbackType` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `output` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalizeDownloadAsset`, `collectRichDownloads`. Those calls show that `collectDownloadsFromValue` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `collectDownloadsFromValue`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `collectDownloadsFromValue` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**asset:** `asset` stores a computed value for lookup collection. It is initialized from `normalizeDownloadAsset(value, fallbackType)`. Within the function it is passed into `push`; assigned or updated later in the function.

### **function uniqueDownloads(items = [])**

Read `uniqueDownloads` first as a named idea, not as syntax. `uniqueDownloads` belongs to resources and links. This function decides whether content is a website link, local download, uploaded file, or resource card. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function uniqueDownloads(items = []) { ... }
```



There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `items`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `items` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `uniqueDownloads` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `uniqueDownloads`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `uniqueDownloads` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**seen:** `seen` stores a constructed object for lookup collection. It is initialized from `new Set()`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `add`, `has`; assigned or updated later in the function.

**key:** `key` stores a string value for temporary calculation. It is initialized from ``${item.url}::${item.title}``. Within the function it is passed into `add`; assigned or updated later in the function.

## Appearance and responsive presentation

The functions in this group all support appearance and responsive presentation. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### **function cleanFontFamily(value = "")**

Begin with the role of `cleanFontFamily`. `cleanFontFamily` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function cleanFontFamily(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and

generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `cleanFontFamily` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `cleanFontFamily`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `cleanFontFamily`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function normalizeFontPx(value)**

Begin with the role of `normalizeFontPx`. `normalizeFontPx` standardizes `FontPx` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeFontPx(value) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `normalizeFontPx` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `normalizeFontPx`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `normalizeFontPx` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**size:** size stores a computed value for temporary calculation. It is initialized from `Number(value)`. Within the function it is passed into `String`, `isFinite`; assigned or updated later in the function.

### **function normalizeTextColor(value = "")**

To understand `normalizeTextColor`, start with the problem it owns. `normalizeTextColor` standardizes `TextColor` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeTextColor(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `normalizeTextColor` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeTextColor`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `normalizeTextColor` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**color:** color stores a computed value for appearance configuration. It is initialized from `String(value || "").trim()`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function rgbColorToHex(value = "")**

Read `rgbColorToHex` first as a named idea, not as syntax. `rgbColorToHex` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function rgbColorToHex(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract

between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `rgbColorToHex` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `rgbColorToHex`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `rgbColorToHex` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**color:** `color` stores a computed value for appearance configuration. It is initialized from `String(value || "").trim()`. Within the function it is returned to the caller; accesses properties/methods `match`, `toLowerCase`; assigned or updated later in the function.

**shortHex:** `shortHex` stores a computed value for temporary calculation. It is initialized from `color.match(/^#([0-9a-f])([0-9a-f])([0-9a-f])$/i)`. Within the function it is assigned or updated later in the function.

**rgb:** `rgb` stores a computed value for temporary calculation. It is initialized from `color.match(/^rgba?\s*(\d+)\s*(\d+)\s*(\d+)\s*(.*)$/i)`. Within the function it is assigned or updated later in the function.

### **function normalizePlainFieldStyle(style = {})**

Read `normalizePlainFieldStyle` first as a named idea, not as syntax. `normalizePlainFieldStyle` standardizes `PlainFieldStyle` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizePlainFieldStyle(style = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `style`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `style` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `cleanFontFamily`, `normalizeFontPx`, `normalizeTextColor`, `rgbColorToHex`. Those calls show that `normalizePlainFieldStyle` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `normalizePlainFieldStyle`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `normalizePlainFieldStyle` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**normalized:** `normalized` stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is returned to the caller; accesses properties/methods `bold`, `color`, `fontFamily`, `fontPx`, `italic`, `underline`; assigned or updated later in the function.

**fontFamily:** `fontFamily` stores a computed value for temporary calculation. It is initialized from `cleanFontFamily(style.fontFamily || "")`. Within the function it is passed into `cleanFontFamily`; assigned or updated later in the function.

**fontPx:** `fontPx` stores a computed value for temporary calculation. It is initialized from `normalizeFontPx(style.fontPx || style.fontSize || "")`. Within the function it is passed into `normalizeFontPx`; assigned or updated later in the function.

**color:** `color` stores a computed value for appearance configuration. It is initialized from `normalizeTextColor(style.color || "")`. Within the function it is passed into `normalizeTextColor`, `rgbColorToHex`; assigned or updated later in the function.

### **function normalizeFieldStyles(styles = {})**

To understand `normalizeFieldStyles`, start with the problem it owns. `normalizeFieldStyles` standardizes `FieldStyles` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeFieldStyles(styles = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `styles`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `styles` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalizePlainFieldStyle`. Those calls show that `normalizeFieldStyles` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeFieldStyles`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `normalizeFieldStyles`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function plainFieldStyleToCss(style = {})**

To understand `plainFieldStyleToCss`, start with the problem it owns. `plainFieldStyleToCss` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function plainFieldStyleToCss(style = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `style`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `style` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalizePlainFieldStyle`. Those calls show that `plainFieldStyleToCss` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `plainFieldStyleToCss`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `plainFieldStyleToCss` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**normalized:** normalized stores a computed value for temporary calculation. It is initialized from `normalizePlainFieldStyle(style)`. Within the function it is accesses properties/methods `bold`, `color`, `fontFamily`, `fontPx`, `italic`, `underline`; passed into `push`; assigned or updated later in the function.

**rules:** rules stores a list for temporary calculation. It is initialized from `[]`. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods `join`, `push`; assigned or updated later in the function.

#### **function applyPlainFieldStyle(element, fieldId = "")**

The best entry point for `applyPlainFieldStyle` is its responsibility in the surrounding workflow. `applyPlainFieldStyle` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function applyPlainFieldStyle(element, fieldId = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `element`, `fieldId`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `element` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered. `fieldId` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Its connection to the rest of the file matters as much as its local code. It works with `normalizePlainFieldStyle`. Those calls show that `applyPlainFieldStyle` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `applyPlainFieldStyle`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `applyPlainFieldStyle` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**style:** style stores a computed value for appearance configuration. It is initialized from `normalizePlainFieldStyle(fieldStyles[fieldId] || {})`. Within the function it is accesses properties/methods `bold`, `color`, `fontFamily`, `fontPx`, `fontSize`, `fontStyle`, `fontWeight`, `italic`; assigned or updated later in the function.

#### **function plainFieldStyleAttribute(fieldId = "")**

To understand `plainFieldStyleAttribute`, start with the problem it owns. `plainFieldStyleAttribute` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.



The syntax of the function is:

```
function plainFieldStyleAttribute(fieldId = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `fieldId`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `fieldId` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `plainFieldStyleToCss`, `escapeHtml`. Those calls show that `plainFieldStyleAttribute` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `plainFieldStyleAttribute`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `plainFieldStyleAttribute` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**css:** `css` stores a computed value for temporary calculation. It is initialized from `plainFieldStyleToCss(fieldStyles[fieldId] || {})`. Within the function it is returned to the caller; passed into `escapeHtml`; assigned or updated later in the function.

**function normalizeCropAspect(value = "")**

To understand `normalizeCropAspect`, start with the problem it owns. `normalizeCropAspect` standardizes `CropAspect` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCropAspect(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `normalizeCropAspect` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeCropAspect`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `normalizeCropAspect`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function normalizeCropZoom(value = 1)**

Read `normalizeCropZoom` first as a named idea, not as syntax. `normalizeCropZoom` standardizes `CropZoom` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCropZoom(value = 1) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `normalizeCropZoom` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `normalizeCropZoom`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `normalizeCropZoom` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**zoom:** zoom stores a computed value for temporary calculation. It is initialized from Number(value). Within the function it is passed into isFinite, min; assigned or updated later in the function.

### **function normalizeCropPosition(value = 50)**

Read normalizeCropPosition first as a named idea, not as syntax. normalizeCropPosition standardizes CropPosition data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCropPosition(value = 50) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: value. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. value is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means normalizeCropPosition performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without normalizeCropPosition, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside normalizeCropPosition show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**position:** position stores a computed value for temporary calculation. It is initialized from Number(value). Within the function it is passed into isFinite, min; assigned or updated later in the function.

### **function canonicalTemplateId(id)**

To understand canonicalTemplateId, start with the problem it owns. canonicalTemplateId is a decision predicate. It answers one yes/no question so the larger workflow can branch clearly instead of embedding the same condition in several places.

The syntax of the function is:

```
function canonicalTemplateId(id) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: id. Read those names as the contract

between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `id` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `canonicalTemplateId` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `canonicalTemplateId`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `canonicalTemplateId`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function projectTemplateId(project)**

Begin with the role of `projectTemplateId`. `projectTemplateId` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectTemplateId(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `canonicalTemplateId`. Those calls show that `projectTemplateId` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `projectTemplateId`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `projectTemplateId`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function projectTemplateClass(project)**

To understand `projectTemplateClass`, start with the problem it owns. `projectTemplateClass` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectTemplateClass(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `projectTemplateId`. Those calls show that `projectTemplateClass` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `projectTemplateClass`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `projectTemplateClass` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**templateId:** `templateId` stores a computed value for appearance configuration. It is initialized from `projectTemplateId(project)`. Within the function it is returned to the caller; assigned or updated later in the function.

### **function responsiveClassName(value, fallback)**

To understand `responsiveClassName`, start with the problem it owns. `responsiveClassName` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function responsiveClassName(value, fallback) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`, `fallback`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `fallback` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `responsiveClassName` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `responsiveClassName`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `responsiveClassName`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function projectResponsiveClass(project)**

The best entry point for `projectResponsiveClass` is its responsibility in the surrounding workflow. `projectResponsiveClass` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectResponsiveClass(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `projectResponsiveProfile`, `responsiveClassName`. Those calls show that `projectResponsiveClass` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `projectResponsiveClass`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `projectResponsiveClass` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**profile:** profile stores a computed value for portfolio data state. It is initialized from `projectResponsiveProfile(project)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods `cardLayout`, `density`; passed into `responsiveClassName`; assigned or updated later in the function.

### **function responsiveStyleValues(project)**

The best entry point for `responsiveStyleValues` is its responsibility in the surrounding workflow. `responsiveStyleValues` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function responsiveStyleValues(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Its connection to the rest of the file matters as much as its local code. It works with `projectResponsiveProfile`. Those calls show that `responsiveStyleValues` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `responsiveStyleValues`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `responsiveStyleValues` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**profile:** profile stores a computed value for portfolio data state. It is initialized from `projectResponsiveProfile(project)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods `contentMax`, `mediaMax`, `sectionCardMin`, `touchTarget`; assigned or updated later in the function.

### **function hasPublicTemplate(project)**

The best entry point for `hasPublicTemplate` is its responsibility in the surrounding workflow. `hasPublicTemplate` is a decision predicate. It answers one yes/no question so the larger workflow can branch clearly instead of embedding the same condition in several places.

The syntax of the function is:



```
function hasPublicTemplate(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `projectTemplateId`. Those calls show that `hasPublicTemplate` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `hasPublicTemplate`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `hasPublicTemplate`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function projectTemplateVisual(project)**

The best entry point for `projectTemplateVisual` is its responsibility in the surrounding workflow. `projectTemplateVisual` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectTemplateVisual(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `projectTemplateVisual` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `projectTemplateVisual`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `projectTemplateVisual`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function projectTemplateStyle(project, accent)**

Begin with the role of `projectTemplateStyle`. `projectTemplateStyle` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectTemplateStyle(project, accent) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`, `accent`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `accent` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `projectTemplateVisual`, `responsiveStyleValues`. Those calls show that `projectTemplateStyle` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `projectTemplateStyle`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local state inside `projectTemplateStyle` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**visual:** `visual` stores a computed value for temporary calculation. It is initialized from `projectTemplateVisual(project) || {}`. Within the function it accesses properties/methods `accent`, `background`, `click`, `clickText`, `hover`, `line`, `panel`, `text`; assigned or updated later in the function.

**values:** `values` stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is passed into entries; assigned or updated later in the function.

### **function applyProjectTemplateToElement(element, project, accent)**

Read `applyProjectTemplateToElement` first as a named idea, not as syntax. `applyProjectTemplateToElement` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function applyProjectTemplateToElement(element, project, accent) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `element`, `project`, `accent`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `element` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `accent` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The function also has to be read through the helpers it calls. It works with `projectTemplateClass`, `projectResponsiveClass`, `projectTemplateVisual`, `responsiveStyleValues`. Those calls show that `applyProjectTemplateToElement` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `applyProjectTemplateToElement`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `applyProjectTemplateToElement` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**visual:** `visual` stores a computed value for temporary calculation. It is initialized from `projectTemplateVisual(project) || {}`. Within the function it is accesses properties/methods `accent`, `background`, `click`, `clickText`, `hover`, `line`, `panel`, `text`; assigned or updated later in the function.

**values:** `values` stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is passed into entries; assigned or updated later in the function.

### **function responsiveChildren(node)**

The best entry point for `responsiveChildren` is its responsibility in the surrounding workflow. `responsiveChildren` belongs to appearance and responsive presentation. This function applies templates, visual style, font/color/crop settings, or responsive layout decisions. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function responsiveChildren(node) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `node`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `node` is a DOM object or parsed content node. The function reads from it, writes into it, or uses it as the current place where UI or project content should be rendered.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `responsiveChildren` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `responsiveChildren`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

Inside `responsiveChildren`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **function normalizeTextColor(value = "")**

To understand `normalizeTextColor`, start with the problem it owns. `normalizeTextColor` standardizes `TextColor` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeTextColor(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `normalizeTextColor` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `normalizeTextColor`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `normalizeTextColor` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**color:** `color` stores a computed value for appearance configuration. It is initialized from `String(value || "").trim()`. Within the function it is passed into `test`; assigned or updated later in the function.

## Project directory

The functions in this group all support project directory. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### function `flattenProject(project)`

To understand `flattenProject`, start with the problem it owns. `flattenProject` gathers scattered nested data into a shape that another function can loop over. This is important because portfolio data is hierarchical while rendering, search, and AI context often need a flat list.

The syntax of the function is:

```
function flattenProject(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Next, trace the helper calls that surround the function. It works with `slugLabel`. Those calls show that `flattenProject` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `flattenProject`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `flattenProject` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**categoryLabel:** `categoryLabel` stores a computed value for portfolio data state. It is initialized from `slugLabel(project.category)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is passed into `electronicsSearchKeywords`; assigned or updated later in the function.

**electronicsKeywords:** `electronicsKeywords` stores a computed value for temporary calculation. It is initialized from `window.electronicsSearchKeywords`. Within the function it is assigned or updated later in the function.

**richSummaryTerms:** richSummaryTerms stores a function or callback value for lookup collection. It is initialized from (project.summaryRich?.blocks || []).flatMap((block) => [. Within the function it is assigned or updated later in the function.

**parsedChildTerms:** parsedChildTerms stores a function or callback value for lookup collection. It is initialized from (project.portfolioView?.sections || []).flatMap((section) =>. Within the function it is assigned or updated later in the function.

**design:** design stores a computed value for temporary calculation. It is initialized from project.design || {}. Within the function it is accesses properties/methods brief, documentation, simulation; assigned or updated later in the function.

**designTerms:** designTerms stores a list for temporary calculation. It is initialized from [. Within the function it is assigned or updated later in the function.

### function projectVisible(project, query)

To understand projectVisible, start with the problem it owns. projectVisible handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectVisible(project, query) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: project, query. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. project identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. query is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with projectMatches. Those calls show that projectVisible is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without projectVisible, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside projectVisible reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**categoryMatch:** categoryMatch stores a computed value for portfolio data state. It is initialized from activeFilter === "all" || project.category === activeFilter. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is returned to the caller; assigned or updated later in the function.

### **function normalizeCategory(category = {})**

The best entry point for `normalizeCategory` is its responsibility in the surrounding workflow. `normalizeCategory` standardizes `Category` data before another part of the app uses it. The point is not decoration; it prevents different callers from interpreting the same input differently. A normalizer usually accepts loose user, catalog, or file data and turns it into the one shape the rest of the subsystem expects.

The syntax of the function is:

```
function normalizeCategory(category = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `category`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `category` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Its connection to the rest of the file matters as much as its local code. It works with `normalizeTextColor`. Those calls show that `normalizeCategory` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `normalizeCategory`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The easiest way to follow the body of `normalizeCategory` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**label:** `label` stores a computed value for temporary calculation. It is initialized from `String(category.label || category.title || category.id || "Project category").trim() || "Project category"`. Within the function it accesses properties/methods `toLowerCase`; passed into `String`; assigned or updated later in the function.

### **function hydrateProjectCategories(rawCategories = [], rawProjects = [])**

Read `hydrateProjectCategories` first as a named idea, not as syntax. `hydrateProjectCategories` belongs to project directory. This function normalizes, flattens, filters, or renders projects and categories. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function hydrateProjectCategories(rawCategories = [], rawProjects = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `rawCategories`, `rawProjects`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.



Those syntax pieces become practical once you connect them to the data being passed in. `rawCategories` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `rawProjects` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalizeCategory`. Those calls show that `hydrateProjectCategories` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `hydrateProjectCategories`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside `hydrateProjectCategories` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**normalized:** `normalized` stores a computed value for lookup collection. It is initialized from `rawCategories.map(normalizeCategory)`. Within the function it is returned to the caller; iterated or transformed as a collection; mutated as a collection or DOM container; accesses properties/methods `map`, `push`; passed into `Set`.

**knownIds:** `knownIds` stores a constructed object for lookup collection. It is initialized from `new Set(normalized.map((category) => category.id))`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `add`, `has`; assigned or updated later in the function.

**id:** `id` stores a computed value for temporary calculation. It is initialized from `String(project?.category || "").trim()`. Within the function it is passed into `add`, `push`; assigned or updated later in the function.

**label:** `label` stores a computed value for temporary calculation. It is initialized from `id`. Within the function it is passed into `push`; assigned or updated later in the function.

### function projectCard(project)

To understand `projectCard`, start with the problem it owns. `projectCard` handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectCard(project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Next, trace the helper calls that surround the function. It works with `sectionHasRenderableContent`, `projectTemplateClass`, `projectResponsiveClass`, `projectTemplateStyle`, `renderParsedBriefBlock`, `parsedSection`, `renderMultilineInlineText`, `detailBlock`. Those calls show that `projectCard` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `projectCard`, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

After the main flow, the working values inside `projectCard` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**category:** `category` stores a function or callback value for portfolio data state. It is initialized from `categories.find((item) => item.id === project.category) || {}`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods `accent`; assigned or updated later in the function.

**accent:** `accent` stores a computed value for temporary calculation. It is initialized from `category.accent || "#117c7a"`. Within the function it is passed into `projectTemplateStyle`; assigned or updated later in the function.

**sections:** `sections` stores a computed value for portfolio data state. It is initialized from `project.portfolioView.sections || []`. Within the function it is iterated or transformed as a collection; accesses properties/methods `filter`, `find`; assigned or updated later in the function.

**briefSection:** `briefSection` stores a function or callback value for portfolio data state. It is initialized from `sections.find((section) => section.id === "brief")`. Within the function it is passed into `renderParsedBriefBlock`; assigned or updated later in the function.

**otherSections:** `otherSections` stores a function or callback value for portfolio data state. It is initialized from `sections.filter((section) => section.id !== "brief" && sectionHasRenderableContent(section))`. Within the function it is iterated or transformed as a collection; accesses properties/methods `map`; assigned or updated later in the function.

**category:** `category` stores a function or callback value for portfolio data state. It is initialized from `categories.find((item) => item.id === project.category) || {}`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods `accent`; assigned or updated later in the function.

**accent:** `accent` stores a computed value for temporary calculation. It is initialized from `category.accent || "#117c7a"`. Within the function it is passed into `projectTemplateStyle`; assigned or updated later in the function.

## function loadProjectCatalog()

Read loadProjectCatalog first as a named idea, not as syntax. loadProjectCatalog is the public website startup pipeline. It chooses embedded preview data when the builder is previewing, otherwise it fetches projects.json from the published site. After that it hydrates the major runtime state: categories, projects, profile, site content, fun facts, custom sections, text styles, and rich text variants. Then it renders the visible page and restores any deep-linked project window. This function is async because fetching projects.json can take time and can fail. It is the reason the static HTML shell becomes a real portfolio. Without it, the page would show placeholders but no saved projects, profile details, resume, fun facts, project categories, search index, or assistant context.

The syntax of the function is:

```
function loadProjectCatalog() { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it calls a network resource and updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: throws explicit errors when a required safety or compatibility condition fails; writes generated markup into the page, which is why escaping and sanitizing helpers around it matter; performs a fetch and therefore must handle network delay and failure.

The function also has to be read through the helpers it calls. It works with hydrateProjectCategories, normalizeFieldStyles, normalizeSiteContent, normalizeProfile, normalizeFunFacts, renderProfile, renderSiteContent, renderFunFacts. Those calls show that loadProjectCatalog is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without loadProjectCatalog, the specific rule it owns would disappear from the named workflow. Callers would either skip that behavior or reimplement it inconsistently, which is exactly how rendering, publishing, compiling, searching, or AI context drifts over time.

The local objects and variables inside loadProjectCatalog show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**catalogUrl:** catalogUrl stores a constructed object for network or routing value. It is initialized from new URL("projects.json", window.location.href). It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods href, searchParams; passed into fetch; assigned or updated later in the function.

## Search and filtering

The functions in this group all support search and filtering. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### **function projectMatches(project, query)**

Begin with the role of projectMatches. projectMatches handles project-directory structure. It helps turn saved catalog data into visible categories, cards, filters, responsive layout, or searchable project text.

The syntax of the function is:

```
function projectMatches(project, query) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`, `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `flattenProject`. Those calls show that `projectMatches` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `projectMatches`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `projectMatches`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function setActiveFilter(filter = "all")**

Read `setActiveFilter` first as a named idea, not as syntax. `setActiveFilter` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function setActiveFilter(filter = "all") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `filter`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `filter` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM,

writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `setActiveFilter` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `setActiveFilter`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `setActiveFilter`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function bindFilterButtons()**

Read `bindFilterButtons` first as a named idea, not as syntax. `bindFilterButtons` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function bindFilterButtons() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: attaches event handlers that turn user actions into behavior.

The function also has to be read through the helpers it calls. It works with `setActiveFilter`, `renderProjects`, `updateSearchDropdown`. Those calls show that `bindFilterButtons` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `bindFilterButtons`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `bindFilterButtons`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderCategoryFilters()**

The best entry point for `renderCategoryFilters` is its responsibility in the surrounding workflow.

`renderCategoryFilters` turns saved data for `CategoryFilters` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderCategoryFilters() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

Its connection to the rest of the file matters as much as its local code. It works with `escapeHtml`, `bindFilterButtons`, `setActiveFilter`. Those calls show that `renderCategoryFilters` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `renderCategoryFilters`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `renderCategoryFilters`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function flattenSearchText(values = [])**

To understand `flattenSearchText`, start with the problem it owns. `flattenSearchText` gathers scattered nested data into a shape that another function can loop over. This is important because portfolio data is hierarchical while rendering, search, and AI context often need a flat list.

The syntax of the function is:

```
function flattenSearchText(values = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `values`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `values` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point

is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: serializes structured data before sending or saving it.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `flattenSearchText` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `flattenSearchText`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `flattenSearchText`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function addSearchEntry(entries, entry)**

To understand `addSearchEntry`, start with the problem it owns. `addSearchEntry` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function addSearchEntry(entries, entry) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `entries`, `entry`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `entries` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `entry` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `flattenSearchText`, `normalize`. Those calls show that `addSearchEntry` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `addSearchEntry`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

After the main flow, the working values inside `addSearchEntry` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**text:** `text` stores a computed value for temporary calculation. It is initialized from `flattenSearchText([entry.title, entry.type, entry.context, entry.text, entry.url])`. Within the function it is passed into `flattenSearchText`, `normalize`; assigned or updated later in the function.



**storedEntry:** storedEntry stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is returned to the caller; passed into push; assigned or updated later in the function.

**function searchSnippet(text = "", query = "")**

To understand searchSnippet, start with the problem it owns. searchSnippet belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function searchSnippet(text = "", query = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `text`, `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `text` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Next, trace the helper calls that surround the function. It works with `normalize`. Those calls show that `searchSnippet` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `searchSnippet`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

After the main flow, the working values inside `searchSnippet` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** clean stores a computed value for temporary calculation. It is initialized from `String(text || "").replace(/s+/g, " ").trim()`. Within the function it is returned to the caller; accesses properties/methods `length`, `slice`; passed into `min`, `normalize`; assigned or updated later in the function.

**index:** index stores a computed value for temporary calculation. It is initialized from `normalize(clean).indexOf(normalize(query))`. Within the function it is passed into `max`, `min`; assigned or updated later in the function.

**start:** start stores a computed value for temporary calculation. It is initialized from `Math.max(0, index - 54)`. Within the function it is passed into `slice`; assigned or updated later in the function.

**end:** end stores a computed value for temporary calculation. It is initialized from `Math.min(clean.length, index + query.length + 80)`. Within the function it is passed into `slice`; assigned or updated later in the function.

## **function entryScore(entry, query)**

Read `entryScore` first as a named idea, not as syntax. `entryScore` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function entryScore(entry, query) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `entry`, `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `entry` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalize`. Those calls show that `entryScore` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `entryScore`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local objects and variables inside `entryScore` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**normalizedQuery:** `normalizedQuery` stores a computed value for temporary calculation. It is initialized from `normalize(query).trim()`. Within the function it is accesses properties/methods `split`; assigned or updated later in the function.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `normalize(entry.title)`. Within the function it is accesses properties/methods `includes`, `startsWith`; passed into `normalize`; assigned or updated later in the function.

**type:** `type` stores a computed value for temporary calculation. It is initialized from `normalize(entry.type)`. Within the function it is accesses properties/methods `includes`; passed into `normalize`; assigned or updated later in the function.

**context:** `context` stores a computed value for temporary calculation. It is initialized from `normalize(entry.context)`. Within the function it is accesses properties/methods `includes`; passed into `normalize`; assigned or updated later in the function.

**text:** `text` stores a computed value for temporary calculation. It is initialized from `entry.normalizedText || ""`. Within the function it is accesses properties/methods `includes`; assigned or updated later in the function.

**words:** words stores a computed value for temporary calculation. It is initialized from `normalizedQuery.split(/\s+/).filter(Boolean)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `forEach`; assigned or updated later in the function.

**score:** score stores a numeric value for temporary calculation. It is initialized from 0. Within the function it is returned to the caller; assigned or updated later in the function.

### **function entryMatches(entry, query)**

The best entry point for `entryMatches` is its responsibility in the surrounding workflow. `entryMatches` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function entryMatches(entry, query) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `entry`, `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `entry` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `entryScore`. Those calls show that `entryMatches` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `entryMatches`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `entryMatches`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function fileIsBrowserSearchable(url = "")**

Read `fileIsBrowserSearchable` first as a named idea, not as syntax. `fileIsBrowserSearchable` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function fileIsBrowserSearchable(url = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `url`. Read those names as the contract

between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `url` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `filesBrowserSearchable` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `filesBrowserSearchable`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local objects and variables inside `filesBrowserSearchable` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(url || "").split(/[?#]/)[0].toLowerCase()`. Within the function it is passed into `test`; assigned or updated later in the function.

### **async function indexSearchableFileText(entry)**

The best entry point for `indexSearchableFileText` is its responsibility in the surrounding workflow. `indexSearchableFileText` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
async function indexSearchableFileText(entry) { ... }
```

The `async` keyword is important here: this function returns a `Promise`, so callers must wait for it before trusting the result or assuming the side effect has completed. The parameter list tells you what the caller must provide: `entry`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `entry` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it waits for asynchronous work and calls a network resource. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. Because it is `async`, callers must await it or handle its `Promise`; otherwise the UI or backend could move on before files, network calls, Git commands, compiler runs, or model responses have finished. A close reading of the body shows these implementation details: performs a fetch and therefore must handle network delay and failure.

Its connection to the rest of the file matters as much as its local code. It works with `filesBrowserSearchable`, `flattenSearchText`, `normalize`, `updateSearchDropdown`. Those calls show that `indexSearchableFileText` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `indexSearchableFileText`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The easiest way to follow the body of `indexSearchableFileText` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**response:** `response` stores data returned from a network call. It is initialized from `await fetch(entry.url)`. Within the function it is accessed via `properties/methods ok, text`; assigned or updated later in the function.

**text:** `text` stores the completed result of an awaited operation. It is initialized from `await response.text()`. Within the function it is accessed via `properties/methods slice`; passed into `flattenSearchText`, `normalize`; assigned or updated later in the function.

### **function addProjectSearchEntries(entries, project)**

Begin with the role of `addProjectSearchEntries`. `addProjectSearchEntries` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function addProjectSearchEntries(entries, project) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `entries`, `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `entries` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `slugLabel`, `richTextTerms`, `addSearchEntry`, `sectionHasRenderableContent`, `displayTitle`, `parsedItemTerms`, `fileNameFromUrl`, `itemUrl`. Those calls show that `addProjectSearchEntries` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `addProjectSearchEntries`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local state inside `addProjectSearchEntries` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**categoryLabel:** `categoryLabel` stores a computed value for portfolio data state. It is initialized from `slugLabel(project.category)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is passed into `addSearchEntry`, `electronicsSearchKeywords`; assigned or updated later in the function.

**baseContext:** `baseContext` stores a string value for temporary calculation. It is initialized from ``${categoryLabel} / ${project.title}``. Within the function it is passed into `addSearchEntry`; assigned or updated later in the function.

**projectTerms:** `projectTerms` stores a list for portfolio data state. It is initialized from `[]`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is passed into `addSearchEntry`; assigned or updated later in the function.

**sections:** `sections` stores a function or callback value for portfolio data state. It is initialized from `(project.portfolioView?.sections || []).filter((section) => section.id !== "brief" && sectionHasRenderableContent(section))`. Within the function it is iterated or transformed as a collection; accesses properties/methods `forEach`; assigned or updated later in the function.

**walkItems:** `walkItems` stores a function or callback value for temporary calculation. It is initialized from `(items = [], path = [], parentTitles = []) => {}`. Within the function it is assigned or updated later in the function.

**nextPath:** `nextPath` stores a list for filesystem location. It is initialized from `[...path, index]`. It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler artifacts, or published assets. Within the function it is passed into `pathToString`; assigned or updated later in the function.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `displayTitle(item.title || item.name || item.label || fileNameFromUrl(item.url || ""), "Project item")`. Within the function it is passed into `addSearchEntry`, `displayTitle`; assigned or updated later in the function.

**url:** `url` stores a computed value for network or routing value. It is initialized from `itemUrl(item)`. Within the function it is passed into `displayTitle`, `fileNameFromUrl`; assigned or updated later in the function.

**context:** `context` stores a list for temporary calculation. It is initialized from `[baseContext, displayTitle(section.title, "Section"), ...parentTitles].filter(Boolean).join(" / ")`. Within the function it is passed into `addSearchEntry`; assigned or updated later in the function.

**entry:** `entry` stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is passed into `addSearchEntry`; assigned or updated later in the function.

**storedEntry:** `storedEntry` stores a computed value for temporary calculation. It is initialized from `addSearchEntry(entries, entry)`. Within the function it is passed into `indexSearchableFileText`; assigned or updated later in the function.

**entry:** `entry` stores a structured object for temporary calculation. It is initialized from `{}`. Within the function it is passed into `addSearchEntry`; assigned or updated later in the function.

**storedEntry:** `storedEntry` stores a computed value for temporary calculation. It is initialized from `addSearchEntry(entries, entry)`. Within the function it is passed into `indexSearchableFileText`; assigned or updated later in the function.

## **function addSiteSectionSearchEntries(entries)**

The best entry point for `addSiteSectionSearchEntries` is its responsibility in the surrounding workflow. `addSiteSectionSearchEntries` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function addSiteSectionSearchEntries(entries) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `entries`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `entries` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Its connection to the rest of the file matters as much as its local code. It works with `normalizeProfile`, `addSearchEntry`, `indexSearchableFileText`, `richTextTerms`. Those calls show that `addSiteSectionSearchEntries` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `addSiteSectionSearchEntries`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The easiest way to follow the body of `addSiteSectionSearchEntries` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**current:** `current` stores a computed value for temporary calculation. It is initialized from `normalizeProfile(profile || {})`. Within the function it is accesses properties/methods `contactIntro`, `displayName`, `email`, `githubUrl`, `linkedinUrl`, `phone`, `portfolioLabel`, `resumeUrl`; passed into `addSearchEntry`; assigned or updated later in the function.

**pageSections:** `pageSections` stores a list for portfolio data state. It is initialized from `.`. Within the function it is iterated or transformed as a collection; accesses properties/methods `forEach`, `splice`; assigned or updated later in the function.

**resumeTextEntry:** `resumeTextEntry` stores a computed value for temporary calculation. It is initialized from `addSearchEntry(entries, {`. Within the function it is passed into `indexSearchableFileText`; assigned or updated later in the function.

## **function rebuildSearchIndex()**

To understand `rebuildSearchIndex`, start with the problem it owns. `rebuildSearchIndex` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.



The syntax of the function is:

```
function rebuildSearchIndex() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Next, trace the helper calls that surround the function. It works with `addSiteSectionSearchEntries`, `addProjectSearchEntries`. Those calls show that `rebuildSearchIndex` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `rebuildSearchIndex`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

After the main flow, the working values inside `rebuildSearchIndex` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**entries:** `entries` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is passed into `addProjectSearchEntries`, `addSiteSectionSearchEntries`; assigned or updated later in the function.

### **function searchResultsFor(query)**

The best entry point for `searchResultsFor` is its responsibility in the surrounding workflow. `searchResultsFor` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function searchResultsFor(query) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller

can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Its connection to the rest of the file matters as much as its local code. It works with `normalize`, `entryScore`. Those calls show that `searchResultsFor` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `searchResultsFor`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The easiest way to follow the body of `searchResultsFor` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**normalizedQuery:** `normalizedQuery` stores a computed value for temporary calculation. It is initialized from `normalize(query).trim()`. Within the function it is passed into `entryScore`; assigned or updated later in the function.

**seen:** `seen` stores a constructed object for lookup collection. It is initialized from `new Set()`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `add`, `has`; assigned or updated later in the function.

**key:** `key` stores a list for temporary calculation. It is initialized from `[entry.kind, entry.projectId, entry.sectionIndex, entry.resourcePath, entry.domTarget, entry.url, entry.title].join("|")`. Within the function it is passed into `add`; assigned or updated later in the function.

### **function highlightQueryHtml(value = "", query = "")**

The best entry point for `highlightQueryHtml` is its responsibility in the surrounding workflow. `highlightQueryHtml` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function highlightQueryHtml(value = "", query = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`, `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `escapeHtml`, `escapeRegExp`. Those calls show that `highlightQueryHtml` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `highlightQueryHtml`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The easiest way to follow the body of `highlightQueryHtml` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**escaped:** `escaped` stores a computed value for temporary calculation. It is initialized from `escapeHtml(value)`. Within the function it is returned to the caller; accesses properties/methods `replace`; assigned or updated later in the function.

**cleanQuery:** `cleanQuery` stores a computed value for temporary calculation. It is initialized from `String(query || "").trim()`. Within the function it is passed into `RegExp`; assigned or updated later in the function.

**pattern:** `pattern` stores a constructed object for temporary calculation. It is initialized from `new RegExp(`${escapeRegExp(escapeHtml(cleanQuery))}`, "ig")`. Within the function it is passed into `replace`; assigned or updated later in the function.

### function clearSearchHighlights()

Begin with the role of `clearSearchHighlights`. `clearSearchHighlights` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function clearSearchHighlights() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The surrounding workflow becomes clearer when you notice its collaborators. No other named helper from this file is detected inside the body. That means `clearSearchHighlights` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This is why the function is worth separating instead of burying the logic somewhere else. Without `clearSearchHighlights`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `clearSearchHighlights`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function `textNodeSearchTargets()`**

The best entry point for `textNodeSearchTargets` is its responsibility in the surrounding workflow. `textNodeSearchTargets` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function textNodeSearchTargets() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it updates browser DOM. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `textNodeSearchTargets` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `textNodeSearchTargets`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `textNodeSearchTargets`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function `applySearchHighlights(query)`**

The best entry point for `applySearchHighlights` is its responsibility in the surrounding workflow. `applySearchHighlights` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function applySearchHighlights(query) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `query`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. query is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with clearSearchHighlights, normalize, textNodeSearchTargets. Those calls show that applySearchHighlights is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without applySearchHighlights, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The easiest way to follow the body of applySearchHighlights is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**cleanQuery:** cleanQuery stores a computed value for temporary calculation. It is initialized from String(query || "").trim(). Within the function it is accesses properties/methods length; passed into indexOf, normalize, setEnd; assigned or updated later in the function.

**ranges:** ranges stores a list for temporary calculation. It is initialized from []. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods length, push; passed into set; assigned or updated later in the function.

**needle:** needle stores a computed value for temporary calculation. It is initialized from normalize(cleanQuery). Within the function it is passed into includes, indexOf; assigned or updated later in the function.

**walker:** walker stores a DOM element reference. It is initialized from document.createTreeWalker(root, NodeFilter.SHOW\_TEXT, {. Within the function it is accesses properties/methods nextNode; assigned or updated later in the function.

**node:** node stores a computed value for temporary calculation. It is initialized from walker.nextNode(). Within the function it is accesses properties/methods parentElement, textContent; passed into createTreeWalker, normalize, setEnd, setStart; assigned or updated later in the function.

**haystack:** haystack stores a computed value for runtime state or cache. It is initialized from normalize(node.textContent). Within the function it is accesses properties/methods indexOf; assigned or updated later in the function.

**index:** index stores a computed value for temporary calculation. It is initialized from haystack.indexOf(needle). Within the function it is passed into indexOf, setEnd, setStart; assigned or updated later in the function.

**range:** range stores a DOM element reference. It is initialized from document.createRange(). Within the function it is accesses properties/methods setEnd, setStart; passed into push; assigned or updated later in the function.

### **function queueSearchHighlights()**

Read `queueSearchHighlights` first as a named idea, not as syntax. `queueSearchHighlights` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function queueSearchHighlights() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it uses a timer to avoid blocking or to wait for another process. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `applySearchHighlights`. Those calls show that `queueSearchHighlights` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `queueSearchHighlights`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `queueSearchHighlights`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function ensureSearchPanel()**

Begin with the role of `ensureSearchPanel`. `ensureSearchPanel` creates or repairs something only when it is needed. This pattern avoids duplicate DOM nodes, duplicate repositories, duplicate handlers, or missing folders while still letting later code assume the resource exists.

The syntax of the function is:

```
function ensureSearchPanel() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details:

creates DOM elements instead of relying only on static HTML; writes generated markup into the page, which is why escaping and sanitizing helpers around it matter; attaches event handlers that turn user actions into behavior.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `updateSearchDropdown`, `goToSearchResult`. Those calls show that `ensureSearchPanel` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `ensureSearchPanel`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local state inside `ensureSearchPanel` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**shell:** `shell` stores a DOM element reference. It is initialized from `document.createElement("div")`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `append`, `className`; passed into `replaceWith`; assigned or updated later in the function.

**button:** `button` stores a computed value for DOM reference or UI state. It is initialized from `event.target.closest("[data-search-result-index]")`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods `dataset`; passed into `Number`; assigned or updated later in the function.

**result:** `result` stores a computed value for temporary calculation. It is initialized from `currentSearchResults[Number(button.dataset.searchResultIndex)]`. Within the function it is passed into `closest`, `goToSearchResult`, `querySelector`; assigned or updated later in the function.

#### **function setSearchMessage(message, state = "neutral")**

The best entry point for `setSearchMessage` is its responsibility in the surrounding workflow. `setSearchMessage` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function setSearchMessage(message, state = "neutral") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `message`, `state`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `message` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `state` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.



Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `setSearchMessage` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `setSearchMessage`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `setSearchMessage`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderSearchResults(query, visibleCount)**

Begin with the role of `renderSearchResults`. `renderSearchResults` turns saved data for `SearchResults` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderSearchResults(query, visibleCount) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `query`, `visibleCount`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `query` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. `visibleCount` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `ensureSearchPanel`, `searchResultsFor`, `setSearchMessage`, `highlightQueryHtml`, `escapeHtml`, `searchSnippet`. Those calls show that `renderSearchResults` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `renderSearchResults`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local state inside `renderSearchResults` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**list:** list stores a DOM element reference. It is initialized from `searchPanel.querySelector(".search-results-list")`. Within the function it is accesses properties/methods `innerHTML`; passed into `querySelector`; assigned or updated later in the function.

**cleanQuery:** cleanQuery stores a computed value for temporary calculation. It is initialized from `String(query || "").trim()`. Within the function it is passed into `highlightQueryHtml`, `searchResultsFor`, `searchSnippet`, `setAttribute`, `setSearchMessage`; assigned or updated later in the function.

**limitedResults:** limitedResults stores a computed value for temporary calculation. It is initialized from `currentSearchResults.slice(0, searchResultLimit)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `map`; passed into `setSearchMessage`; assigned or updated later in the function.

**visibleText:** visibleText stores a computed value for temporary calculation. It is initialized from `visibleCount`. Within the function it is passed into `setSearchMessage`; assigned or updated later in the function.

### function updateSearchDropdown()

Read `updateSearchDropdown` first as a named idea, not as syntax. `updateSearchDropdown` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function updateSearchDropdown() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `applySearchHighlights`, `renderSearchResults`. Those calls show that `updateSearchDropdown` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `updateSearchDropdown`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local objects and variables inside `updateSearchDropdown` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**query:** query stores a computed value for temporary calculation. It is initialized from `searchInput?.value || ""`. Within the function it is passed into `applySearchHighlights`, `renderSearchResults`; assigned or updated later in the function.

**visibleCount:** visibleCount stores a computed value for temporary calculation. It is initialized from `applySearchHighlights(query)`. Within the function it is passed into `renderSearchResults`; assigned or updated later in the function.

### **function showSearchPanelIfNeeded()**

The best entry point for `showSearchPanelIfNeeded` is its responsibility in the surrounding workflow. `showSearchPanelIfNeeded` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function showSearchPanelIfNeeded() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Its connection to the rest of the file matters as much as its local code. It works with `updateSearchDropdown`. Those calls show that `showSearchPanelIfNeeded` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `showSearchPanelIfNeeded`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `showSearchPanelIfNeeded`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function goToSearchResult(result)**

Read `goToSearchResult` first as a named idea, not as syntax. `goToSearchResult` turns a search dropdown item into navigation. Depending on what the result represents, it may open a project window, jump to a main page section, highlight text on the current page, or open a file/link. This makes search feel like a search engine instead of a passive list. The function is necessary because search results can point to many different kinds of things: project titles, section overviews, files, custom sections, contact links, or page text. One navigation rule would not work for all of them.

The syntax of the function is:

```
function goToSearchResult(result) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `result`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. result is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `scrollToDomTarget`, `openParsedSection`, `queueSearchHighlights`, `setActiveFilter`, `renderProjects`. Those calls show that `goToSearchResult` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `goToSearchResult`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

The local objects and variables inside `goToSearchResult` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**target:** target stores a DOM element reference. It is initialized from `document.getElementById(result.projectId)`. Within the function it is assigned or updated later in the function.

### **function handleSearchInput()**

Read `handleSearchInput` first as a named idea, not as syntax. `handleSearchInput` belongs to search and filtering. This function builds, scores, filters, highlights, or navigates search results so the public website behaves like a real searchable portfolio. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function handleSearchInput() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. It works with `renderProjects`, `updateSearchDropdown`. Those calls show that `handleSearchInput` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `handleSearchInput`, search would become less like a search engine. Results could be missing, poorly ranked, not highlighted, or unable to navigate the visitor to the exact project, section, file, or phrase.

Inside `handleSearchInput`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

## Portfolio AI frontend

The functions in this group all support portfolio ai frontend. Read them as a small subsystem: the smaller helpers prepare data and the larger functions coordinate side effects or rendering.

### **function assistantEndpoint()**

The best entry point for `assistantEndpoint` is its responsibility in the surrounding workflow. `assistantEndpoint` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantEndpoint() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it updates browser DOM. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `assistantEndpoint` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantEndpoint`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantEndpoint` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**metaEndpoint:** `metaEndpoint` stores a DOM element reference. It is initialized from `document.querySelector('meta[name="portfolio-ai-endpoint"]')?.content || ""`. Within the function it is accesses properties/methods `trim`; passed into `String`; assigned or updated later in the function.

### **function assistantSourceLabel(result = {})**

The best entry point for `assistantSourceLabel` is its responsibility in the surrounding workflow. `assistantSourceLabel` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantSourceLabel(result = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `result`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `result` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `assistantSourceLabel` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantSourceLabel`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantSourceLabel` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `result.title || "Portfolio item"`. Within the function it is returned to the caller; assigned or updated later in the function.

**contextParts:** `contextParts` stores a computed value for temporary calculation. It is initialized from `String(result.context || "")`. Within the function it accesses properties/methods slice; assigned or updated later in the function.

**specificContext:** `specificContext` stores a resolved filesystem or route value. It is initialized from `contextParts.slice(-2).join(" / ")`. Within the function it is assigned or updated later in the function.

### **function assistantProjectSummary(project = {})**

The best entry point for `assistantProjectSummary` is its responsibility in the surrounding workflow. `assistantProjectSummary` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantProjectSummary(project = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one

portfolio project instead of letting work spill across unrelated projects. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

Its connection to the rest of the file matters as much as its local code. It works with `slugLabel`, `sectionHasRenderableContent`, `displayTitle`. Those calls show that `assistantProjectSummary` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantProjectSummary`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantProjectSummary` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**categoryLabel:** `categoryLabel` stores a computed value for portfolio data state. It is initialized from `slugLabel(project.category)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**sections:** `sections` stores a function or callback value for portfolio data state. It is initialized from `(project.portfolioView?.sections || [])`. Within the function it is assigned or updated later in the function.

#### **function assistantProjectTitleAcronyms(title = "")**

The best entry point for `assistantProjectTitleAcronyms` is its responsibility in the surrounding workflow. `assistantProjectTitleAcronyms` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantProjectTitleAcronyms(title = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `title`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `title` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.



Its connection to the rest of the file matters as much as its local code. It works with `normalize`. Those calls show that `assistantProjectTitleAcronyms` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantProjectTitleAcronyms`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantProjectTitleAcronyms` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**words:** `words` stores a computed value for temporary calculation. It is initialized from `normalize(title)`. Within the function it is accesses properties/methods `length`, `slice`; passed into `add`; assigned or updated later in the function.

**acronyms:** `acronyms` stores a constructed object for lookup collection. It is initialized from `new Set()`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `add`; assigned or updated later in the function.

**start:** `start` stores a numeric value for temporary calculation. It is initialized from 0. Within the function it is passed into `add`; assigned or updated later in the function.

**length:** `length` stores a numeric value for temporary calculation. It is initialized from 2. Within the function it is passed into `add`; assigned or updated later in the function.

**function `assistantProjectIsNamedInQuestion(question = "", project = {})`**

To understand `assistantProjectIsNamedInQuestion`, start with the problem it owns.

`assistantProjectIsNamedInQuestion` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantProjectIsNamedInQuestion(question = "", project = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `project`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `project` identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalize`, `assistantProjectTitleAcronyms`. Those calls show that `assistantProjectIsNamedInQuestion` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantProjectIsNamedInQuestion`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantProjectIsNamedInQuestion` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is accesses properties/methods `includes`, `split`; passed into `Set`; assigned or updated later in the function.

**fillerWords:** `fillerWords` stores a constructed object for lookup collection. It is initialized from `new Set(["with", "from", "into", "using", "project", "design", "system", "systems"])`. Within the function it is accesses properties/methods `has`; assigned or updated later in the function.

**rawTitle:** `rawTitle` stores a computed value for temporary calculation. It is initialized from `project.portfolioView?.title || project.title || project.id`. Within the function it is passed into `assistantProjectTitleAcronyms`, `normalize`; assigned or updated later in the function.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `normalize(rawTitle)`. Within the function it is accesses properties/methods `split`; assigned or updated later in the function.

**cleanWords:** `cleanWords` stores a constructed object for lookup collection. It is initialized from `new Set(clean.split(/\s+/))`. Within the function it is accesses properties/methods `has`; assigned or updated later in the function.

**acronymMatch:** `acronymMatch` stores a computed value for temporary calculation. It is initialized from `assistantProjectTitleAcronyms(rawTitle)`. Within the function it is assigned or updated later in the function.

**words:** `words` stores a function or callback value for temporary calculation. It is initialized from `title.split(/\s+/).filter((word) => word.length > 3 && !fillerWords.has(word))`. Within the function it is returned to the caller; iterated or transformed as a collection; accesses properties/methods `filter`, `length`; passed into `min`; assigned or updated later in the function.

**matches:** `matches` stores a function or callback value for temporary calculation. It is initialized from `words.filter((word) => clean.includes(word)).length`. Within the function it is assigned or updated later in the function.

### **function assistantNamedProjectIds(question = "")**

Read `assistantNamedProjectIds` first as a named idea, not as syntax. `assistantNamedProjectIds` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantNamedProjectIds(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract

between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `assistantProjectIdsNamedInQuestion`. Those calls show that `assistantNamedProjectIds` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `assistantNamedProjectIds`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside `assistantNamedProjectIds`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

#### **function assistantProjectTitleMatches(question = "")**

Begin with the role of `assistantProjectTitleMatches`. `assistantProjectTitleMatches` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantProjectTitleMatches(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `assistantNamedProjectIds`. Those calls show that `assistantProjectTitleMatches` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `assistantProjectTitleMatches`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside `assistantProjectTitleMatches`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function assistantUrlKind(url = "")**

The best entry point for `assistantUrlKind` is its responsibility in the surrounding workflow. `assistantUrlKind` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantUrlKind(url = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `url`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `url` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Its connection to the rest of the file matters as much as its local code. It works with `fileBrowserSearchable`. Those calls show that `assistantUrlKind` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantUrlKind`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantUrlKind` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `String(url || "").split(/[?#]/)[0].toLowerCase()`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function assistantAbsoluteUrl(url = "")**

Begin with the role of `assistantAbsoluteUrl`. `assistantAbsoluteUrl` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantAbsoluteUrl(url = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `url`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `url` points at an external resource, local file, route, Git remote, or public link. The function normally normalizes or validates it before using it so the app does not treat unsafe or malformed paths as trusted input. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `normalizeLinkTarget`, `looksLikeBareWebAddress`. Those calls show that `assistantAbsoluteUrl` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `assistantAbsoluteUrl`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside `assistantAbsoluteUrl` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**target:** `target` stores a computed value for temporary calculation. It is initialized from `normalizeLinkTarget(url, { assumeWeb: looksLikeBareWebAddress(url) || /^https?:\W/i.test(url) })`. Within the function it is returned to the caller; passed into URL; assigned or updated later in the function.

#### **function assistantLinkRecordsFromValue(value, context = "")**

The best entry point for `assistantLinkRecordsFromValue` is its responsibility in the surrounding workflow. `assistantLinkRecordsFromValue` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantLinkRecordsFromValue(value, context = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`, `context`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. `context` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `uniqueDownloads`, `collectDownloadsFromValue`, `assistantUrlKind`, `fileNameFromUrl`, `assistantAbsoluteUrl`. Those calls show that

assistantLinkRecordsFromValue is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without assistantLinkRecordsFromValue, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of assistantLinkRecordsFromValue is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**records:** records stores a list for temporary calculation. It is initialized from []. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods push; assigned or updated later in the function.

**downloads:** downloads stores a computed value for temporary calculation. It is initialized from uniqueDownloads(collectDownloadsFromValue(value, "Portfolio asset", [])). Within the function it is iterated or transformed as a collection; accesses properties/methods forEach; assigned or updated later in the function.

**rawUrl:** rawUrl stores a computed value for network or routing value. It is initialized from download.url || download.href || download.path || "". Within the function it is passed into assistantAbsoluteUrl, fileNameFromUrl, push; assigned or updated later in the function.

#### **function assistantUrlsFromTextValue(value, context = "")**

Read assistantUrlsFromTextValue first as a named idea, not as syntax. assistantUrlsFromTextValue is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantUrlsFromTextValue(value, context = "") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: value, context. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. value is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. context is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

The function also has to be read through the helpers it calls. It works with flattenSearchText, assistantUrlKind, fileNameFromUrl, assistantAbsoluteUrl. Those calls show that assistantUrlsFromTextValue is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `assistantUrlsFromTextValue`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `assistantUrlsFromTextValue` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**text:** text stores a computed value for temporary calculation. It is initialized from `flattenSearchText([value])`. Within the function it is accesses properties/methods `match`; assigned or updated later in the function.

**matches:** matches stores a computed value for temporary calculation. It is initialized from `text.match(/\b(?:https?:VV|www\.)[^s<>"])+/gi) || []`. Within the function it is iterated or transformed as a collection; accesses properties/methods `map`; passed into `Set`; assigned or updated later in the function.

### **function assistantPublicProfileLinks()**

The best entry point for `assistantPublicProfileLinks` is its responsibility in the surrounding workflow. `assistantPublicProfileLinks` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantPublicProfileLinks() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `fileNameFromUrl`, `assistantUrlKind`, `assistantAbsoluteUrl`, `normalizeProfile`. Those calls show that `assistantPublicProfileLinks` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantPublicProfileLinks`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantPublicProfileLinks` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**records:** records stores a list for temporary calculation. It is initialized from `[]`. Within the function it is returned to the caller; iterated or transformed as a collection; mutated as a collection or DOM container; accesses properties/methods `filter`, `push`; assigned or updated later in the function.



**add:** add stores a function or callback value for temporary calculation. It is initialized from (item = {}, context = "") => {}. Within the function it is assigned or updated later in the function.

**rawUrl:** rawUrl stores a computed value for network or routing value. It is initialized from item.url || item.href || item.path || "". Within the function it is passed into assistantAbsoluteUrl, fileNameFromUrl, push; assigned or updated later in the function.

**label:** label stores a computed value for temporary calculation. It is initialized from item.label || item.title || fileNameFromUrl(rawUrl). Within the function it is passed into add; assigned or updated later in the function.

**current:** current stores a computed value for temporary calculation. It is initialized from normalizeProfile(profile || {}). Within the function it is accesses properties/methods githubUrl, linkedinUrl, resumeUrl, websiteUrl; passed into add; assigned or updated later in the function.

### **function assistantProjectEvidence(project = {})**

Begin with the role of assistantProjectEvidence. assistantProjectEvidence is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantProjectEvidence(project = {}) { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: project. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. project identifies or carries the project being built, rendered, compiled, searched, or published. It links this function back to one portfolio project instead of letting work spill across unrelated projects. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

The surrounding workflow becomes clearer when you notice its collaborators. It works with uniqueDownloads, collectDownloadsFromValue, assistantUrlKind, fileNameFromUrl, assistantAbsoluteUrl, assistantUrlsFromTextValue. Those calls show that assistantProjectEvidence is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without assistantProjectEvidence, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside assistantProjectEvidence explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**downloads:** downloads stores a computed value for temporary calculation. It is initialized from uniqueDownloads(collectDownloadsFromValue(project, "Project asset", [])). Within the function it is accesses properties/methods slice; assigned or updated later in the function.

**downloadRecords:** downloadRecords stores a function or callback value for lookup collection. It is initialized from `downloads.slice(0, 26).map((download) => {`. Within the function it is assigned or updated later in the function.

**rawUrl:** rawUrl stores a computed value for network or routing value. It is initialized from `download.url || ""`. Within the function it is passed into `assistantAbsoluteUrl`, `assistantUrlKind`, `fileNameFromUrl`; assigned or updated later in the function.

**textLinkRecords:** textLinkRecords stores a computed value for temporary calculation. It is initialized from `assistantUrlsFromTextValue([project.summary, project.links, project.portfolioView], project.portfolioView?.title || project.title)`. Within the function it is assigned or updated later in the function.

### **function assistantQuestionIsCasual(question = "")**

The best entry point for `assistantQuestionIsCasual` is its responsibility in the surrounding workflow. `assistantQuestionIsCasual` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionIsCasual(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. It works with `normalize`. Those calls show that `assistantQuestionIsCasual` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantQuestionIsCasual`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantQuestionIsCasual` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** clean stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function assistantQuestionHasPortfolioIntent(question = "")**

Read `assistantQuestionHasPortfolioIntent` first as a named idea, not as syntax.

`assistantQuestionHasPortfolioIntent` is one piece of the Ask My Portfolio conversation system. It helps decide

what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionHasPortfolioIntent(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalize`, `assistantProjectTitleMatches`. Those calls show that `assistantQuestionHasPortfolioIntent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `assistantQuestionHasPortfolioIntent`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `assistantQuestionHasPortfolioIntent` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function assistantConversationSuggestsPortfolioContext(history = [])**

To understand `assistantConversationSuggestsPortfolioContext`, start with the problem it owns.

`assistantConversationSuggestsPortfolioContext` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantConversationSuggestsPortfolioContext(history = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `history`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `history` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting

it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalize`. Those calls show that `assistantConversationSuggestsPortfolioContext` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantConversationSuggestsPortfolioContext`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside `assistantConversationSuggestsPortfolioContext`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **`function assistantQuestionAllowsPublicLookup(question = "", intent = assistantQuestionIntent(question))`**

To understand `assistantQuestionAllowsPublicLookup`, start with the problem it owns.

`assistantQuestionAllowsPublicLookup` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionAllowsPublicLookup(question = "", intent =
  assistantQuestionIntent(question)) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `intent`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `intent` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalize`. Those calls show that `assistantQuestionAllowsPublicLookup` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantQuestionAllowsPublicLookup`, the assistant would either answer with less context, display unrelated

links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantQuestionAllowsPublicLookup` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

#### **function assistantQuestionLooksConceptual(question = "")**

Begin with the role of `assistantQuestionLooksConceptual`. `assistantQuestionLooksConceptual` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionLooksConceptual(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `normalize`, `assistantQuestionIsEngineeringRelated`. Those calls show that `assistantQuestionLooksConceptual` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `assistantQuestionLooksConceptual`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside `assistantQuestionLooksConceptual` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

#### **function assistantQuestionIntent(question = "", history = assistantChatHistory)**

Begin with the role of `assistantQuestionIntent`. `assistantQuestionIntent` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionIntent(question = "", history = assistantChatHistory) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `history`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `history` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `assistantQuestionIsCasual`, `normalize`, `assistantQuestionHasPortfolioIntent`, `assistantConversationSuggestsPortfolioContext`, `assistantQuestionIsEngineeringRelated`, `assistantQuestionLooksConceptual`. Those calls show that `assistantQuestionIntent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `assistantQuestionIntent`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside `assistantQuestionIntent` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**hasPortfolioIntent:** `hasPortfolioIntent` stores a computed value for temporary calculation. It is initialized from `assistantQuestionHasPortfolioIntent(question)`. Within the function it is assigned or updated later in the function.

```
function assistantQuestionTokens(question = "")
```

Read `assistantQuestionTokens` first as a named idea, not as syntax. `assistantQuestionTokens` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionTokens(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalize`. Those calls show that `assistantQuestionTokens` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `assistantQuestionTokens`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `assistantQuestionTokens` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**stopWords:** `stopWords` stores a constructed object for lookup collection. It is initialized from `new Set([`. Within the function it is accessed via `properties/methods` `has`; assigned or updated later in the function.

**function `assistantSpecificSourceTokens(question = "")`**

To understand `assistantSpecificSourceTokens`, start with the problem it owns. `assistantSpecificSourceTokens` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantSpecificSourceTokens(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `assistantQuestionTokens`. Those calls show that `assistantSpecificSourceTokens` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantSpecificSourceTokens`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantSpecificSourceTokens` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.



**genericTokens:** genericTokens stores a constructed object for lookup collection. It is initialized from new Set([. Within the function it is accesses properties/methods has; assigned or updated later in the function.

**function assistantPromptTargetsProjectLanding(question = "")**

Read assistantPromptTargetsProjectLanding first as a named idea, not as syntax.

assistantPromptTargetsProjectLanding is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantPromptTargetsProjectLanding(question = "") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: question. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. question is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with normalize. Those calls show that assistantPromptTargetsProjectLanding is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without assistantPromptTargetsProjectLanding, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside assistantPromptTargetsProjectLanding show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** clean stores a computed value for temporary calculation. It is initialized from normalize(question). Within the function it is passed into test; assigned or updated later in the function.

**function assistantEntryRelationScore(result = {}, question = "")**

Read assistantEntryRelationScore first as a named idea, not as syntax. assistantEntryRelationScore is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantEntryRelationScore(result = {}, question = "") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: result, question. Read those names as the

contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `result` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `assistantSpecificSourceTokens`, `normalize`. Those calls show that `assistantEntryRelationScore` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `assistantEntryRelationScore`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `assistantEntryRelationScore` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**tokens:** `tokens` stores a computed value for temporary calculation. It is initialized from `assistantSpecificSourceTokens(question)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `filter`, `length`; assigned or updated later in the function.

**haystack:** `haystack` stores a resolved filesystem or route value. It is initialized from `normalize([result.title, result.context, result.type, result.text, result.url].filter(Boolean).join(" "))`. Within the function it is accesses properties/methods `includes`; assigned or updated later in the function.

**tokenHits:** `tokenHits` stores a function or callback value for temporary calculation. It is initialized from `tokens.filter((token) => haystack.includes(token)).length`. Within the function it is assigned or updated later in the function.

**relation:** `relation` stores a computed value for temporary calculation. It is initialized from `tokenHits * 30`. Within the function it is returned to the caller; assigned or updated later in the function.

**function `assistantNamedProjectTopicTokens(question = "", namedIds = assistantNamedProjectIds(question))`**

To understand `assistantNamedProjectTopicTokens`, start with the problem it owns.

`assistantNamedProjectTopicTokens` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantNamedProjectTopicTokens(question = "", namedIds =
assistantNamedProjectIds(question)) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `namedIds`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `namedIds` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalize`, `assistantProjectTitleAcronyms`, `assistantSpecificSourceTokens`. Those calls show that `assistantNamedProjectTopicTokens` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantNamedProjectTopicTokens`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantNamedProjectTopicTokens` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**projectVocabulary:** `projectVocabulary` stores a constructed object for portfolio data state. It is initialized from `new Set()`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is mutated as a collection or DOM container; accesses properties/methods `add`, `has`; assigned or updated later in the function.

**title:** `title` stores a computed value for temporary calculation. It is initialized from `project.portfolioView?.title || project.title || project.id`. Within the function it is passed into `assistantProjectTitleAcronyms`, `normalize`; assigned or updated later in the function.

**function `assistantSourcesForDisplay(question = "", results = [], intent = assistantQuestionIntent(question))`**

The best entry point for `assistantSourcesForDisplay` is its responsibility in the surrounding workflow. `assistantSourcesForDisplay` chooses which portfolio links appear under an assistant answer. The user complained earlier that related links were generic; this function is the fix-point for that behavior. It scores source relation to the prompt, respects named projects and topic tokens, and avoids showing source buttons that do not belong to the answer context. The function matters because sources are a promise. A displayed link should help verify or explore the answer. If this function is too loose, the assistant looks random even when the text answer is good.

The syntax of the function is:

```
function assistantSourcesForDisplay(question = "", results = [], intent =
assistantQuestionIntent(question)) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `results`, `intent`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `results` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `intent` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Its connection to the rest of the file matters as much as its local code. It works with `assistantSpecificSourceTokens`, `assistantNamedProjectIds`, `assistantNamedProjectTopicTokens`, `normalize`, `assistantPromptTargetsProjectLanding`, `assistantEntryRelationScore`. Those calls show that `assistantSourcesForDisplay` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantSourcesForDisplay`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantSourcesForDisplay` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**tokens:** `tokens` stores a computed value for temporary calculation. It is initialized from `assistantSpecificSourceTokens(question)`. Within the function it is iterated or transformed as a collection; accesses properties/methods `filter`, `length`; assigned or updated later in the function.

**namedProjectIdList:** `namedProjectIdList` stores a computed value for portfolio data state. It is initialized from `assistantNamedProjectIds(question)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is passed into `Set`, `assistantNamedProjectTopicTokens`; assigned or updated later in the function.

**namedProjectIds:** `namedProjectIds` stores a constructed object for portfolio data state. It is initialized from `new Set(namedProjectIdList)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it accesses properties/methods `has`, `size`; assigned or updated later in the function.

**namedProjectTopicTokens:** `namedProjectTopicTokens` stores a computed value for portfolio data state. It is initialized from `assistantNamedProjectTopicTokens(question, namedProjectIdList)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is iterated or transformed as a collection; accesses properties/methods `filter`, `length`; assigned or updated later in the function.

**asksForLinks:** asksForLinks stores a computed value for temporary calculation. It is initialized from `\b(open|show|where|link|links|github|linkedin|resume|download|file|files|repo|repository|source\s+code)\b/.test(normalize(question))`. Within the function it is passed into slice; assigned or updated later in the function.

**asksForProjectLanding:** asksForProjectLanding stores a computed value for portfolio data state. It is initialized from `assistantPromptTargetsProjectLanding(question)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**pageLinkIntent:** pageLinkIntent stores a computed value for temporary calculation. It is initialized from `\b(resume|github|linkedin|contact|email|phone)\b/.test(normalize(question))`. Within the function it is assigned or updated later in the function.

**asksForSpecificArtifact:** asksForSpecificArtifact stores a computed value for temporary calculation. It is initialized from `\b(code|source|repo|repository|github|file|files|document|documents|test|tests|result|results|pcb|schematic|simulation|tool|tools)\b/.test(normalize(question))`. Within the function it is assigned or updated later in the function.

**filtered:** filtered stores a computed value for temporary calculation. It is initialized from results. Within the function it is returned to the caller; assigned or updated later in the function.

**haystack:** haystack stores a resolved filesystem or route value. It is initialized from `normalize([result.title, result.context, result.type, result.text, result.url].filter(Boolean).join(" "))`. Within the function it is accesses properties/methods includes; assigned or updated later in the function.

**tokenHits:** tokenHits stores a function or callback value for temporary calculation. It is initialized from `tokens.filter((token) => haystack.includes(token)).length`. Within the function it is assigned or updated later in the function.

**topicHits:** topicHits stores a function or callback value for temporary calculation. It is initialized from `namedProjectTopicTokens.filter((token) => haystack.includes(token)).length`. Within the function it is assigned or updated later in the function.

**function assistantContextForQuestion(question = "", intent = assistantQuestionIntent(question), knownResults = null)**

Read `assistantContextForQuestion` first as a named idea, not as syntax. `assistantContextForQuestion` builds the packet of portfolio evidence sent to the AI backend. It uses the question intent, search results, named project matches, profile links, project summaries, and available evidence to decide what the model should see. It is selective by design so a generic greeting does not receive a pile of unrelated project context. Without this function, the assistant would either know too little about the portfolio or always overload the model with the same generic links. This is the function that makes the AI choose context based on the conversation.

The syntax of the function is:

```
function assistantContextForQuestion(question = "", intent = assistantQuestionIntent(question),
knownResults = null) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `intent`, `knownResults`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing,

or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. intent is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. knownResults is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a structured object. That object is the contract the next layer uses instead of trying to interpret raw strings or side effects.

The function also has to be read through the helpers it calls. It works with searchResultsFor, assistantNamedProjectIds, normalize, assistantPublicProfileLinks, assistantUrlKind, assistantAbsoluteUrl, assistantProjectEvidence, normalizeProfile. Those calls show that assistantContextForQuestion is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without assistantContextForQuestion, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside assistantContextForQuestion show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**directResults:** directResults stores a list for filesystem location. It is initialized from ["general\_engineering", "general\_knowledge", "general\_conversation"].includes(intent). It controls a boundary or destination for file work, so an incorrect value can redirect uploads, previews, compiler artifacts, or published assets. Within the function it is iterated or transformed as a collection; accesses properties/methods map; passed into Set; assigned or updated later in the function.

**projectIds:** projectIds stores a list for portfolio data state. It is initialized from [...new Set(]. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**matchedProjects:** matchedProjects stores a computed value for portfolio data state. It is initialized from projectIds. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is iterated or transformed as a collection; accesses properties/methods flatMap, map; assigned or updated later in the function.

**includePortfolioLinkContext:** includePortfolioLinkContext stores a computed value for temporary calculation. It is initialized from intent === "portfolio\_specific". Within the function it is assigned or updated later in the function.

**wantsPublicCode:** wantsPublicCode stores a computed value for temporary calculation. It is initialized from `\b(github|repo|repository|repositories|source\s+code|code|snippet|firmware|verilog|systemverilog|python|javascript)\b/.test(normalize(question))`. Within the function it is assigned or updated later in the function.

**publicProfileFetches:** publicProfileFetches stores a computed value for portfolio data state. It is initialized from includePortfolioLinkContext. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**aGithub:** aGithub stores a computed value for Git publishing and authorization state. It is initialized from `/github/i.test(`${a.kind} || "" ${a.url} || ""`) ? 1 : 0`. It affects publishing, where mistakes can change the public website or expose the wrong files. Within the function it is assigned or updated later in the function.

**bGithub:** bGithub stores a computed value for Git publishing and authorization state. It is initialized from `/github/i.test(`${b.kind} || "" ${b.url} || ""`) ? 1 : 0`. It affects publishing, where mistakes can change the public website or expose the wrong files. Within the function it is returned to the caller; assigned or updated later in the function.

**resultFetches:** resultFetches stores a function or callback value for lookup collection. It is initialized from `directResults.map((result) => {`. Within the function it is assigned or updated later in the function.

**projectFetches:** projectFetches stores a function or callback value for portfolio data state. It is initialized from `matchedProjects.flatMap((project) => assistantProjectEvidence(project)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**sourceFetches:** sourceFetches stores a list for temporary calculation. It is initialized from `[`. Within the function it is iterated or transformed as a collection; accesses properties/methods filter; assigned or updated later in the function.

**uniqueSourceFetches:** uniqueSourceFetches stores a function or callback value for temporary calculation. It is initialized from `sourceFetches.filter((item, index, array) => (`. Within the function it is assigned or updated later in the function.

### **function assistantRemember(role, content = "")**

To understand assistantRemember, start with the problem it owns. assistantRemember is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantRemember(role, content = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `role`, `content`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `role` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `content` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means assistantRemember performs its rule directly from parameters, module state, standard APIs, or local calculations.



The practical importance of the function is easiest to see by imagining it removed. Without `assistantRemember`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantRemember` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**cleanRole:** `cleanRole` stores a computed value for temporary calculation. It is initialized from `role === "assistant" ? "assistant" : "user"`. Within the function it is passed into `push`; assigned or updated later in the function.

**cleanContent:** `cleanContent` stores a computed value for temporary calculation. It is initialized from `String(content || "").trim()`. Within the function it is accesses properties/methods slice; passed into `push`; assigned or updated later in the function.

**function assistantConversationForRequest(currentQuestion = "", priorConversation = assistantChatHistory)**

To understand `assistantConversationForRequest`, start with the problem it owns.

`assistantConversationForRequest` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantConversationForRequest(currentQuestion = "", priorConversation =
assistantChatHistory) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `currentQuestion`, `priorConversation`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `currentQuestion` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `priorConversation` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser route/history. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `assistantConversationForRequest` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantConversationForRequest`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantConversationForRequest` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**history:** history stores a function or callback value for runtime state or cache. It is initialized from `priorConversation.slice(-8).map((item) => ({`. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods `push`; assigned or updated later in the function.

**function assistantFallbackResults(question = "", intent = assistantQuestionIntent(question))**

Begin with the role of `assistantFallbackResults`. `assistantFallbackResults` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantFallbackResults(question = "", intent = assistantQuestionIntent(question))
{ ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `intent`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `intent` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

The body is where the contract above becomes behavior. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `normalize`, `assistantNamedProjectIds`, `searchResultsFor`, `slugLabel`, `flattenSearchText`. Those calls show that `assistantFallbackResults` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `assistantFallbackResults`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside `assistantFallbackResults` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question).trim()`. Within the function it is accesses properties/methods `includes`; passed into `searchResultsFor`; assigned or updated later in the function.

**namedProjectIds:** `namedProjectIds` stores a computed value for portfolio data state. It is initialized from `assistantNamedProjectIds(question)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is returned to the caller; accesses properties/methods `includes`, `length`; assigned or updated later in the function.

**results:** `results` stores a computed value for temporary calculation. It is initialized from `searchResultsFor(clean)`. Within the function it is returned to the caller; iterated or transformed as a collection; accesses properties/methods `filter`, `length`, `slice`; assigned or updated later in the function.

**focusedResults:** focusedResults stores a function or callback value for temporary calculation. It is initialized from `results.filter((result) => namedProjectIds.includes(result.projectId))`. Within the function it is returned to the caller; accesses properties/methods `length`, `slice`; assigned or updated later in the function.

**category:** category stores a function or callback value for portfolio data state. It is initialized from `categories.find((item) => clean.includes(normalize(item.label)) || clean.includes(normalize(item.id)))`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is accesses properties/methods `id`, `label`; passed into `slugLabel`; assigned or updated later in the function.

### **function assistantQuestionIsEngineeringRelated(question = "")**

Read `assistantQuestionIsEngineeringRelated` first as a named idea, not as syntax.

`assistantQuestionIsEngineeringRelated` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantQuestionIsEngineeringRelated(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

The function also has to be read through the helpers it calls. It works with `normalize`. Those calls show that `assistantQuestionIsEngineeringRelated` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `assistantQuestionIsEngineeringRelated`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `assistantQuestionIsEngineeringRelated` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

### **function assistantGeneralEngineeringAnswer(question = "")**

The best entry point for `assistantGeneralEngineeringAnswer` is its responsibility in the surrounding workflow.

`assistantGeneralEngineeringAnswer` is one piece of the Ask My Portfolio conversation system. It helps decide

what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantGeneralEngineeringAnswer(question = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Its connection to the rest of the file matters as much as its local code. It works with `normalize`, `assistantQuestionIsEngineeringRelated`. Those calls show that `assistantGeneralEngineeringAnswer` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantGeneralEngineeringAnswer`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantGeneralEngineeringAnswer` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

**function `assistantLocalAnswer(question = "", results = [], intent = assistantQuestionIntent(question))`**

To understand `assistantLocalAnswer`, start with the problem it owns. `assistantLocalAnswer` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantLocalAnswer(question = "", results = [], intent =  
  assistantQuestionIntent(question)) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `results`, `intent`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing,

or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. results is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. intent is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns an array. The caller usually iterates over that list to render entries, process files, or choose candidates.

Next, trace the helper calls that surround the function. It works with `normalizeProfile`, `normalize`, `assistantGeneralEngineeringAnswer`, `searchSnippet`. Those calls show that `assistantLocalAnswer` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantLocalAnswer`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantLocalAnswer` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**ownerName:** `ownerName` stores a computed value for temporary calculation. It is initialized from `normalizeProfile(profile || {}).displayName || "the portfolio owner"`. Within the function it is assigned or updated later in the function.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

**generalAnswer:** `generalAnswer` stores a computed value for temporary calculation. It is initialized from `assistantGeneralEngineeringAnswer(question)`. Within the function it is returned to the caller; assigned or updated later in the function.

**generalAnswer:** `generalAnswer` stores a computed value for temporary calculation. It is initialized from `assistantGeneralEngineeringAnswer(question)`. Within the function it is returned to the caller; assigned or updated later in the function.

**projectIds:** `projectIds` stores a list for portfolio data state. It is initialized from `[...new Set(results.map((item) => item.projectId).filter(Boolean))]`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is assigned or updated later in the function.

**matchedProjects:** `matchedProjects` stores a computed value for portfolio data state. It is initialized from `projectIds`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it is iterated or transformed as a collection; accesses properties/methods `map`; assigned or updated later in the function.

**projectNames:** `projectNames` stores a function or callback value for portfolio data state. It is initialized from `matchedProjects.map((project) => project.portfolioView?.title || project.title).slice(0, 3)`. It carries portfolio content, so rendering, search, AI context, and publishing expect its shape to stay consistent. Within the function it accesses properties/methods `join`, `length`; passed into `push`; assigned or updated later in the function.

**fileCount:** fileCount stores a function or callback value for temporary calculation. It is initialized from `results.filter((item) => item.kind === "file").length`. Within the function it is passed into push; assigned or updated later in the function.

**sectionCount:** sectionCount stores a function or callback value for portfolio data state. It is initialized from `results.filter((item) => item.kind === "section" || item.kind === "subsection").length`. Within the function it is passed into push; assigned or updated later in the function.

**pageCount:** pageCount stores a function or callback value for temporary calculation. It is initialized from `results.filter((item) => item.kind === "page").length`. Within the function it is passed into push; assigned or updated later in the function.

**parts:** parts stores a list for temporary calculation. It is initialized from `[]`. Within the function it is returned to the caller; mutated as a collection or DOM container; accesses properties/methods join, push; assigned or updated later in the function.

**firstSnippet:** firstSnippet stores a computed value for temporary calculation. It is initialized from `searchSnippet(results[0].text, question)`. Within the function it is passed into push; assigned or updated later in the function.

### **function assistantSourcesMarkup(results = [])**

Read assistantSourcesMarkup first as a named idea, not as syntax. assistantSourcesMarkup is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantSourcesMarkup(results = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `results`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `results` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

The function also has to be read through the helpers it calls. It works with `escapeHtml`, `assistantSourceLabel`. Those calls show that assistantSourcesMarkup is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without assistantSourcesMarkup, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside assistantSourcesMarkup show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**sources:** sources stores a computed value for temporary calculation. It is initialized from results.slice(0, 5). Within the function it is iterated or transformed as a collection; accesses properties/methods length, map; assigned or updated later in the function.

**sourceButtons:** sourceButtons stores a function or callback value for DOM reference or UI state. It is initialized from sources.map((source) => {). It anchors UI behavior to a specific browser element or window state. Within the function it is assigned or updated later in the function.

**id:** id stores a computed value for temporary calculation. It is initialized from String(assistantSourceCounter++). Within the function it is passed into set; assigned or updated later in the function.

### **function setAssistantStatus(message, state = "ready")**

Read setAssistantStatus first as a named idea, not as syntax. setAssistantStatus belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function setAssistantStatus(message, state = "ready") { ... }
```

There is no async keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: message, state. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. message is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. state is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means setAssistantStatus performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without setAssistantStatus, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside setAssistantStatus, there are no local const, let, or var values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function assistantDesktopDockedHeight()**

The best entry point for assistantDesktopDockedHeight is its responsibility in the surrounding workflow. assistantDesktopDockedHeight is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:



```
function assistantDesktopDockedHeight() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

After the syntax, move into the body and follow the work in the order the function performs it. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Its connection to the rest of the file matters as much as its local code. No other named helper from this file is detected inside the body. That means `assistantDesktopDockedHeight` performs its rule directly from parameters, module state, standard APIs, or local calculations.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `assistantDesktopDockedHeight`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `assistantDesktopDockedHeight` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**value:** value stores a computed value for temporary calculation. It is initialized from `getComputedStyle(aiAssistantPanel).getPropertyValue("--ai-console-rest-height").trim()`. Within the function it is accesses properties/methods endsWith; passed into `parseFloat`; assigned or updated later in the function.

### **function updateAssistantPanelGrowth()**

Begin with the role of `updateAssistantPanelGrowth`. `updateAssistantPanelGrowth` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function updateAssistantPanelGrowth() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `assistantDesktopDockedHeight`. Those calls show that `updateAssistantPanelGrowth` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `updateAssistantPanelGrowth`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside `updateAssistantPanelGrowth` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**isDesktop:** `isDesktop` stores a computed value for temporary calculation. It is initialized from `window.matchMedia("(min-width: 901px)").matches`. Within the function it is assigned or updated later in the function.

**dockedHeight:** `dockedHeight` stores a computed value for temporary calculation. It is initialized from `assistantDesktopDockedHeight()`. Within the function it is passed into `round`, `toggle`; assigned or updated later in the function.

**formHeight:** `formHeight` stores a computed value for temporary calculation. It is initialized from `aiAssistantForm.getBoundingClientRect().height || 46`. Within the function it is assigned or updated later in the function.

**clearHeight:** `clearHeight` stores a computed value for temporary calculation. It is initialized from `aiClearChatButton?.getBoundingClientRect().height || 38`. Within the function it is assigned or updated later in the function.

**titleHeight:** `titleHeight` stores a DOM element reference. It is initialized from `document.querySelector("#ai-assistant-title")?.getBoundingClientRect().height || 34`. Within the function it is assigned or updated later in the function.

**chromeHeight:** `chromeHeight` stores a computed value for temporary calculation. It is initialized from `titleHeight + formHeight + clearHeight + 92`. Within the function it is assigned or updated later in the function.

**messages:** `messages` stores a list for temporary calculation. It is initialized from `[...aiAssistantLog.querySelectorAll(".ai-message")]`. Within the function it is iterated or transformed as a collection; accesses properties/methods `length`, `reduce`; assigned or updated later in the function.

**messageGap:** `messageGap` stores a computed value for temporary calculation. It is initialized from `messages.length > 1 ? (messages.length - 1) * 10 : 0`. Within the function it is assigned or updated later in the function.

**messageHeight:** `messageHeight` stores a function or callback value for temporary calculation. It is initialized from `messages.reduce((sum, message) => sum + message.getBoundingClientRect().height, 0) + messageGap`. Within the function it is assigned or updated later in the function.

**contentHeight:** `contentHeight` stores a computed value for temporary calculation. It is initialized from `messageHeight + chromeHeight`. Within the function it is passed into `round`; assigned or updated later in the function.

**maxHeight:** `maxHeight` stores a computed value for temporary calculation. It is initialized from `Math.min(window.innerHeight * 0.76, 720)`. Within the function it is passed into `round`; assigned or updated later in the function.

**nextHeight:** nextHeight stores a computed value for temporary calculation. It is initialized from `Math.round(Math.max(dockedHeight, Math.min(contentHeight, maxHeight)))`. Within the function it is passed into `setProperty, toggle`; assigned or updated later in the function.

### **function renderAssistantInline(value = "")**

To understand `renderAssistantInline`, start with the problem it owns. `renderAssistantInline` turns saved data for `AssistantInline` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderAssistantInline(value = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `value`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `value` is human-facing text. The function cleans, limits, displays, compares, or turns it into a safe identifier so the UI and generated site remain readable. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `renderInlineMath`. Those calls show that `renderAssistantInline` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `renderAssistantInline`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside `renderAssistantInline`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function renderAssistantAnswerContent(content = "")**

To understand `renderAssistantAnswerContent`, start with the problem it owns. `renderAssistantAnswerContent` turns saved data for `AssistantAnswerContent` into visible UI. It is a presentation function, but it also enforces public-site rules: empty content should not appear, rich content must be converted into safe HTML, and the generated markup must match the class names that `styles.css` expects.

The syntax of the function is:

```
function renderAssistantAnswerContent(content = "") { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `content`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `content` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a string that is already normalized for display, command output, a URL, a CSS value, or a file/path label.

Next, trace the helper calls that surround the function. It works with `renderAssistantInline`, `escapeHtml`. Those calls show that `renderAssistantAnswerContent` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `renderAssistantAnswerContent`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `renderAssistantAnswerContent` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**lines:** `lines` stores a computed value for temporary calculation. It is initialized from `String(content || "").replace(/r\n?/g, "\n").split("\n")`. Within the function it is iterated or transformed as a collection; mutated as a collection or DOM container; accesses properties/methods `forEach`, `join`, `push`; passed into `push`; assigned or updated later in the function.

**html:** `html` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is mutated as a collection or DOM container; accesses properties/methods `join`, `push`; assigned or updated later in the function.

**bulletItems:** `bulletItems` stores a list for temporary calculation. It is initialized from `[]`. Within the function it is iterated or transformed as a collection; mutated as a collection or DOM container; accesses properties/methods `length`, `map`, `push`; passed into `push`; assigned or updated later in the function.

**codeBlock:** `codeBlock` stores a computed value for temporary calculation. It is initialized from `null`. Within the function it is accesses properties/methods `language`, `lines`; passed into `escapeHtml`, `push`; assigned or updated later in the function.

**flushBullets:** `flushBullets` stores a function or callback value for temporary calculation. It is initialized from `() => {}`. Within the function it is assigned or updated later in the function.

**flushCode:** `flushCode` stores a function or callback value for temporary calculation. It is initialized from `() => {}`. Within the function it is assigned or updated later in the function.

**language:** `language` stores a computed value for compiler or language configuration. It is initialized from `codeBlock.language ? `data-language="${escapeHtml(codeBlock.language)}" : ""`. It supports the compile workspace by describing language rules, executable locations, tool status, or compiler output. Within the function it is passed into `escapeHtml`, `push`; assigned or updated later in the function.

**trimmed:** `trimmed` stores a computed value for temporary calculation. It is initialized from `line.trim()`. Within the function it is accesses properties/methods `match`, `slice`; passed into `push`; assigned or updated later in the function.

**fence:** fence stores a computed value for temporary calculation. It is initialized from `trimmed.match(/^``([a-z0-9_+.#-]*)s*$/i)`. Within the function it is assigned or updated later in the function.

**bullet:** bullet stores a computed value for temporary calculation. It is initialized from `trimmed.match(/^[-*\u2022]\s+(.+)/u)`. Within the function it is passed into `push`; assigned or updated later in the function.

**function `appendAssistantMessage(role, content, sources = [])`**

The best entry point for `appendAssistantMessage` is its responsibility in the surrounding workflow.

`appendAssistantMessage` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function appendAssistantMessage(role, content, sources = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `role`, `content`, `sources`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `role` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `content` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `sources` is source text supplied by the builder, project catalog, GitHub, or an uploaded file. The function inspects, formats, saves, compiles, highlights, or extracts meaning from this text. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: creates DOM elements instead of relying only on static HTML; writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

Its connection to the rest of the file matters as much as its local code. It works with `renderAssistantAnswerContent`, `assistantSourcesMarkup`, `renderMultilineInlineText`, `updateAssistantPanelGrowth`. Those calls show that `appendAssistantMessage` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `appendAssistantMessage`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The easiest way to follow the body of `appendAssistantMessage` is to watch its local values. They are included here because they explain the implementation rather than merely naming syntax.

**article:** article stores a DOM element reference. It is initialized from `document.createElement("article")`. Within the function it is returned to the caller; accesses properties/methods `className`, `innerHTML`; passed into `append`, `createElement`; assigned or updated later in the function.

### **function scrollAssistantAnswerToStart(article)**

Read `scrollAssistantAnswerToStart` first as a named idea, not as syntax. `scrollAssistantAnswerToStart` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function scrollAssistantAnswerToStart(article) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `article`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `article` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it uses a timer to avoid blocking or to wait for another process. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

The function also has to be read through the helpers it calls. No other named helper from this file is detected inside the body. That means `scrollAssistantAnswerToStart` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The application keeps this behavior isolated because other features rely on it being consistent. Without `scrollAssistantAnswerToStart`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `scrollAssistantAnswerToStart` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**scrollToAnswer:** `scrollToAnswer` stores a function or callback value for temporary calculation. It is initialized from `() => {}`. Within the function it is passed into `requestAnimationFrame`, `setTimeout`; assigned or updated later in the function.

### **function appendAssistantPendingMessage()**

Begin with the role of `appendAssistantPendingMessage`. `appendAssistantPendingMessage` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function appendAssistantPendingMessage() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

The body is where the contract above becomes behavior. Inside the body, it updates browser DOM. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. A close reading of the body shows these implementation details: toggles CSS classes to express UI state.

The surrounding workflow becomes clearer when you notice its collaborators. It works with `appendAssistantMessage`. Those calls show that `appendAssistantPendingMessage` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This is why the function is worth separating instead of burying the logic somewhere else. Without `appendAssistantPendingMessage`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local state inside `appendAssistantPendingMessage` explains how the larger task is split into smaller steps. They are included here because they explain the implementation rather than merely naming syntax.

**article:** article stores a computed value for temporary calculation. It is initialized from `appendAssistantMessage("assistant", "Thinking", [])`. Within the function it is returned to the caller; assigned or updated later in the function.

#### **function replaceAssistantMessage(article, content, sources = [])**

The best entry point for `replaceAssistantMessage` is its responsibility in the surrounding workflow. `replaceAssistantMessage` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function replaceAssistantMessage(article, content, sources = []) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `article`, `content`, `sources`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `article` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `content` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything. `sources` is source text supplied by the builder, project catalog, GitHub, or an uploaded file. The function inspects, formats, saves, compiles, highlights, or extracts meaning from this text. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

After the syntax, move into the body and follow the work in the order the function performs it. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.



Its connection to the rest of the file matters as much as its local code. It works with `appendAssistantMessage`, `renderAssistantAnswerContent`, `assistantSourcesMarkup`, `updateAssistantPanelGrowth`, `scrollAssistantAnswerToStart`. Those calls show that `replaceAssistantMessage` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

This function earns its place in the file because it protects one behavior from being rewritten differently elsewhere. Without `replaceAssistantMessage`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside `replaceAssistantMessage`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

### **function setAssistantBusy(isBusy)**

To understand `setAssistantBusy`, start with the problem it owns. `setAssistantBusy` belongs to `portfolio ai frontend`. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function setAssistantBusy(isBusy) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `isBusy`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `isBusy` is the caller-supplied input used by this function. The function interprets it according to its local responsibility, often converting it into safer, more predictable data before returning or rendering anything.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state.

Next, trace the helper calls that surround the function. No other named helper from this file is detected inside the body. That means `setAssistantBusy` performs its rule directly from parameters, module state, standard APIs, or local calculations.

The practical importance of the function is easiest to see by imagining it removed. Without `setAssistantBusy`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `setAssistantBusy` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**submitButton:** `submitButton` stores a DOM element reference. It is initialized from `aiAssistantForm?.querySelector('button[type="submit"]')`. It anchors UI behavior to a specific browser element or window state. Within the function it is accesses properties/methods disabled; assigned or updated later in the function.

## **function clearAssistantChat()**

To understand `clearAssistantChat`, start with the problem it owns. `clearAssistantChat` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
function clearAssistantChat() { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The empty parameter list means the function works from module-level state, browser state, local constants, or values it creates internally. That is common for UI helpers that operate on known page elements and backend helpers that read already configured paths.

Those syntax pieces become practical once you connect them to the data being passed in. This function accepts no explicit parameters. It reads controlled module state, browser state, local configuration, or constants that are already available in the file.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it updates browser DOM. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. A close reading of the body shows these implementation details: writes generated markup into the page, which is why escaping and sanitizing helpers around it matter.

Next, trace the helper calls that surround the function. It works with `setAssistantStatus`, `updateAssistantPanelGrowth`. Those calls show that `clearAssistantChat` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `clearAssistantChat`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

Inside `clearAssistantChat`, there are no local `const`, `let`, or `var` values to discuss. The function mainly works from its parameters, module-level state, direct expressions, or helpers it calls.

## **function assistantRemoteTimeoutMs(question = "", context = {}, options = {})**

To understand `assistantRemoteTimeoutMs`, start with the problem it owns. `assistantRemoteTimeoutMs` is one piece of the Ask My Portfolio conversation system. It helps decide what the user means, what evidence belongs in context, how answers should be displayed, or how the chat panel should behave.

The syntax of the function is:

```
function assistantRemoteTimeoutMs(question = "", context = {}, options = {}) { ... }
```

There is no `async` keyword, so the function completes synchronously unless it calls another asynchronous helper indirectly. The parameter list tells you what the caller must provide: `question`, `context`, `options`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `context` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated

decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field. `options` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. The body is mostly deterministic transformation logic: it reads its inputs, normalizes or scores them, and returns a value the caller can trust. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere.

Next, trace the helper calls that surround the function. It works with `normalize`. Those calls show that `assistantRemoteTimeoutMs` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `assistantRemoteTimeoutMs`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `assistantRemoteTimeoutMs` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**clean:** `clean` stores a computed value for temporary calculation. It is initialized from `normalize(question)`. Within the function it is passed into `test`; assigned or updated later in the function.

**sourceFetches:** `sourceFetches` stores a computed value for temporary calculation. It is initialized from `Array.isArray(context.sourceFetches) ? context.sourceFetches : []`. Within the function it is accesses properties/methods `some`; passed into `isArray`; assigned or updated later in the function.

**usesGitHub:** `usesGitHub` stores a function or callback value for Git publishing and authorization state. It is initialized from `sourceFetches.some((source) => String(source.kind || "").toLowerCase() === "github" || /github\.com/i.test(source.url || ""))`. It affects publishing, where mistakes can change the public website or expose the wrong files. Within the function it is assigned or updated later in the function.

### **async function askRemoteAssistant(question, context, options = {})**

To understand `askRemoteAssistant`, start with the problem it owns. `askRemoteAssistant` belongs to portfolio ai frontend. This function belongs to the public chatbot. It classifies questions, gathers context, renders answers, or manages the chat UI. In this role, it owns one named operation that later code depends on.

The syntax of the function is:

```
async function askRemoteAssistant(question, context, options = {}) { ... }
```

The `async` keyword is important here: this function returns a `Promise`, so callers must wait for it before trusting the result or assuming the side effect has completed. The parameter list tells you what the caller must provide: `question`, `context`, `options`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. `context` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. `options` is a structured object rather than one small value. It lets the caller pass several related settings or pieces of state together so the function can make one coordinated decision. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

Once the inputs are clear, the implementation shows how the function turns those inputs into useful program state. Inside the body, it waits for asynchronous work, calls a network resource and uses a timer to avoid blocking or to wait for another process. It returns a computed value used immediately by a caller. The important point is that the caller does not repeat this rule elsewhere. Because it is async, callers must await it or handle its Promise; otherwise the UI or backend could move on before files, network calls, Git commands, compiler runs, or model responses have finished. A close reading of the body shows these implementation details: serializes structured data before sending or saving it; throws explicit errors when a required safety or compatibility condition fails; performs a fetch and therefore must handle network delay and failure.

Next, trace the helper calls that surround the function. It works with `assistantEndpoint`, `assistantRemoteTimeoutMs`. Those calls show that `askRemoteAssistant` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The practical importance of the function is easiest to see by imagining it removed. Without `askRemoteAssistant`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

After the main flow, the working values inside `askRemoteAssistant` reveal what the function must remember while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**endpoint:** `endpoint` stores a computed value for network or routing value. It is initialized from `assistantEndpoint()`. Within the function it is passed into `fetch`; assigned or updated later in the function.

**controller:** `controller` stores a constructed object for temporary calculation. It is initialized from `new AbortController()`. Within the function it is accesses properties/methods `abort`, `signal`; passed into `fetch`; assigned or updated later in the function.

**didTimeout:** `didTimeout` stores a boolean value for temporary calculation. It is initialized from `false`. Within the function it is assigned or updated later in the function.

**timer:** `timer` stores a function or callback value for temporary calculation. It is initialized from `setTimeout(() => {`. Within the function it is passed into `clearTimeout`; assigned or updated later in the function.

**response:** `response` starts without an immediate value. In this file that usually means it is filled by a loop, branch, callback, or later assignment inside the same function. Within the function it is accesses properties/methods `json`, `ok`; assigned or updated later in the function.

**data:** `data` stores the completed result of an awaited operation. It is initialized from `await response.json()`. Within the function it is accesses properties/methods `answer`, `message`, `model`, `usedWebSearch`; passed into `Boolean`, `String`; assigned or updated later in the function.

**answer:** `answer` stores a computed value for temporary calculation. It is initialized from `String(data.answer || data.message || "").trim()`. Within the function it is returned to the caller; passed into `String`; assigned or updated later in the function.

### **async function answerAssistantQuestion(question = "")**

Read `answerAssistantQuestion` first as a named idea, not as syntax. `answerAssistantQuestion` is the public chatbot orchestrator. It receives the user's typed question, appends it to chat history, classifies the intent, gathers relevant portfolio context, creates a pending assistant answer, calls the backend AI when appropriate, falls back to local answers when needed, renders sources only when relevant, and restores focus to the input. This function is async because the AI backend may need network time and because context gathering can include indexed file text. It is the reason Ask My Portfolio behaves like a conversation instead of a static search form.

The syntax of the function is:

```
async function answerAssistantQuestion(question = "") { ... }
```

The `async` keyword is important here: this function returns a Promise, so callers must wait for it before trusting the result or assuming the side effect has completed. The parameter list tells you what the caller must provide: `question`. Read those names as the contract between this function and the code that calls it. The body should make sense only after those inputs are understood.

Those syntax pieces become practical once you connect them to the data being passed in. `question` is user language: a search query, AI question, or classified assistant intent. It drives scoring, context selection, routing, or AI response behavior. The default value is intentional: it lets the function fail softly or produce a predictable empty result when the caller omits that field.

With the signature understood, the implementation can be read as a sequence of decisions and side effects. Inside the body, it waits for asynchronous work. It mostly works through side effects, such as updating the DOM, writing files, sending an HTTP response, changing history, or mutating controlled runtime state. Because it is `async`, callers must await it or handle its Promise; otherwise the UI or backend could move on before files, network calls, Git commands, compiler runs, or model responses have finished.

The function also has to be read through the helpers it calls. It works with `appendAssistantMessage`, `assistantRemember`, `appendAssistantPendingMessage`, `setAssistantBusy`, `setAssistantStatus`, `assistantEndpoint`, `assistantQuestionIntent`, `assistantFallbackResults`. Those calls show that `answerAssistantQuestion` is not isolated; it is one piece of a larger workflow where helpers clean inputs, perform side effects, render results, or protect boundaries before the next step runs.

The application keeps this behavior isolated because other features rely on it being consistent. Without `answerAssistantQuestion`, the assistant would either answer with less context, display unrelated links, lose conversation memory, fail to format messages cleanly, or expose model/provider details in the wrong layer.

The local objects and variables inside `answerAssistantQuestion` show how the implementation holds temporary work while it runs. They are included here because they explain the implementation rather than merely naming syntax.

**cleanQuestion:** `cleanQuestion` stores a computed value for temporary calculation. It is initialized from `String(question || "").trim()`. Within the function it is passed into `appendAssistantMessage`, `askRemoteAssistant`, `assistantContextForQuestion`, `assistantConversationForRequest`, `assistantFallbackResults`, `assistantLocalAnswer`, `assistantQuestionIntent`, `assistantRemember`; assigned or updated later in the function.

**priorConversation:** `priorConversation` stores a computed value for temporary calculation. It is initialized from `assistantChatHistory.slice()`. Within the function it is passed into `assistantConversationForRequest`, `assistantQuestionIntent`; assigned or updated later in the function.

**pendingMessage:** `pendingMessage` stores a computed value for temporary calculation. It is initialized from `appendAssistantPendingMessage()`. Within the function it is passed into `replaceAssistantMessage`; assigned or updated later in the function.

**intent:** `intent` stores a computed value for temporary calculation. It is initialized from `assistantQuestionIntent(cleanQuestion, priorConversation)`. Within the function it is passed into `askRemoteAssistant`, `assistantContextForQuestion`, `assistantFallbackResults`, `assistantLocalAnswer`, `assistantSourcesForDisplay`; assigned or updated later in the function.

**rawSources:** `rawSources` stores a computed value for temporary calculation. It is initialized from `assistantFallbackResults(cleanQuestion, intent)`. Within the function it is passed into `assistantContextForQuestion`, `assistantSourcesForDisplay`; assigned or updated later in the function.

**sources:** sources stores a computed value for temporary calculation. It is initialized from `assistantSourcesForDisplay(cleanQuestion, rawSources, intent)`. Within the function it is passed into `assistantLocalAnswer`, `replaceAssistantMessage`; assigned or updated later in the function.

**context:** context stores a computed value for temporary calculation. It is initialized from `assistantContextForQuestion(cleanQuestion, intent, rawSources)`. Within the function it is passed into `askRemoteAssistant`; assigned or updated later in the function.

**conversation:** conversation stores a computed value for temporary calculation. It is initialized from `assistantConversationForRequest(cleanQuestion, priorConversation)`. Within the function it is assigned or updated later in the function.

**remoteAnswer:** remoteAnswer stores the completed result of an awaited operation. It is initialized from `await askRemoteAssistant(cleanQuestion, context, {`. Within the function it is accesses properties/methods `answer`, `model`; passed into `assistantRemember`, `replaceAssistantMessage`, `setAssistantStatus`; assigned or updated later in the function.

**fallbackAnswer:** fallbackAnswer stores a computed value for temporary calculation. It is initialized from `assistantLocalAnswer(cleanQuestion, sources, intent)`. Within the function it is passed into `assistantRemember`, `replaceAssistantMessage`; assigned or updated later in the function.

**errorAnswer:** errorAnswer stores a string value for temporary calculation. It is initialized from `"I had trouble processing that question. Try rephrasing it with a project name, tool name, or electronics topic."`. Within the function it is passed into `assistantRemember`, `replaceAssistantMessage`; assigned or updated later in the function.